

Dieter Schmalstieg


# The Graphics Pipeline

# What do we want?

- Computer-generated imagery (CGI) of complex 3D scenes in real-time
- Computationally extremely demanding
  - Full HD at 60 Hz:  
$$1920 \times 1080 \times 60 \text{ Hz} = 124 \text{ Mpx/s}$$
  - And that's just output data!
- Requires specialized hardware

## FROM SD TO 8K - EXPONENTIAL PIXEL GROWTH

Image Courtesy: Rose Adler, Leighana Ginther, Jackie Osterday




# Image Synthesis

## Object Order

Go through all objects

„What do I see where“




→ Rasterization

## Image Order

Go through all pixels

„Where do I see what“




→ Raytracing

# Rasterization Hardware

Most of real-time graphics is based on

- Rasterization of graphic *primitives*

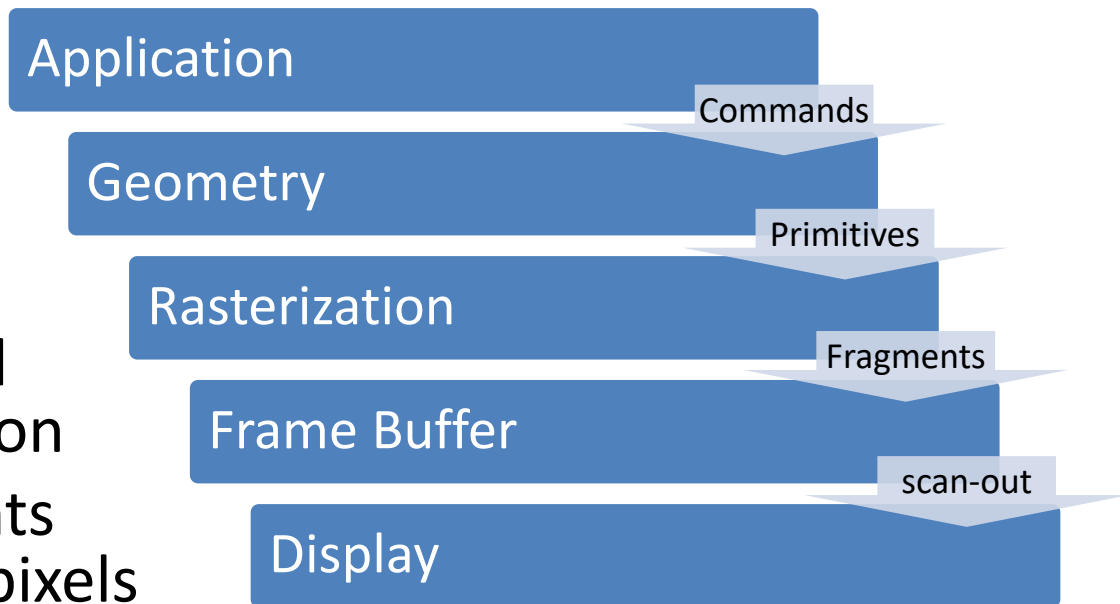
- Points
- Lines
- Triangles
- ...



- Implemented in hardware
  - *Graphics processing unit* (GPU)

# The Graphics Pipeline

- High-level view:
- “Fragment”:
  - Sample produced during rasterization
  - Multiple fragments are *merged* into pixels



# Application Stage

- Generate database
  - Usually only once
  - Load from disk or generate algorithmically
  - Build acceleration structures (hierarchy, ...)
- Repeat main loop
  - Input event handlers
  - Simulation → modify data structures
  - Database traversal → issue **graphics commands**
- Until exit



# Graphics Commands

- Graphics command stream from CPU to GPU
  - Specify primitives
  - Manage resources
  - Modify GPU state
- Commands are abstracted as Graphics API
  - OpenGL, DirectX, Vulkan, Metal

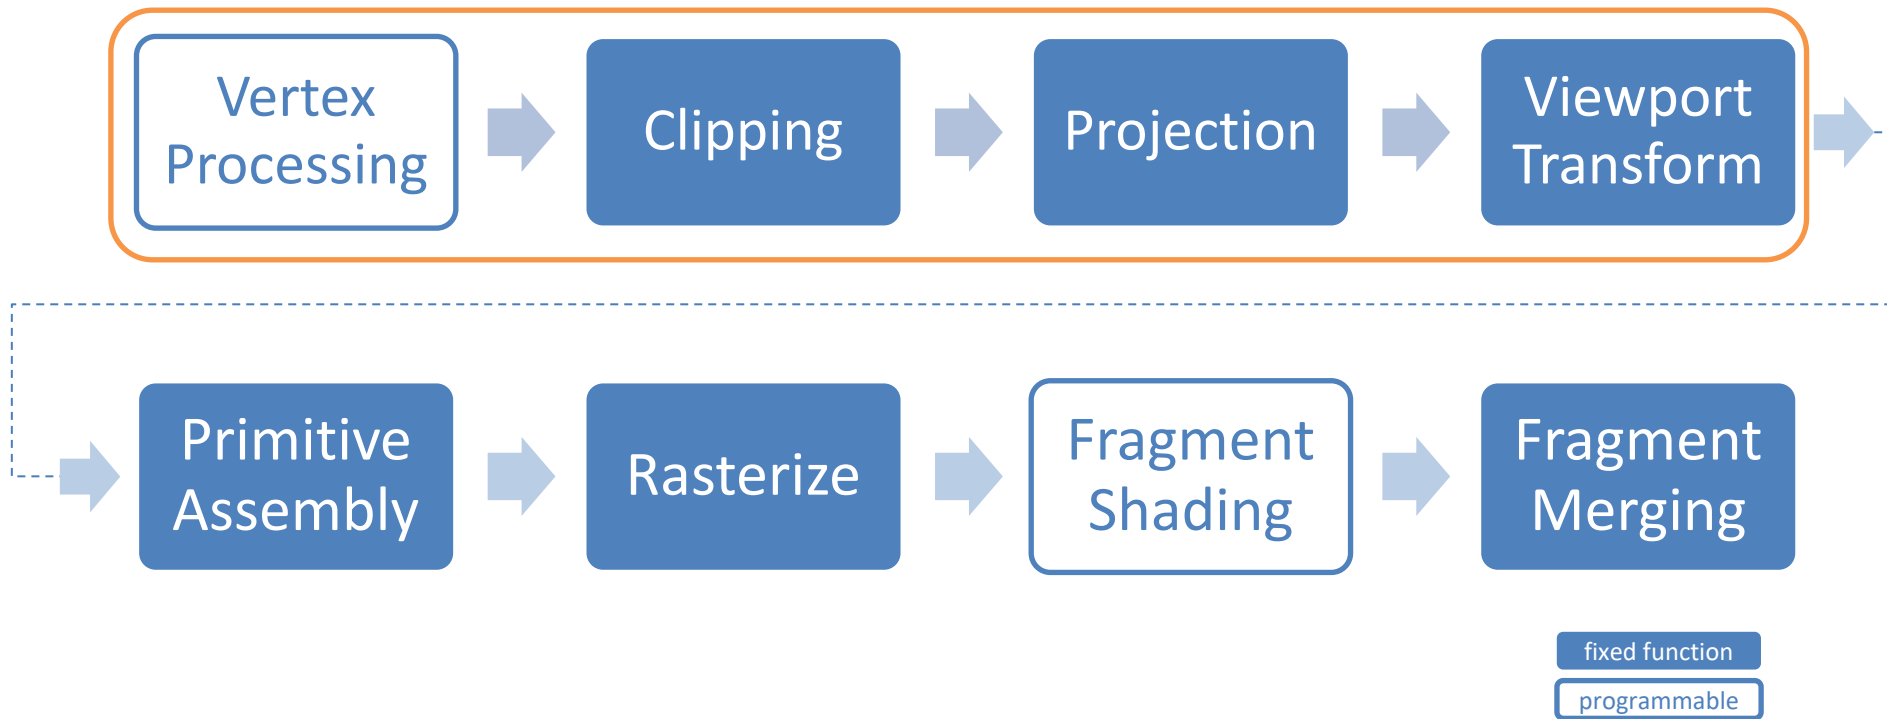


# Graphics Driver

- Graphics hardware is shared resource
- Large user mode graphics driver
  - Prepares command buffers
- Graphics kernel subsystem
  - Schedule access to hardware
- Small kernel mode graphics driver
  - Submit command buffers to hardware


This is were new APIs  
(Apple Metal, DirextX 12, Vulkan)  
try to be more efficient

# Geometry Stage



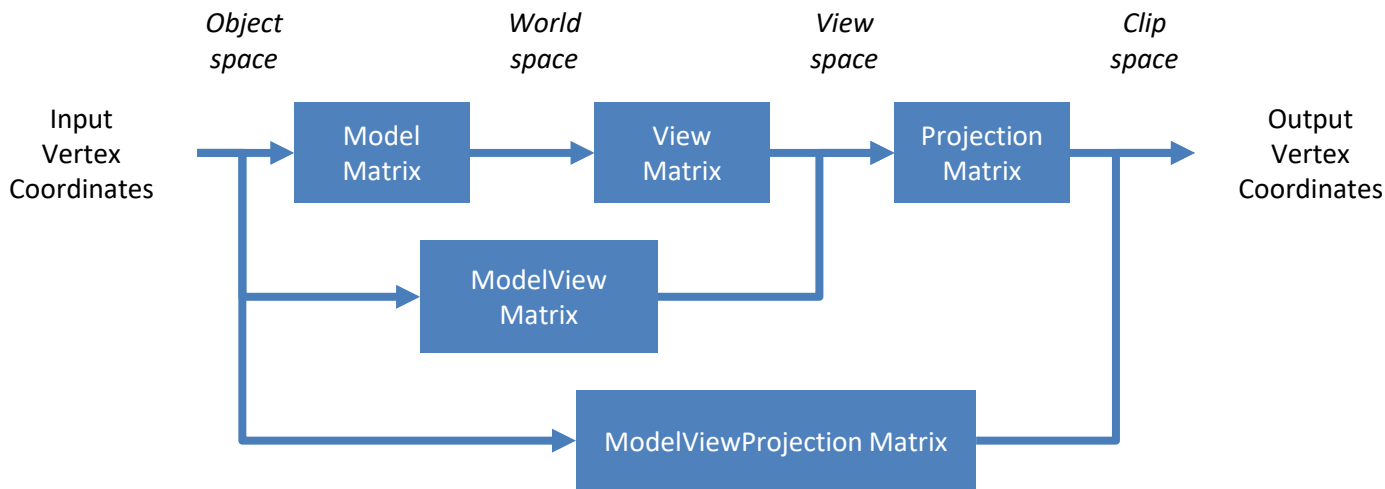
# Vertex Processing

- Input vertex stream
  - Composed of arbitrary *vertex attributes* (position, color, ...)
- Is transformed into stream of vertices mapped onto the screen
  - Composed of their *clip space* coordinates and additional user-defined attributes (color, texture coordinates, ...)
  - Clip space: homogeneous coordinates
- By the *vertex shader*
  - GPU program that implements this mapping
  - Historically, “shaders” were small programs performing lighting calculations




# Vertex Coordinate Transformation

Common model in rasterization-based 3D graphics

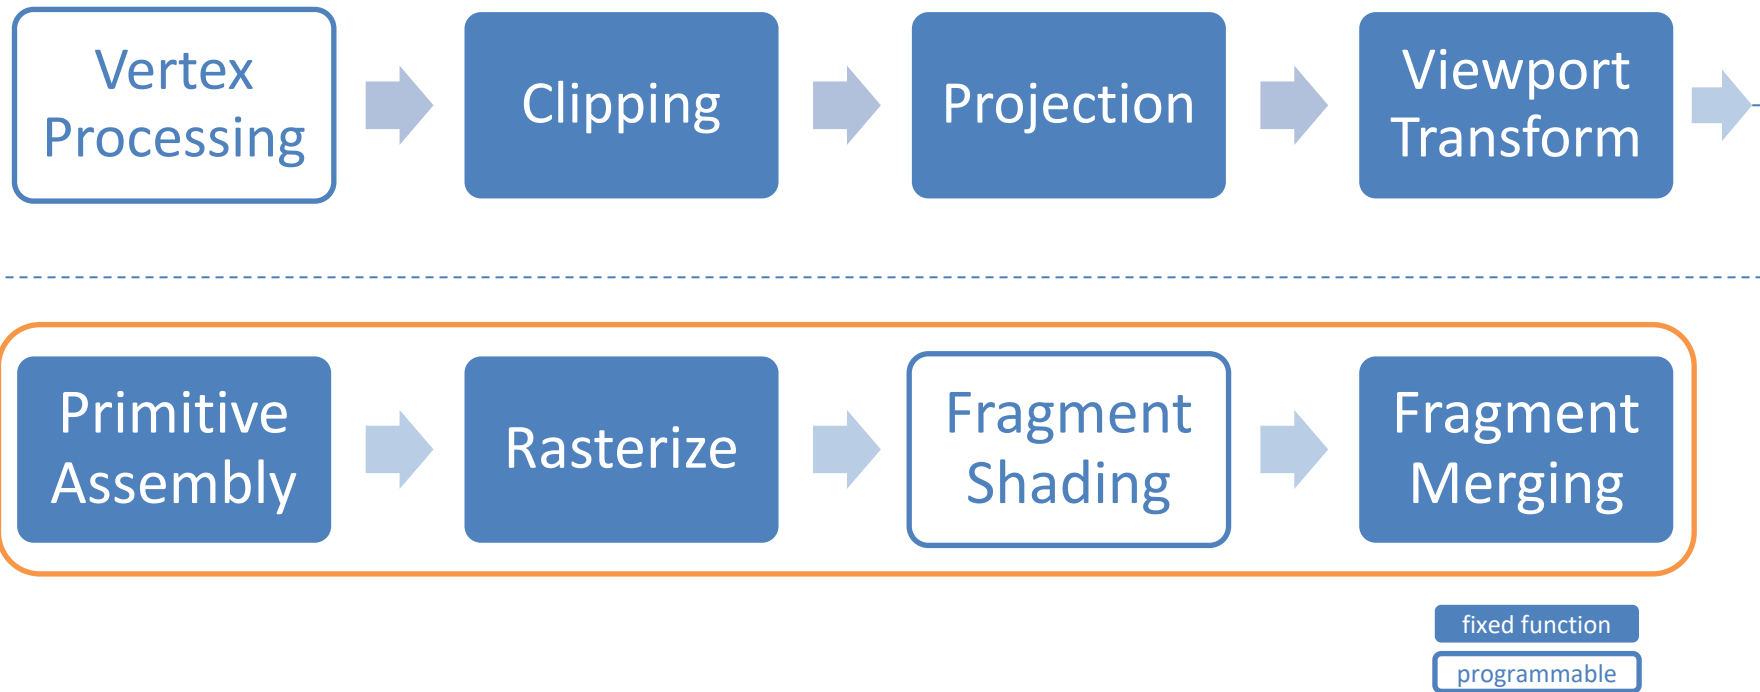


# Geometry Stage Tasks

- Clipping
  - Primitives not entirely in view are clipped to avoid projection errors
- Projection
  - Projects clip space coordinates to the image plane
  - Primitives in *normalized device coordinates*
- Viewport Transform
  - Maps resolution-independent normalized device coordinates to a rectangular window in the frame buffer, the *viewport*
  - Primitives in window (pixel) coordinates




# Rasterization Stage




# Rasterization Stage Tasks

- Primitive assembly
  - Backface culling
  - Setup primitive for traversal
- Rasterization (“primitive traversal”, “scan conversion”)
  - Sampling (triangle → fragments)
  - Interpolation of vertex attributes (depth, color, ...)
- Fragment shading
  - Compute fragment colors
- Fragment merging
  - Compute pixel colors from fragments
  - Depth test, blending, ...






# Rasterization – Coordinates




# Rasterization – Rules

- Different rules apply for each primitive type
  - “Fill convention”
- Non-ambiguous!
  - Avoids artifacts
- Polygons
  - Pixel center contained in polygon
  - Pixels on edge: only one rasterized




# Fragment Shading

- Aka *pixel shader*
- Given the interpolated vertex attributes
  - Output by the vertex shader
- The *fragment shader* computes color values for each fragment
  - Apply textures
  - Lighting calculations
  - ...

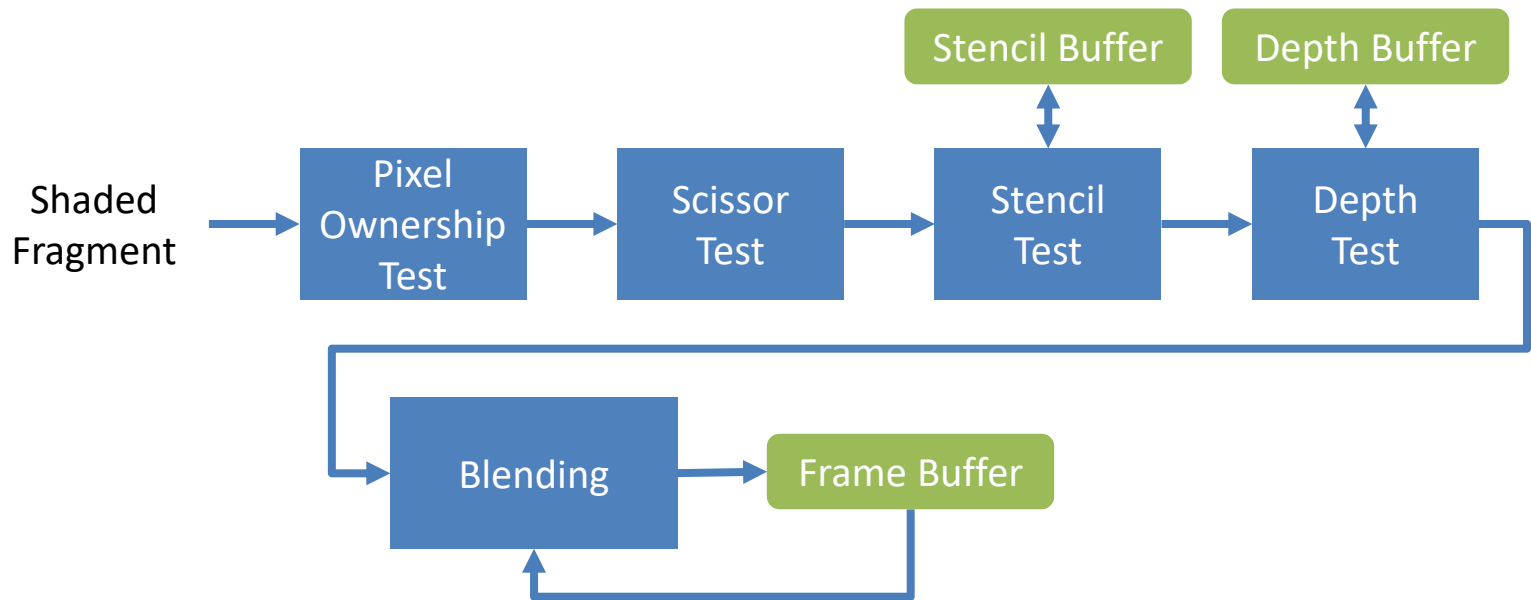


# Fragment Merging

- Also known as “raster operations” (ROP)
- Multiple primitives can cover the same pixel
- Their fragments need to be composed to form the final pixel values
  - Blending
  - Resolve visibility via depth buffering

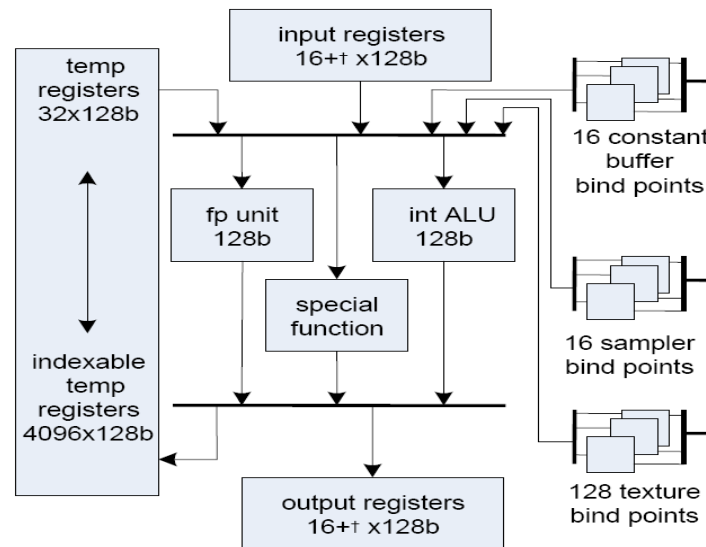


# Fragment Merging Workflow



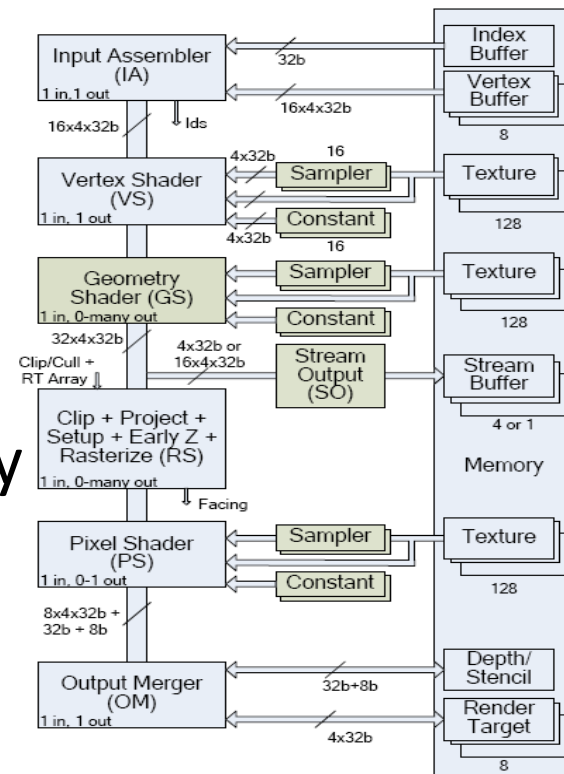
# Unified Shader Model

- Same instruction set and capabilities for all shaders
- Dynamic load balancing between vertex and fragment shaders
- IEEE-754 floating point
- Enables new GPGPU languages like CUDA



# Geometry Shader


- Between vertex and pixel shader
- Can generate primitives dynamically
- Procedural geometry
  - E.g., growing plants
- Geometry shader can write to memory
  - Called „stream output“
  - Enables multi-pass for geometry







# Geometry Shader Examples 1

Cube Map: GS instances  
every triangle 6x



Plants created as  
parameterized instances




Particles produced by  
GS as stream-out

# Geometry Shader Examples 2



Shadow volume created  
by GS extrusion

Displacements created by GS



# Tessellation Shader

- Since DirectX11
- New shader stage: Tessellation
- Input: low-detail mesh
- Output: high-detail mesh

Vertex Shader

Hull Shader

Tessellator

Domain Shader

Geometry Shader

Rasterizer

Pixel Shader




[https://www.youtube.com/watch?v=p\\_VpAMaxwpY](https://www.youtube.com/watch?v=p_VpAMaxwpY)

# Authoring without Tessellation

Sub-D modeling




Animation




Displacement map



Polygon mesh



Generate LODs



# Authoring with Tessellation


Sub-D modeling




Animation



Displacement map

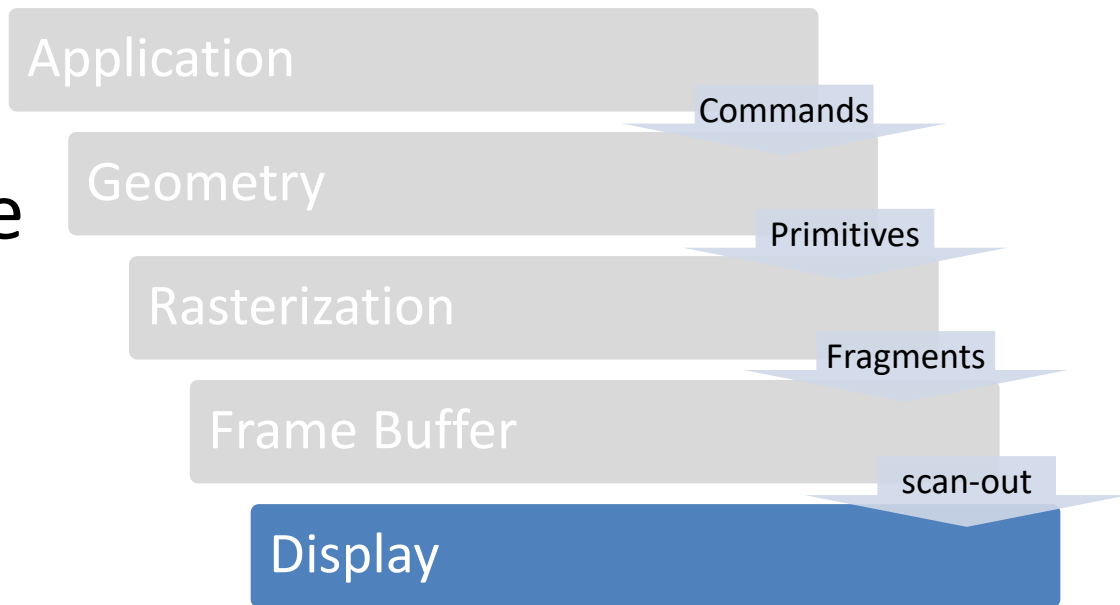


Optimally tessellated mesh




# Display Stage

- What happens after fragments are retired to the frame buffer?



# Display Stage Tasks

- Gamma correction
- Historically: digital to analog conversion
- Today: digital scan-out, HDMI encryption?



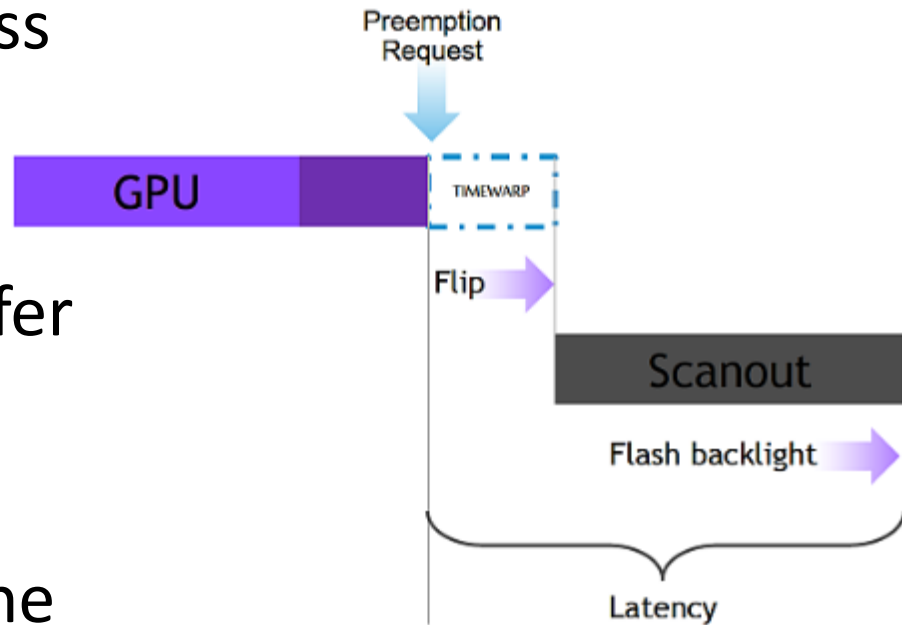


# Display Format

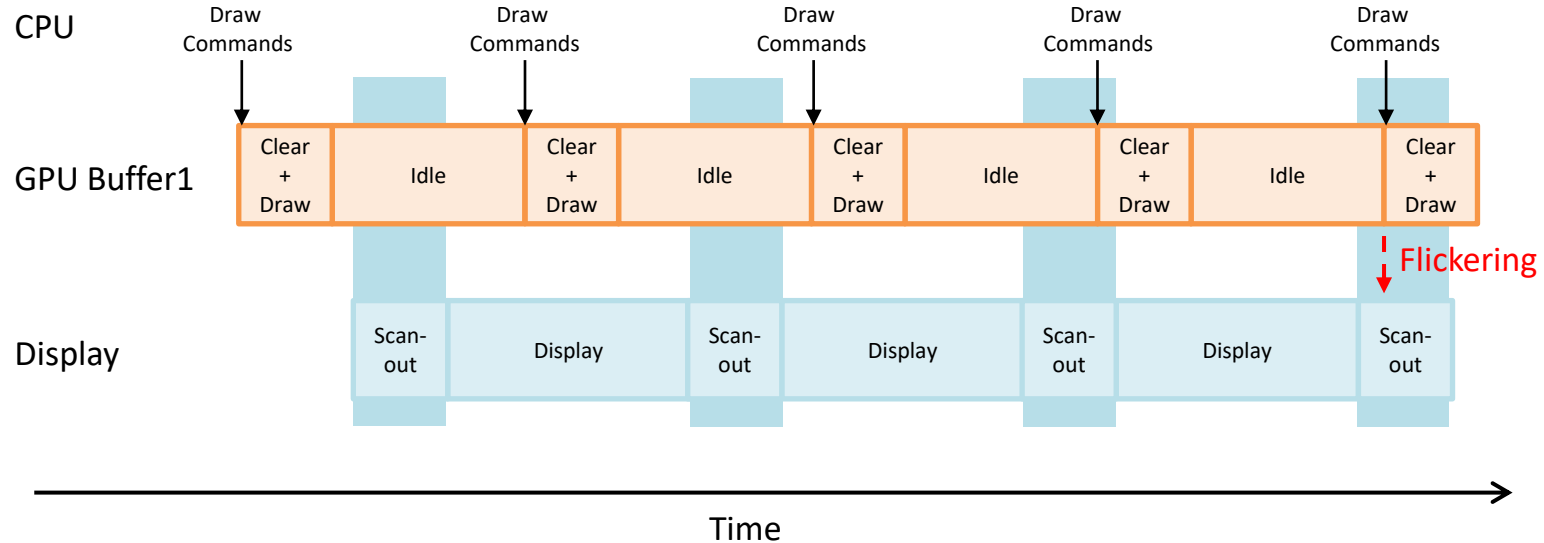
- Frame buffer pixel format:  
RGBA vs. index (obsolete)
- Bits: 16, 32, 64, 128 bit floating point, ...
- Double buffering for smooth animation
- Quad-buffering for stereo graphics
- Overlays (extra bitplanes)
- Auxilliary buffers: alpha, stencil, depth

# What is Display Synchronization?

- Need to synchronize access to frame buffer between GPU and display
- Cannot change frame buffer during display scan-out
- Need proper buffering in the whole graphics pipeline

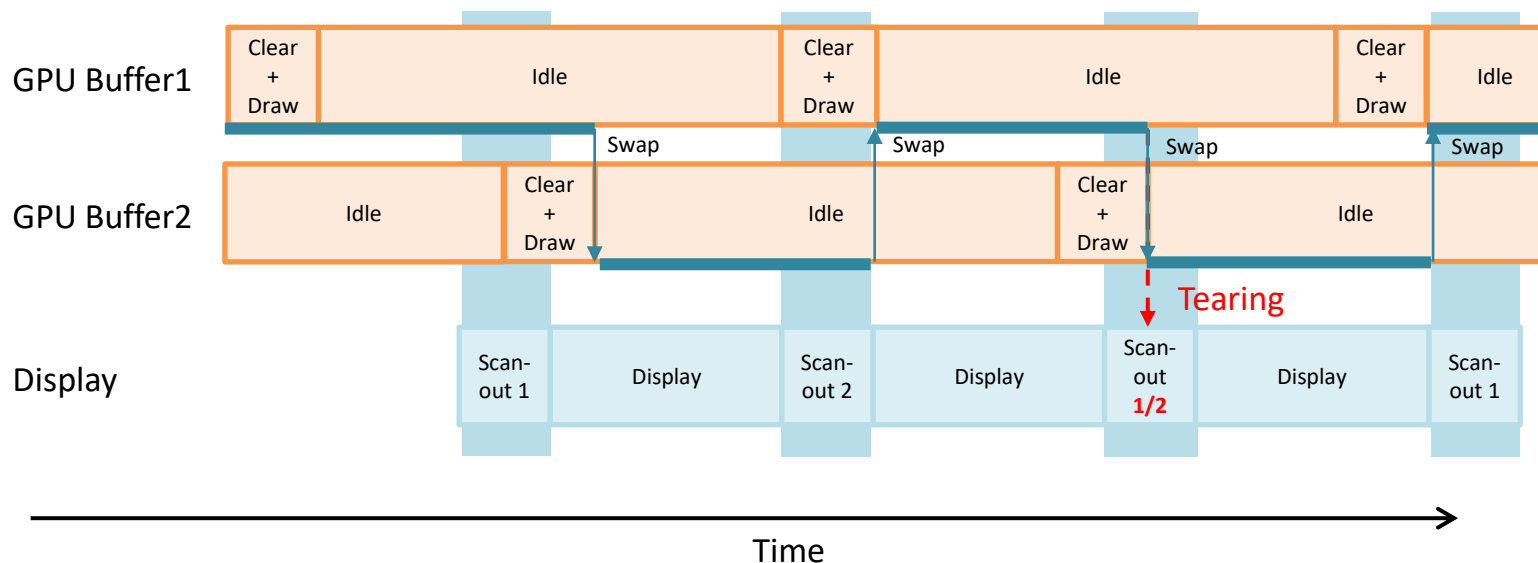


# Single Buffering



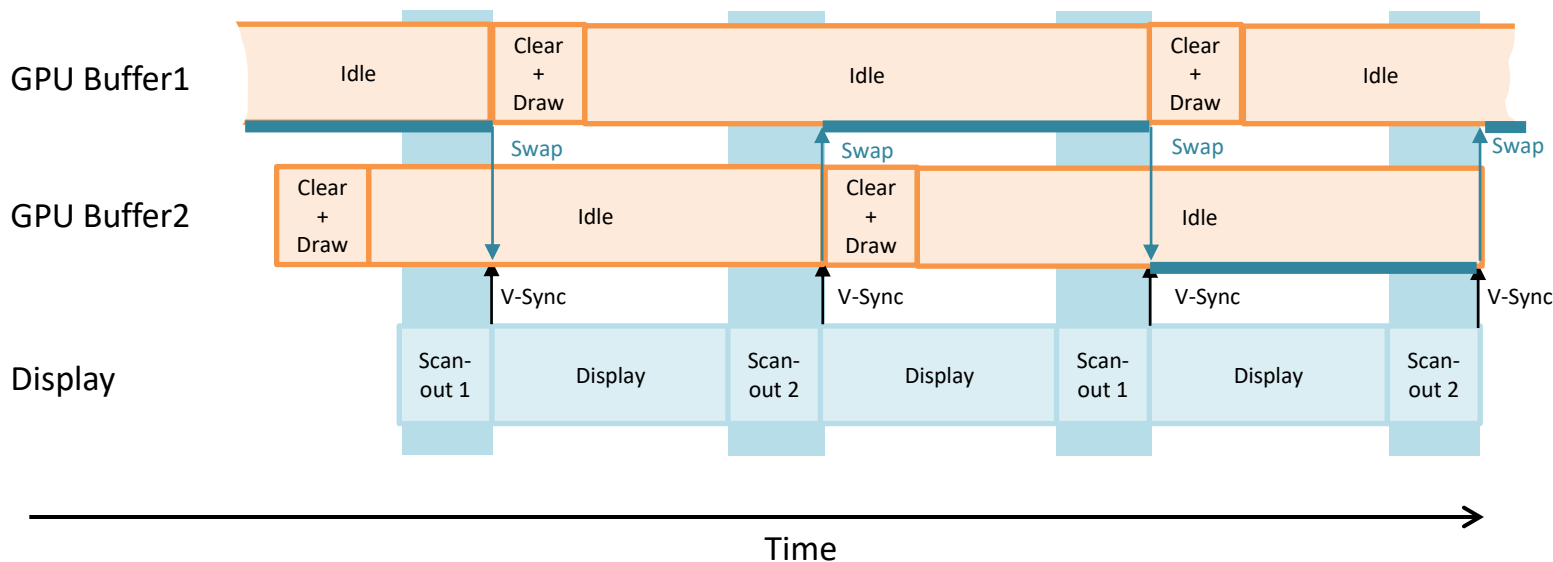
# Double Buffering without V-Sync

- Front buffer used for scan-out GPU → display
- SwapBuffers() changes front and back buffer
- No flickering, but **tearing** artefacts



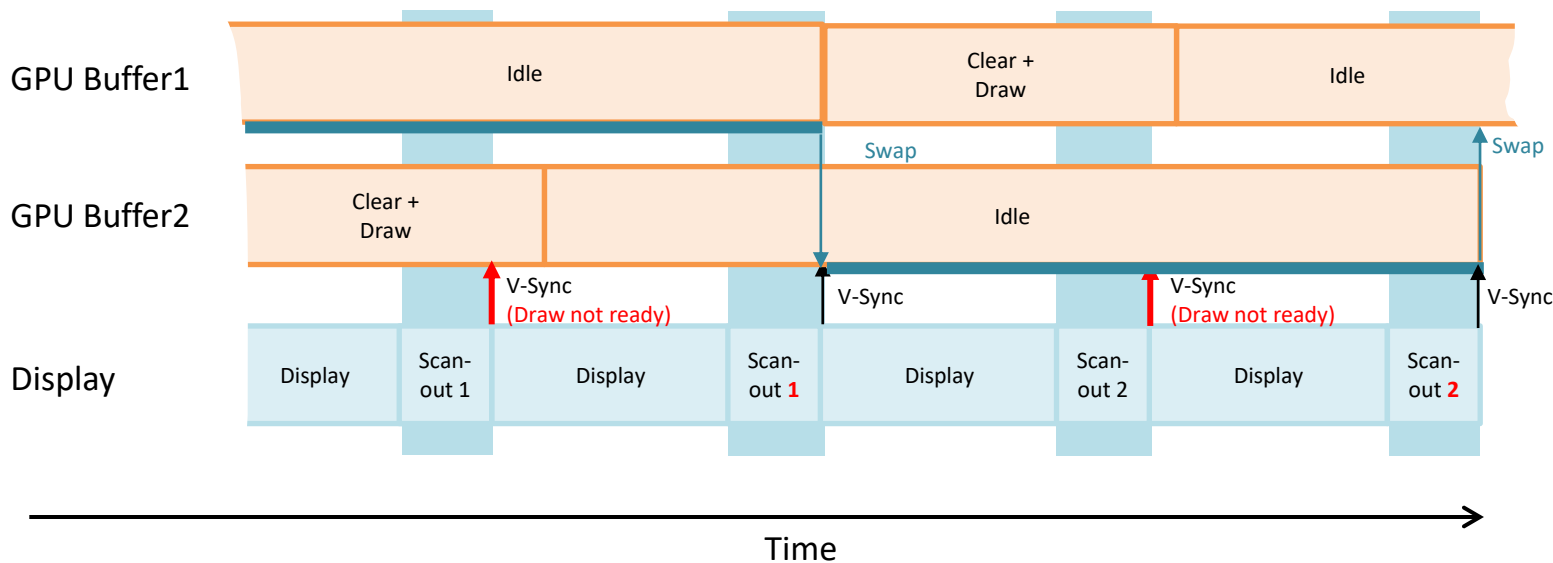
# Double Buffering with V-Sync, Fast

- V-Sync means the display is done with frame scan-out from GPU
- When rendering is **fast**, the frame rate is limited by display rate
- **Additional latency** of max. one frame time



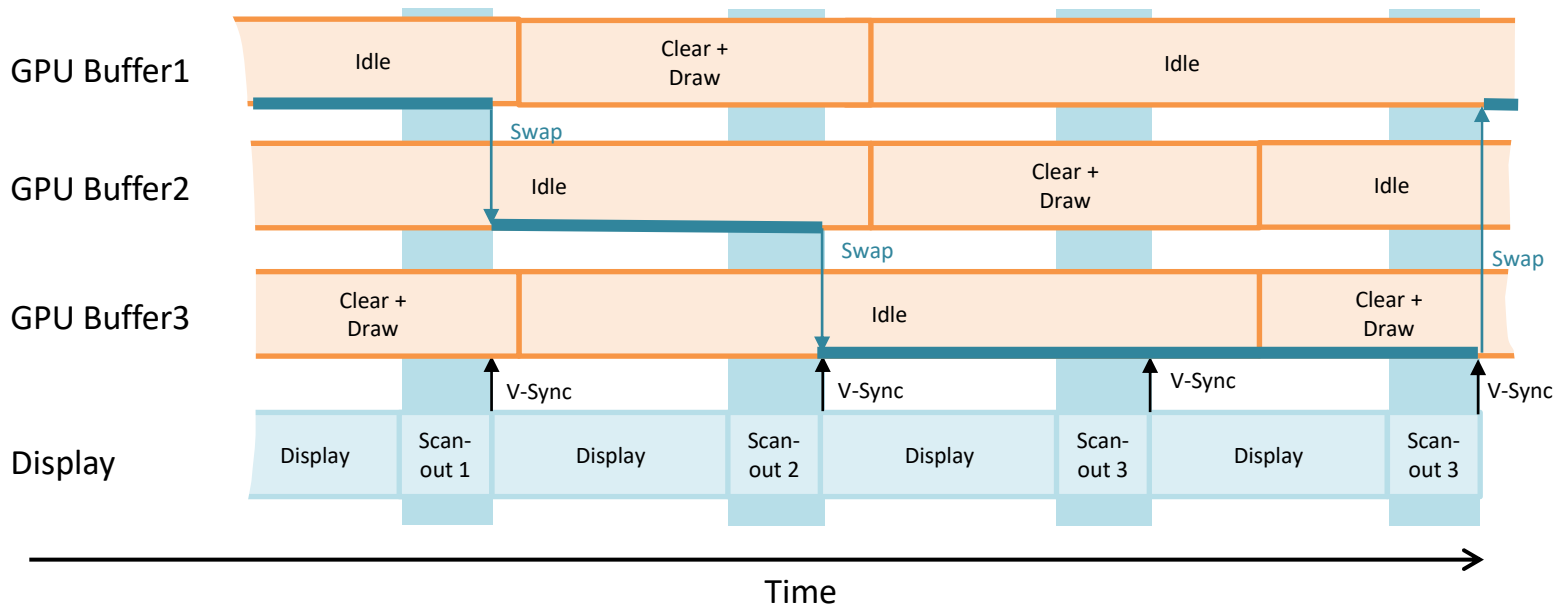
# Double Buffering with V-Sync, Slow

- When rendering is **slow**, frame rates can **only be integer fractions of display rate**
- Display rate 60Hz → frame rates 60, 30, 20, 15, ... (but not 59) possible
- Adaptive V-Sync (NVIDIA) turns off V-Sync automatically if rendering is slow



# Triple Buffering with V-Sync

- Triple buffering only makes sense together with V-Sync
- When rendering is **slow**, third buffer allows **continuous drawing**
- When rendering is **fast**, two back-buffers can be alternated (**wasted** frames!)



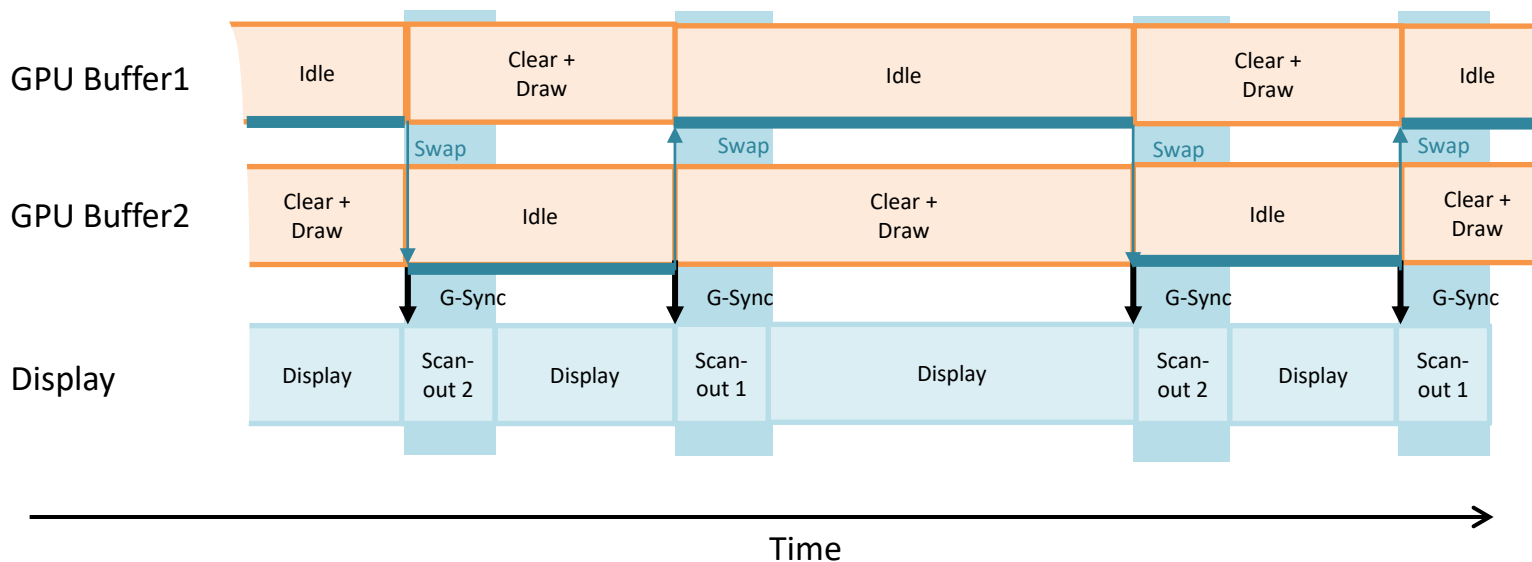


# G-Sync / FreeSync

- Display scan-out can be triggered by GPU
- Requires additional display hardware feature
- G-Sync (NVIDIA), FreeSync (AMD)

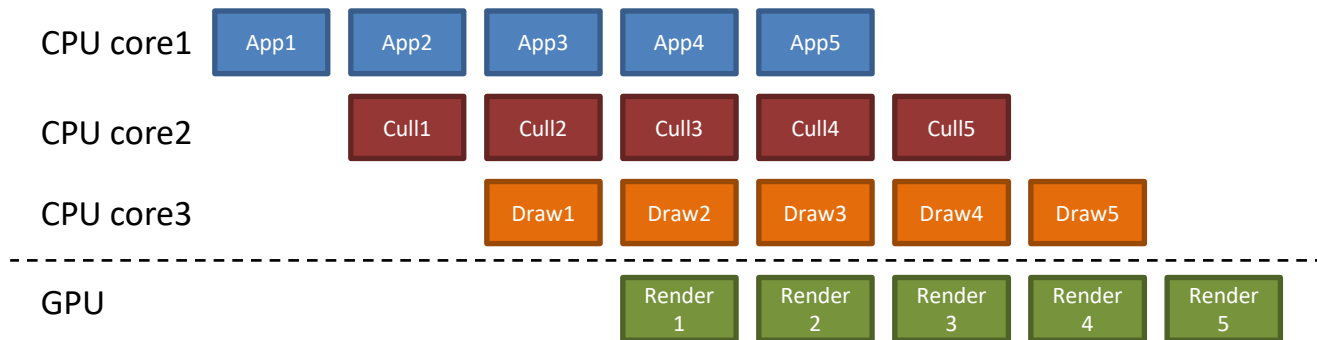


<https://youtu.be/5maHG9Kpjic>



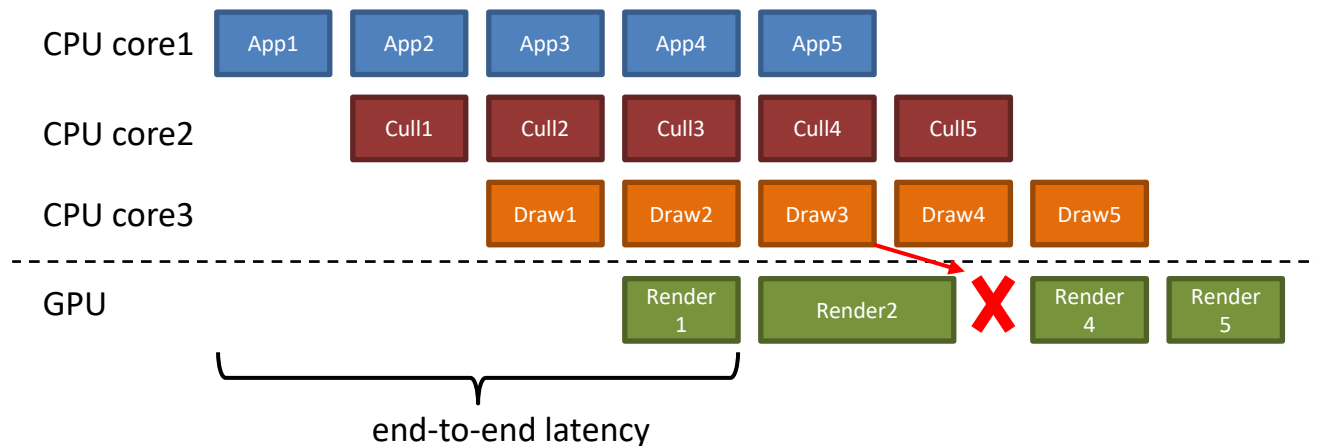
# Multi-Threaded Rendering Pipeline

- App stage: simulate 3D world
- Cull stage: determine object in view frustum
- Draw stage: issue OpenGL commands to driver (includes optimizations such as mode sorting)
- Everything must work at target frame rate!



# Minimizing Pipeline Latency

- Always aim for minimal latency
- Fixed size buffer from stage to stage
- Never wait for (downstream) consumer!



# Thank You!

# Questions ?