

Dieter Schmalstieg

Shader Programming

Today's Agenda

- The GPU execution model
- How to write shader programs

Graphics Application Programmer Interface

- Hardware-vendor independent interface
 - Interface is hardware independent
 - But implementation is hardware dependent
- Defines
 - Abstract rendering device
 - Set of functions to operate the device

API and Vendor Overview

Graphics API

- DirectX (Microsoft)
- OpenGL
- OpenGL ES
- Vulkan (Khronos group)
- Metal (Apple)

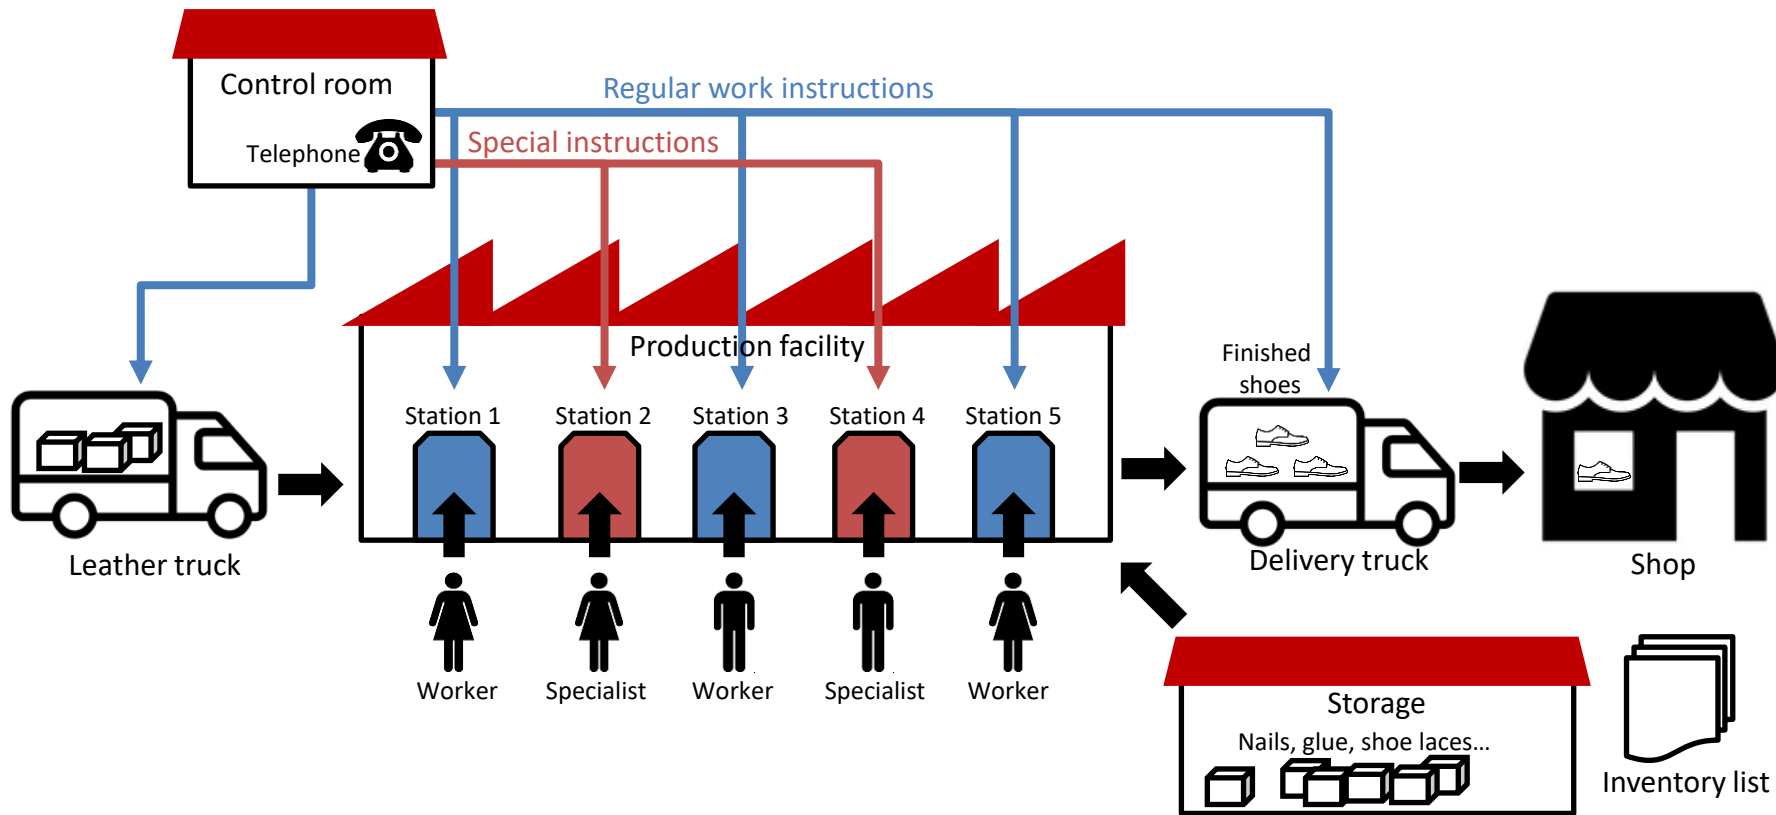
GPU Hardware Vendors

- Intel (Iris, Xe)
- NVIDIA (GeForce)
- AMD (Radeon)
- Qualcomm (Adreno)
- Imagination (PowerVR)
- Apple
- ARM (Mali, only design)

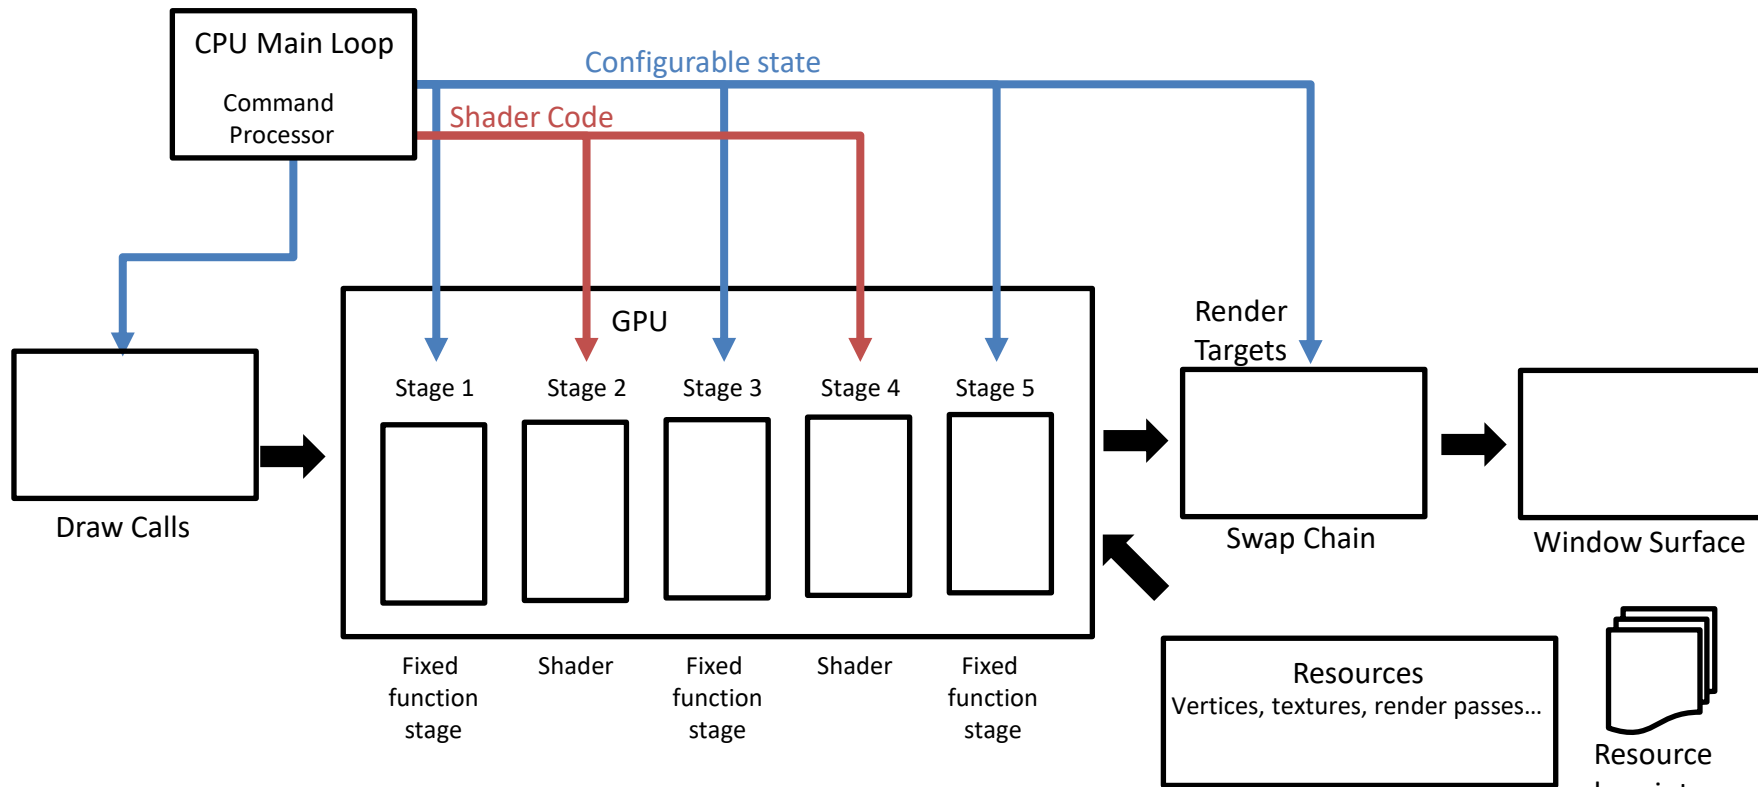
A Story About GPU Programming

- Imagine you want to produce shoes
- CPU, single-threaded
 - The master shoemaker is alone in the shop
 - One task is done after the other
- CPU, multi-threaded
 - The master shoemaker has a few apprentices
 - They talk directly to one another to split the work
- GPU: 10000 worker threads...?

A Big Shoe Factory



A Big Graphics API

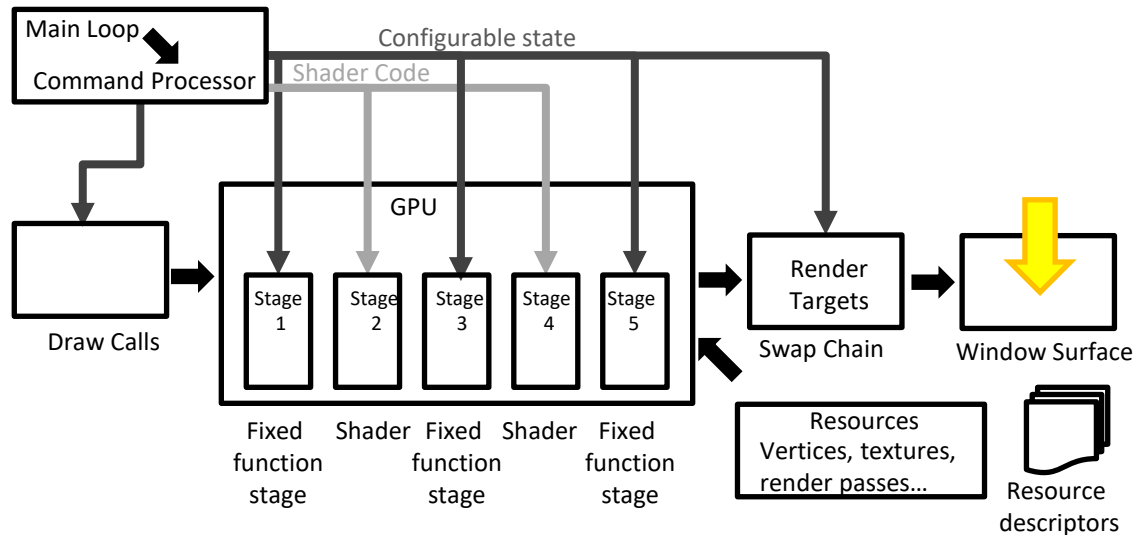


Initialization

- Initialize the graphics system
- Get the “command processor” object
 - Context (OpenGL) or command queue (Vulkan) – see later slides
- Select a render device
 - Which physical GPU
 - If there are multiple
 - Which features are needed
 - Number of viewports, single or double float...
 - Which operation modes are needed
 - Graphics, compute, memory transfer

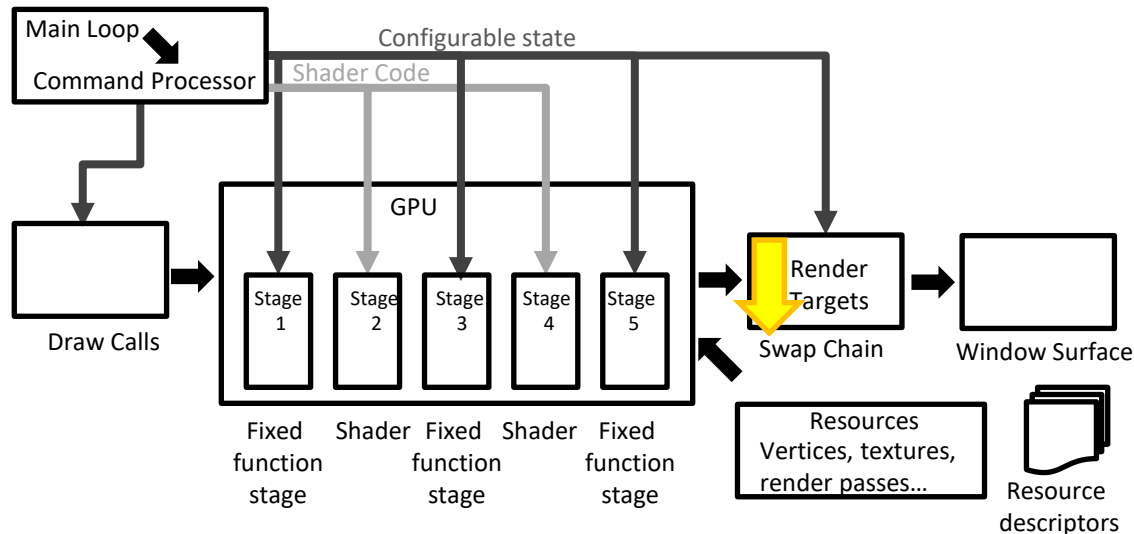
Surface

- Object connected to window or entire screen
- Requires a platform-specific window manager library (e.g., GLFW)



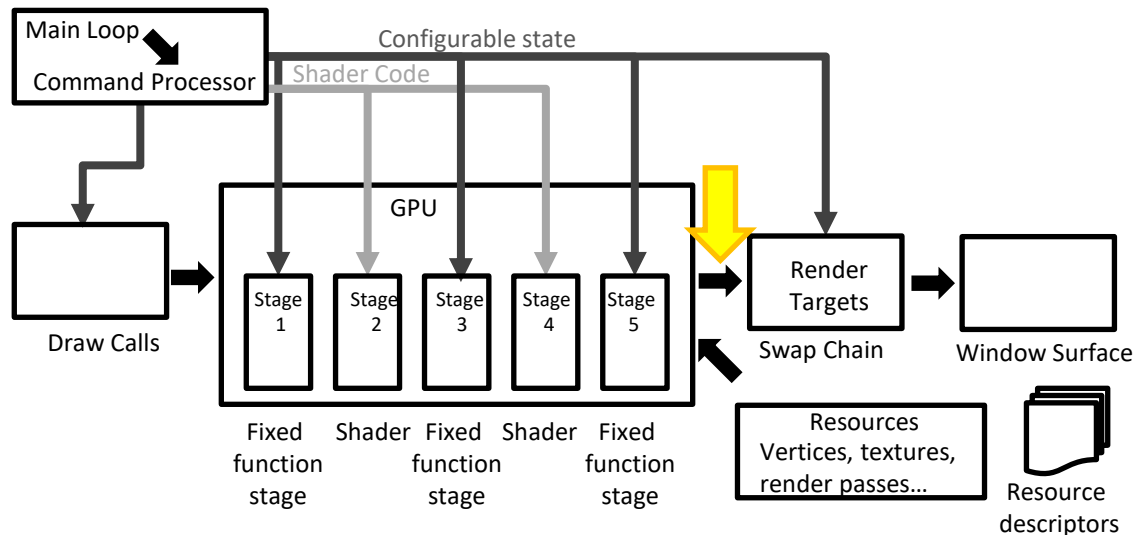
Swap Chain

- Contains render targets (images) waiting to be shown
- Single, double, triple buffering



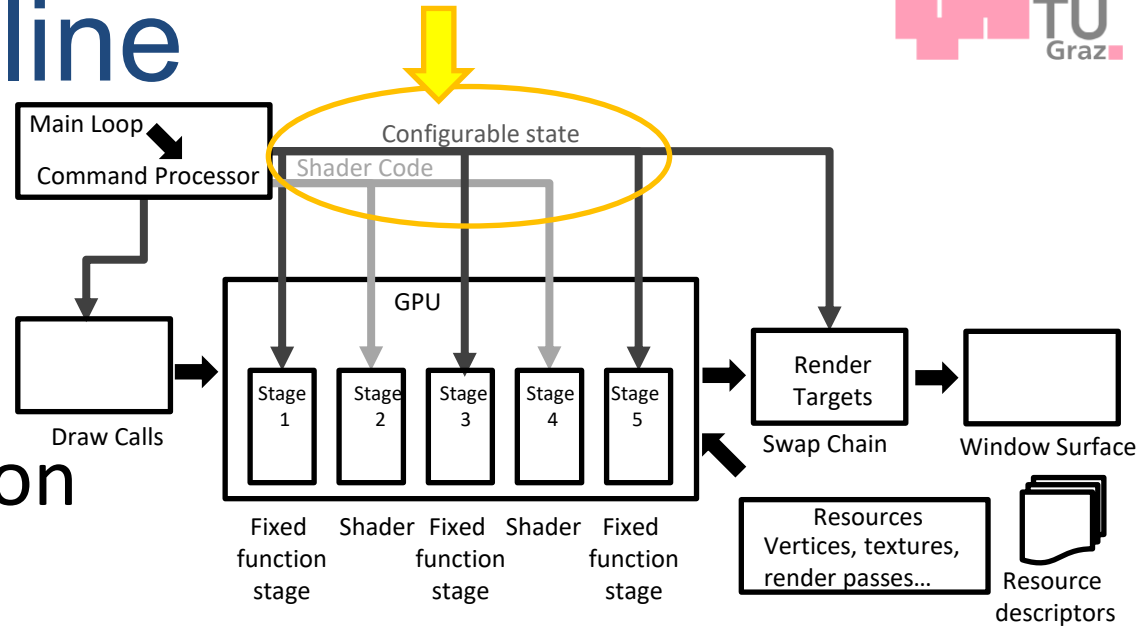
Framebuffer

- Points to render targets for color, depth, stencil



Graphics Pipeline

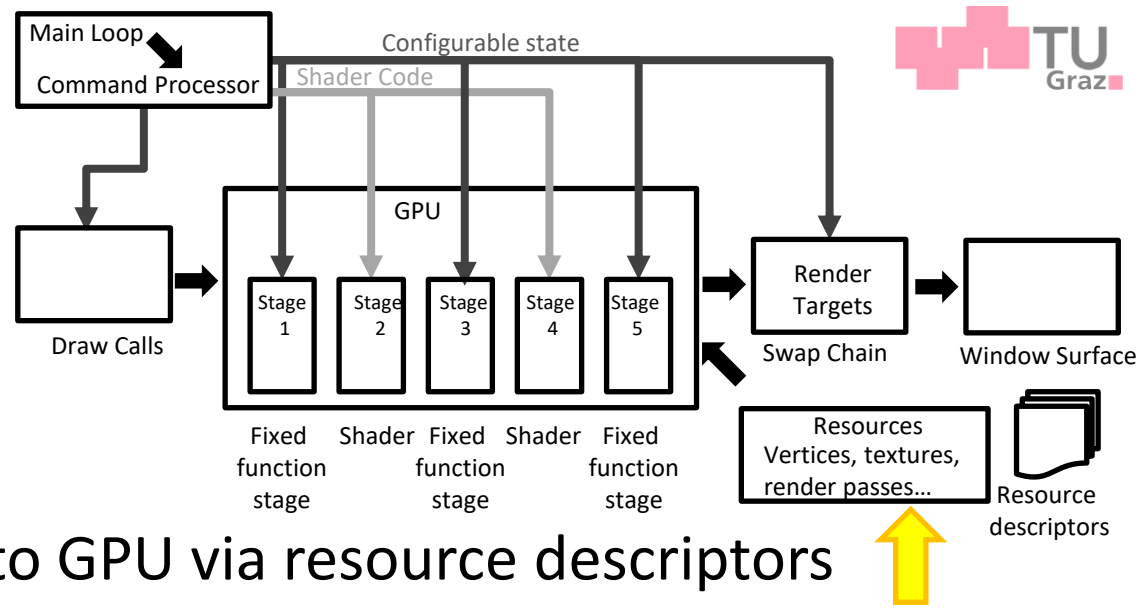
- Pipeline stores specific configuration of render state



- Configurable state: viewport size, depth buffer etc.
- Programmable state: shader programs

Resources

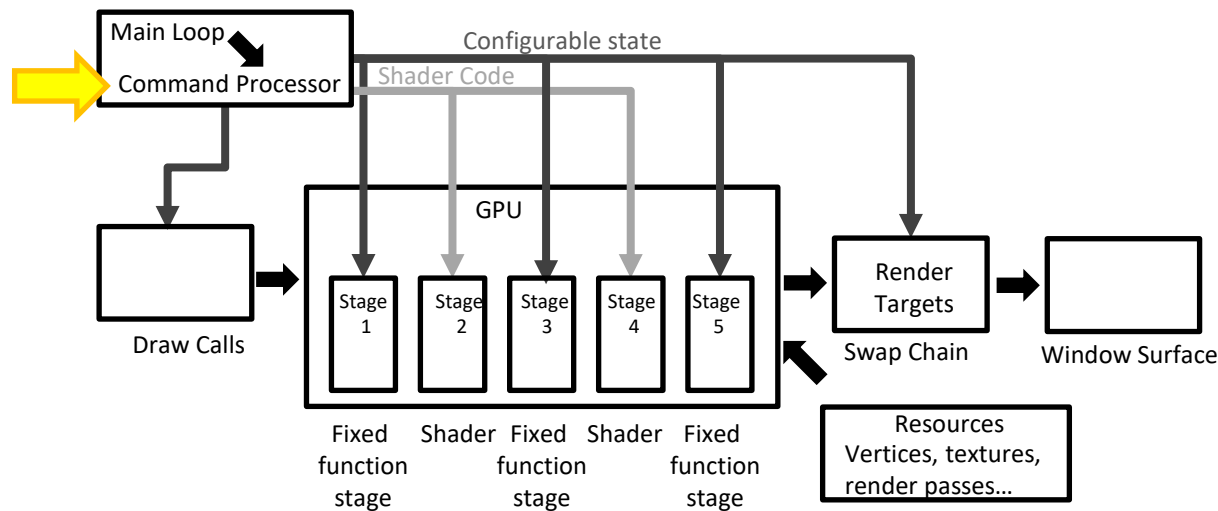
- Data stored in GPU local memory
- E.g, vertices, textures, other buffers
- Resource are described to GPU via resource descriptors
- Before use
 - Generate or download resource
 - Specify resource descriptors
 - Bind (=activate) the chosen resource
- Configurable state has “slots” for binding various resource types



Command Processor

Example command sequence


- Begin render pass
- Bind pipeline (Vulkan) or pipeline components (OpenGL)
- Bind framebuffer
- Draw vertices (=draw call)
- End render pass




Draw Calls

- Specify
 - Which primitive type
 - Which vertex buffer
 - Start index in buffer
 - Stop index in buffer



POINTS




LINES



LINE STRIP



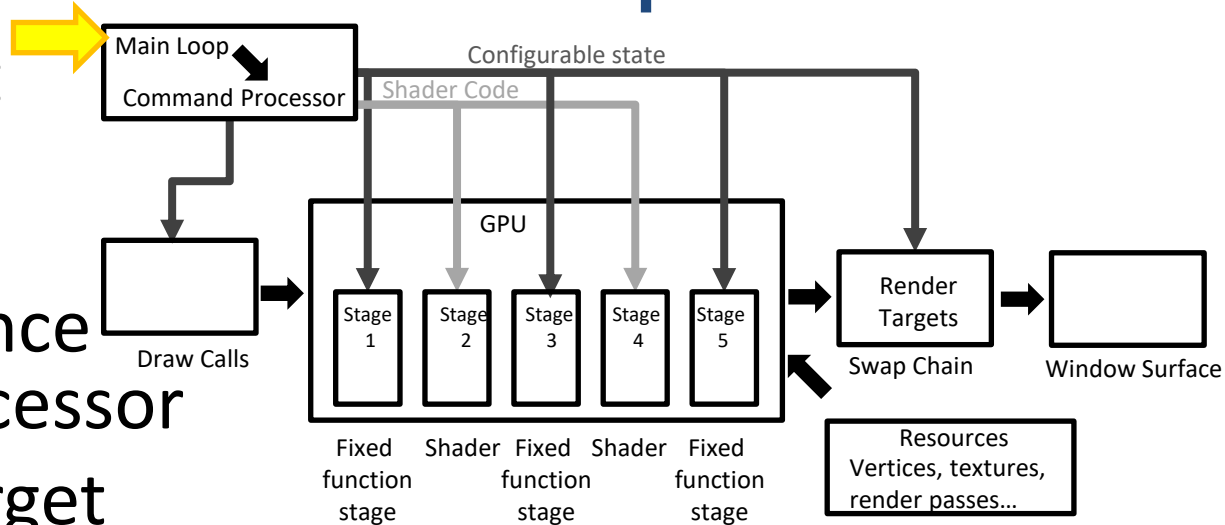
TRIANGLES



TRIANGLE STRIP

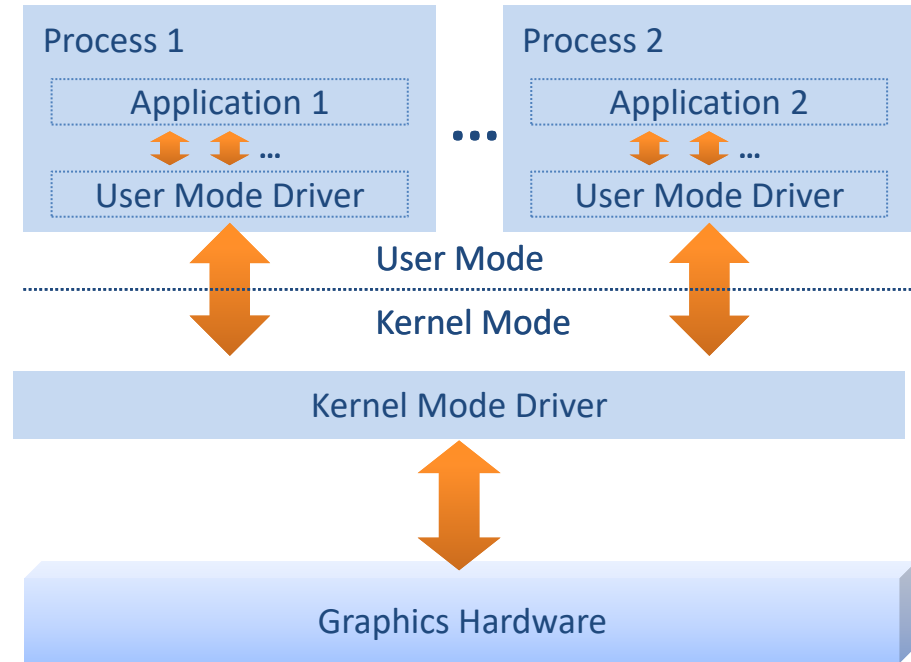
CPU Main Loop

- Get render target from swap chain
- Submit chosen command sequence to command processor
- Return render target to swap chain
- Command submission is very different between traditional and modern API



Graphics Driver Architecture

- User mode graphics driver
 - Minimize number of mode switches
 - Translation of graphics commands to instructions for the hardware
 - Batching, optimization, validation
 - Fine grained memory management
- Kernel mode graphics driver
 - Schedule access to hardware
 - Microkernel pattern, stability
 - Coarse grained memory management
 - Submits command buffers to GPU



Traditional API

- **Graphics context:** the system object in a traditional API
 - OpenGL, DirectX11 and below
 - Represents a virtual GPU
 - One process can have multiple contexts
 - Multiple contexts can share resources
- **Current context** for a given process
 - One to one mapping
 - Maximum of one current context per thread
 - Current context only assigned to one thread at the same time
 - All OpenGL operations work on current context

Pipeline State

- Inside a context, one must use commands to configure GPU **pipeline state** before rendering
- Pipeline state consists of
 - Programmable state (shaders)
 - Configurable state: blend, depth, culling, etc.
 - Layout: how to map settings at each stage's shader

Problems with Traditional API

- Pipeline execution is largely asynchronous
 - CPU sends commands into a “black hole” and
 - CPU does not know exactly when commands are executed
- Configuration of state can only be done incrementally
 - Submitting configuration changes to driver requires immediate validation, conversion, buffering → high cost at runtime
 - Drivers must have per-game optimizations built in
- Pipeline state *not* made explicit in rendering context
 - Switching pipeline configurations is cumbersome
 - Must be done by sequences of state change commands
- Pipeline abstraction is single-threaded on CPU
 - Cannot multi-thread feeding the pipeline
 - Having too many draw calls becomes a bottleneck

Modern API

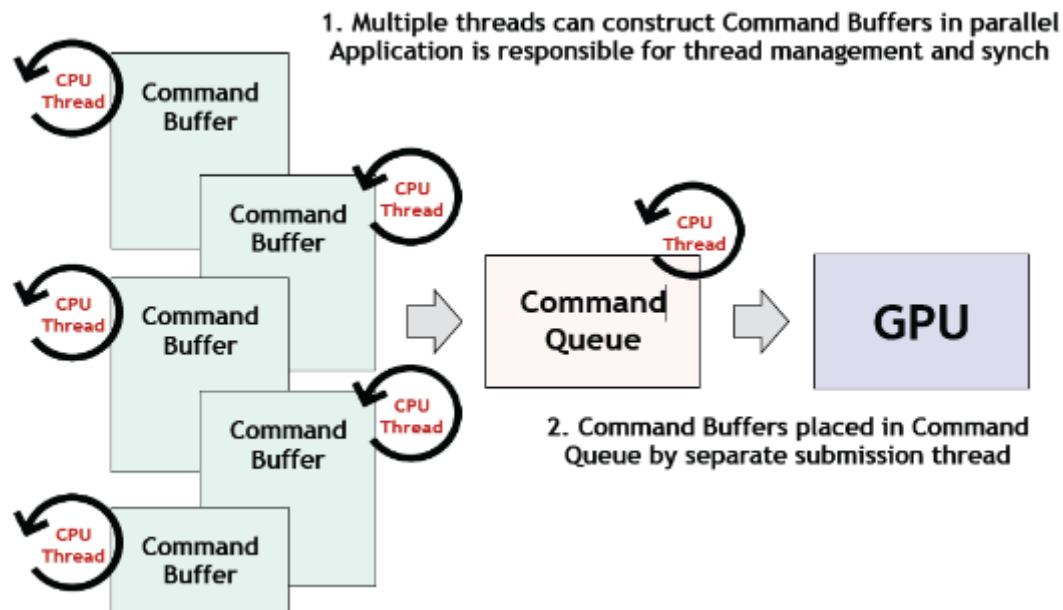
- DirectX 12, Vulkan (=OpenGL successor)
- Designed for modern GPU types (including mobile)
- Make (CPU driver side of) pipeline more programmable
- CPU multi-threading possible (at your own risk)
- Split rendering context into command buffers and queues
- Make pipeline state (and render passes) explicit
- Low overhead
- Fine-grained control (minimal program 1K lines of code)

Command Buffers

- Commands collected in command buffers
 - Optimize and validate command buffers during building, not during submission
 - Yields *immutable*, re-useable pipeline state configurations
 - Selected pipeline state variables can be declared *mutable*
- Command buffer submitted to *queue* for execution
- Build *many* command buffers from *many* threads
 - When the buffers are ready, one can submit them all at once
- Each command buffer just *switches* to its favorite pipeline
- Can use *synchronization* primitives across command buffers
 - Event, barrier, semaphore, fence


Queues

- Queues replace traditional contexts
- Insert command buffer into queue to schedule it



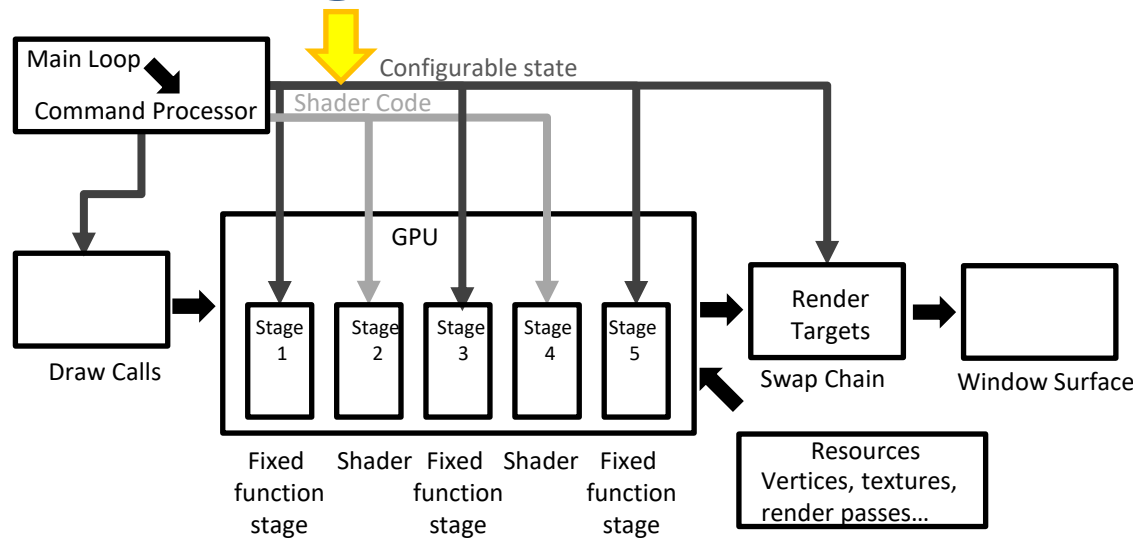
Example: Multi-Threaded Drawing

- ARM Vulkan Benchmark
 - ARM Cortex A-15/7, Mali T-628 MP6
 - 1000 meshes, 3 materials
 - 79% less CPU



Shader Programs

- For the programmable parts of the pipeline
- Compiled from shader language
 - HLSL, GLSL
- C-like syntax
- Stream execution model:
 - Shaders are like callbacks, no “main loop” code



Language Elements

- Skalar data types: float, int, bool...
- Vector/matrix data types: vec2, vec3, vec4, ivec3, mat4...
- Struct, arrays (static size)
- Texture sampler: sample2D, sampler3D, ...
- Control flow: if, else, while, for
- Function calls (no recursion)
- Swizzle operators: color1.rgb = color2.bga
- Mask operator: pos1.z = pos2.x + pos2.y

Build-In Functions

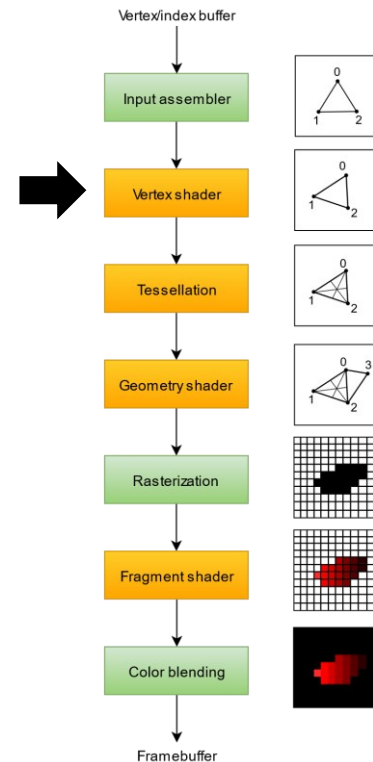
- Trigonometric
 - `sin()`, `cos()`, `radians()`, ...
- Logarithm, exponentiation
 - `log()`, `sqrt()`, ...
- Other
 - `min()`, `max()`, `mod()`, `floor()`, `abs()`, ...
- Geometric
 - `distance()`, `normalize()`, `dot()`, `length()`, ...
- Special functions
 - Linear interpolation
 - Reflection vector
 - Refraction vector

Shader Compilation

- Compilation separated into front-end/back-end
- Front-end: HLSL, GLSL, OpenCL, etc.
- Back-end = Bytecode
 - Microsoft HLSL uses proprietary bytecode
 - Vulkan uses SPIR-V (standard portable intermediate representation)
- Applications ship with bytecode, not shader source
- Just-in-time compilation of bytecode
 - To hardware target platform (NVIDIA, AMD, ...)

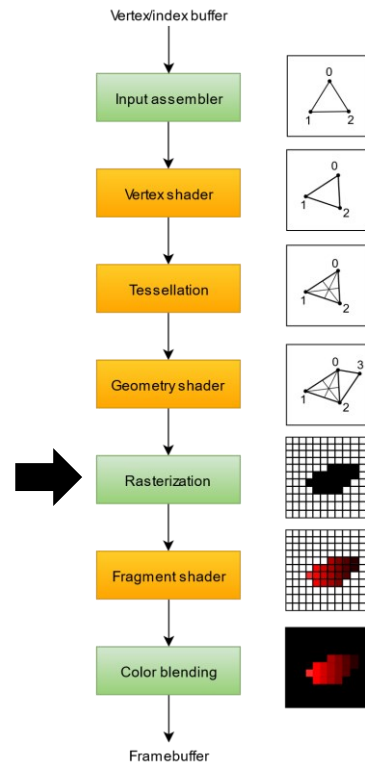
Vertex Shader

- Processes each vertex
- Input: vertex attributes
- Output: vertex attributes
- Mandatory output: `gl_Position`



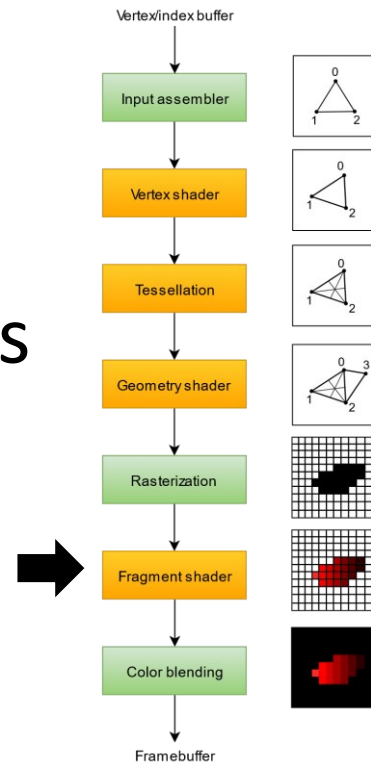
Rasterizer

- Generates fragments covering a primitive
- Fixed-function unit
- Input: primitives, vertex attributes
- Output: fragments with interpolated vertex attributes

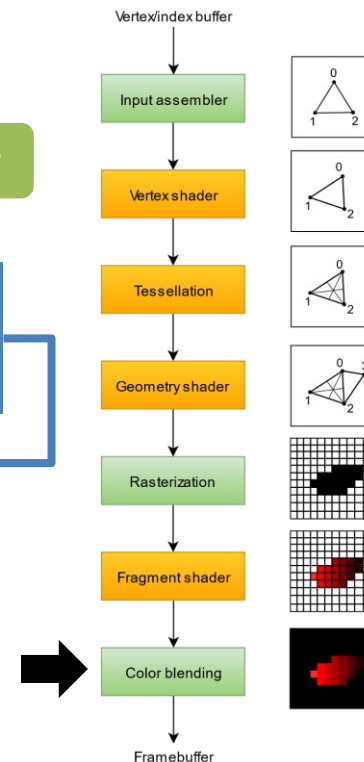
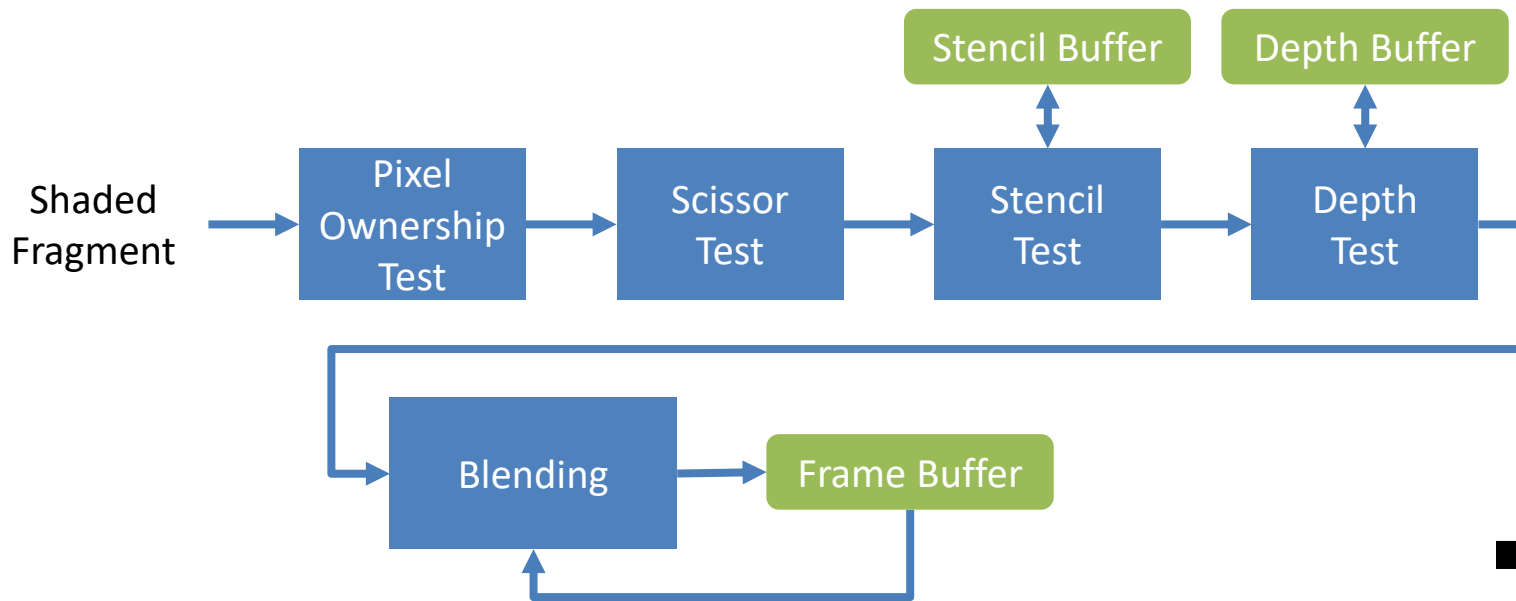


Fragment Shader

- Processes each fragment
- Input: interpolated vertex attributes
- Output: fragment color



Fragment Merging



Anatomy of a GLSL Shader

```
1  #version 330
2  // uniform inputs are constant for all shader invocations
3  uniform vec4 some_uniform;
4  // inputs are varying with each shader invocation
5  layout(location = 0) in vec3 some_input;
6  layout(location = 1) in vec4 another_input;
7  // outputs
8  out vec4 some_output;
9  void main()
10 {
11     //...
12 }
```

Built-In Variables

- Interface to fixed-function parts of pipeline
 - E. g. vertex shader:
 - `in int gl_VertexID;`
 - `out vec4 gl_Position;`
 - E. g. fragment shader:
 - `in vec4 gl_FragCoord;`
 - `out float gl_FragDepth;`

Example: Vertex Shader

```
1 #version 330
2 uniform mat4 mvp_matrix; // model-view-projection matrix
3 layout(location = 0) in vec3 vertex_position;
4 layout(location = 1) in vec4 vertex_color;
5 out vec4 color;
6 void main()
7 {
8     gl_Position = mvp_matrix * vec4(vertex_position, 1.0f);
9     color = vertex_color;
10 }
```

Example: Fragment Shader

```
1 #version 330
2 layout(location = 0) in vec4 color; // interpolated color
3 out vec4 fragment_color;
4 void main()
5 {
6     fragment_color = color;
7 }
```

Thank You!

Questions ?