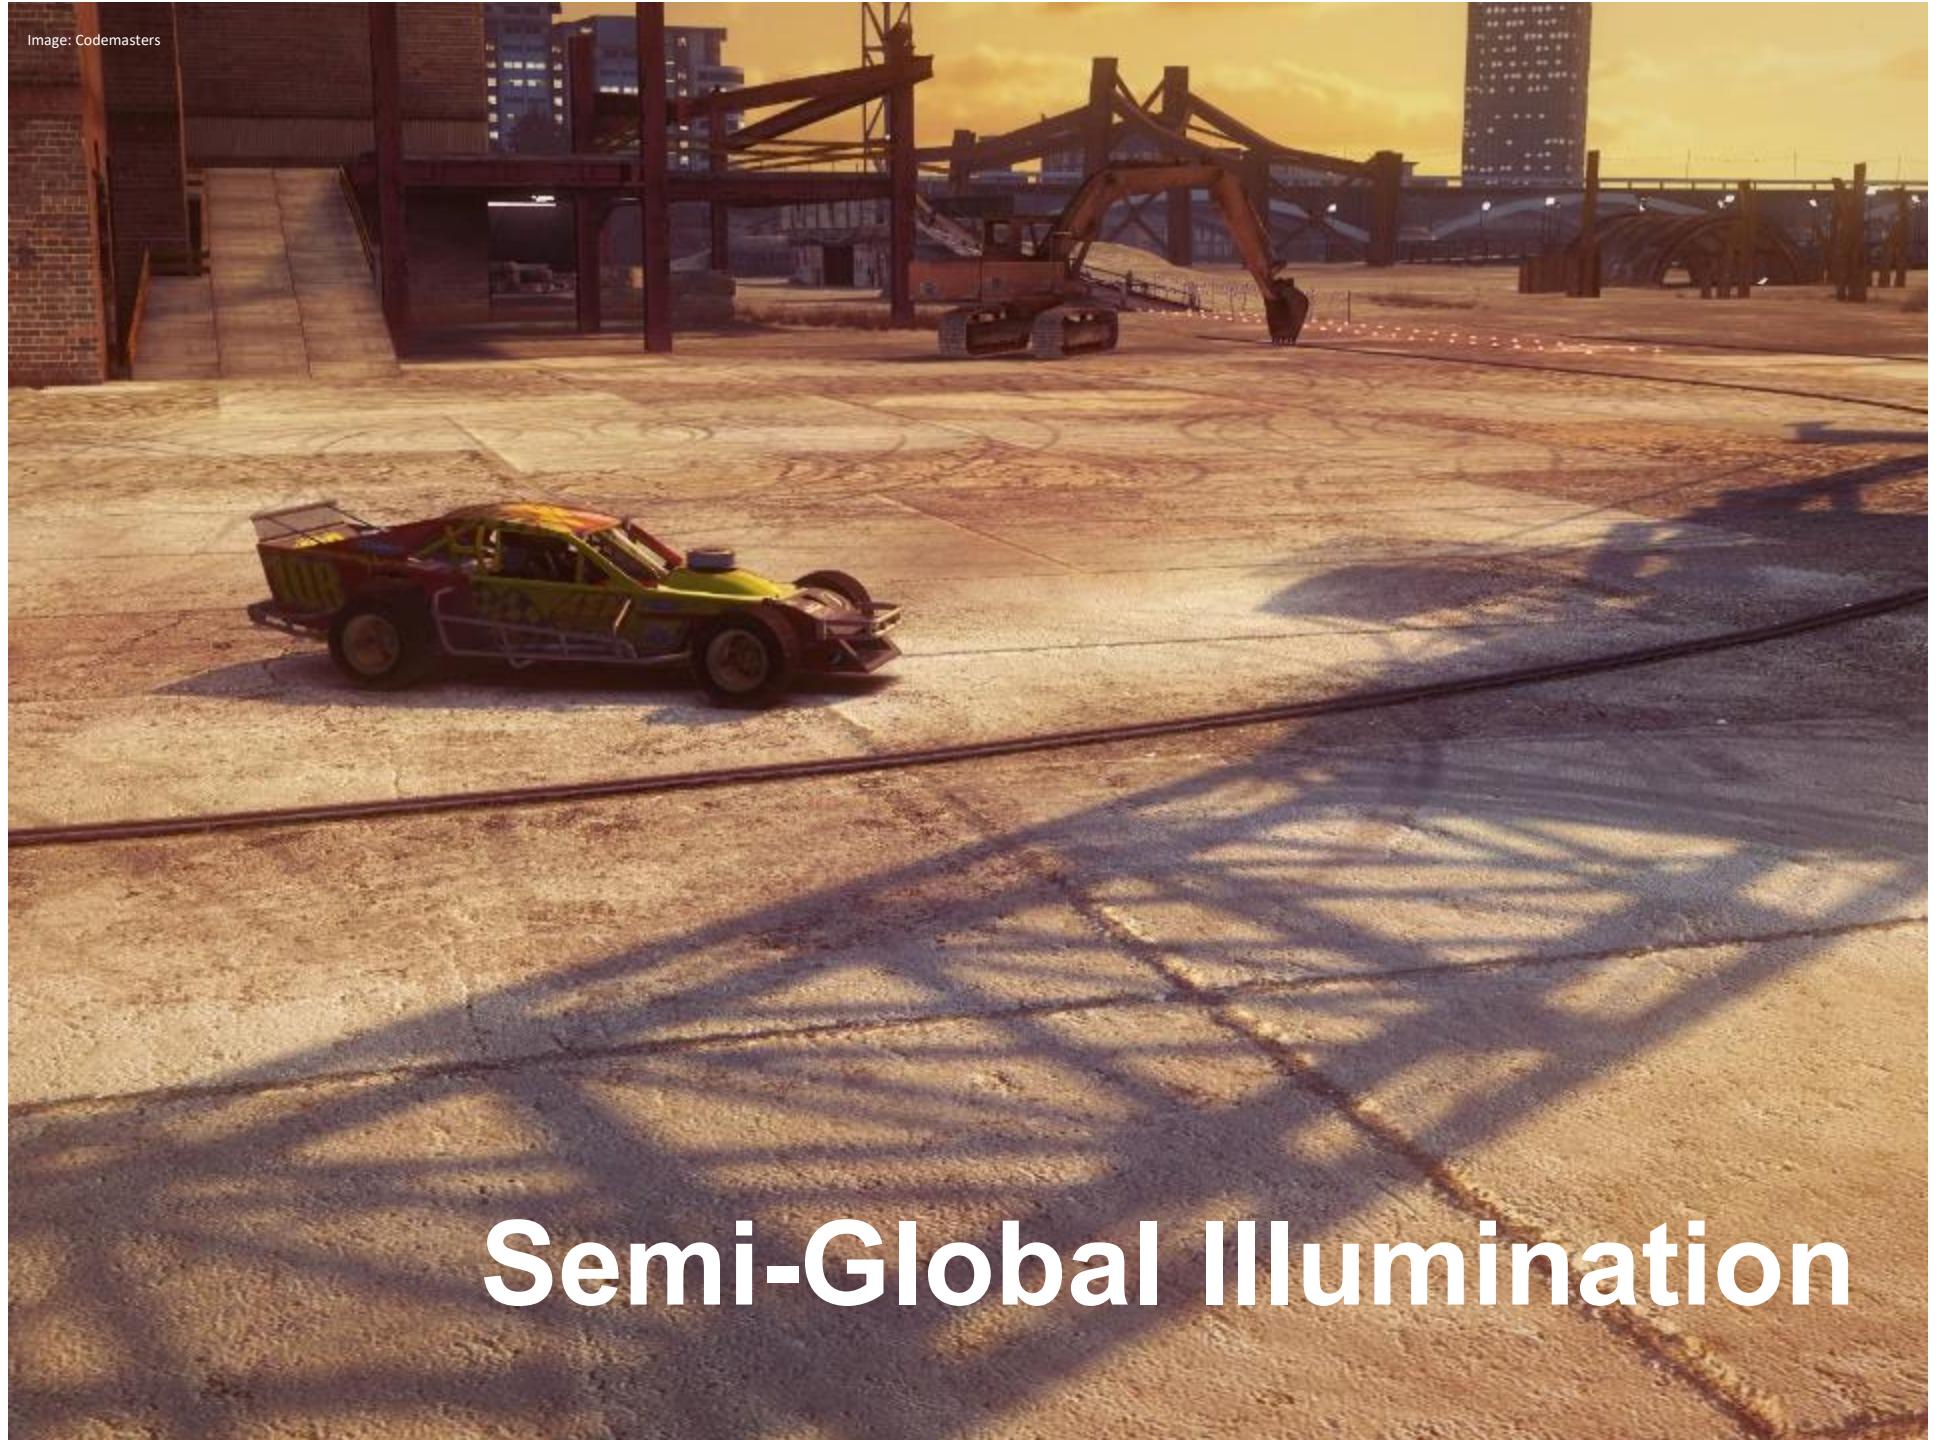


Image: Codemasters

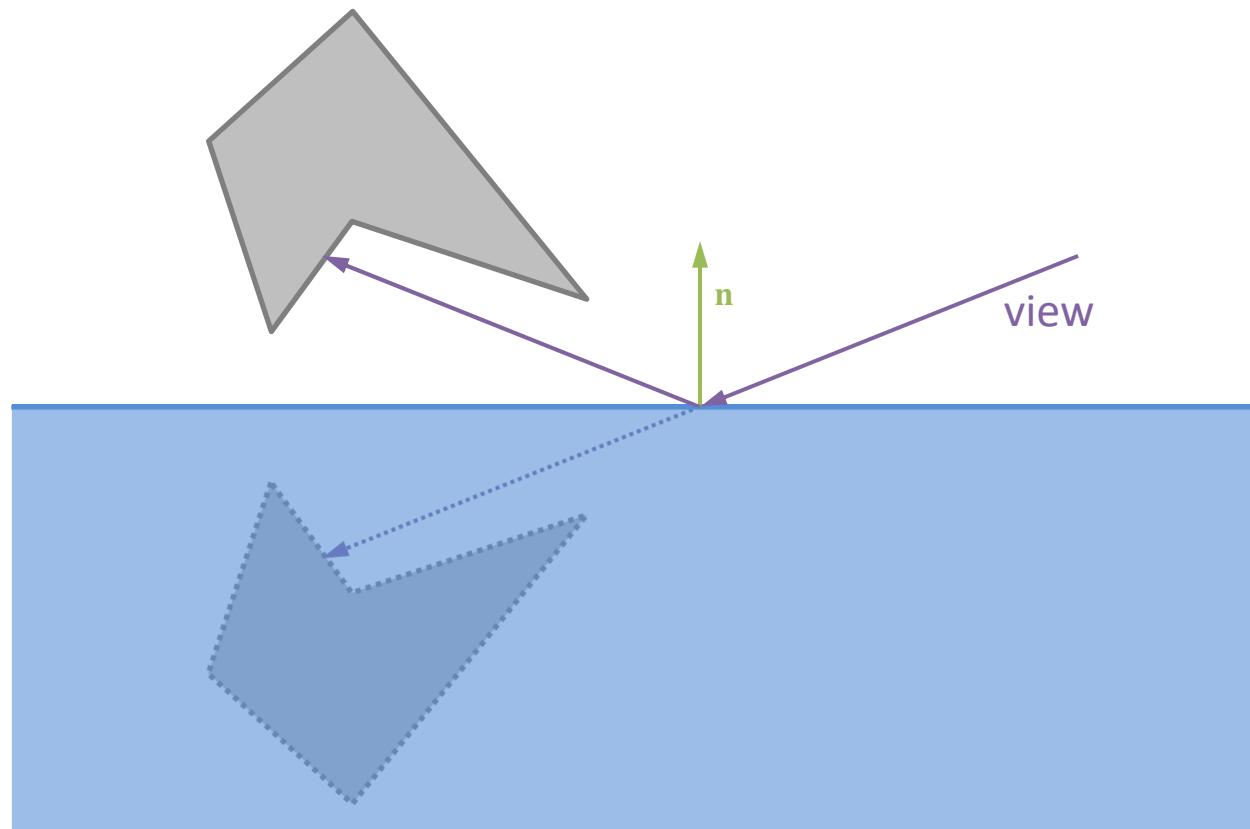


Semi-Global Illumination

Introduction

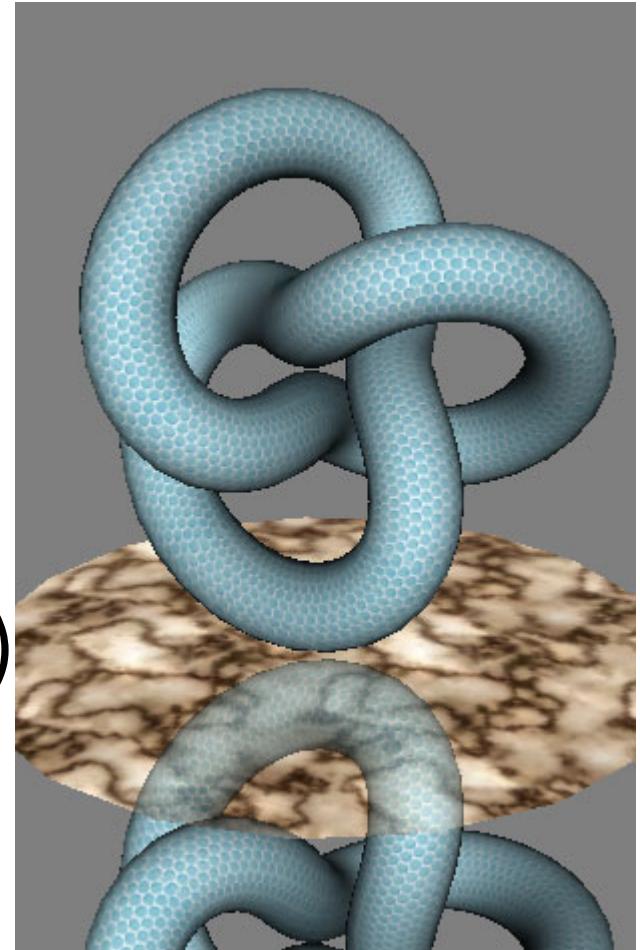
- Full global illumination is expensive
- Many subtle lighting effects may not be visible
- For real time rendering
 - Compute local illumination *everywhere*
 - Compute global illumination only *selectively*
 - Only certain types of light transport
 - Only for objects where it is visually important
- Important categories: reflections, transparency, shadows, ambient occlusion, precomputed radiance

Planar Reflections

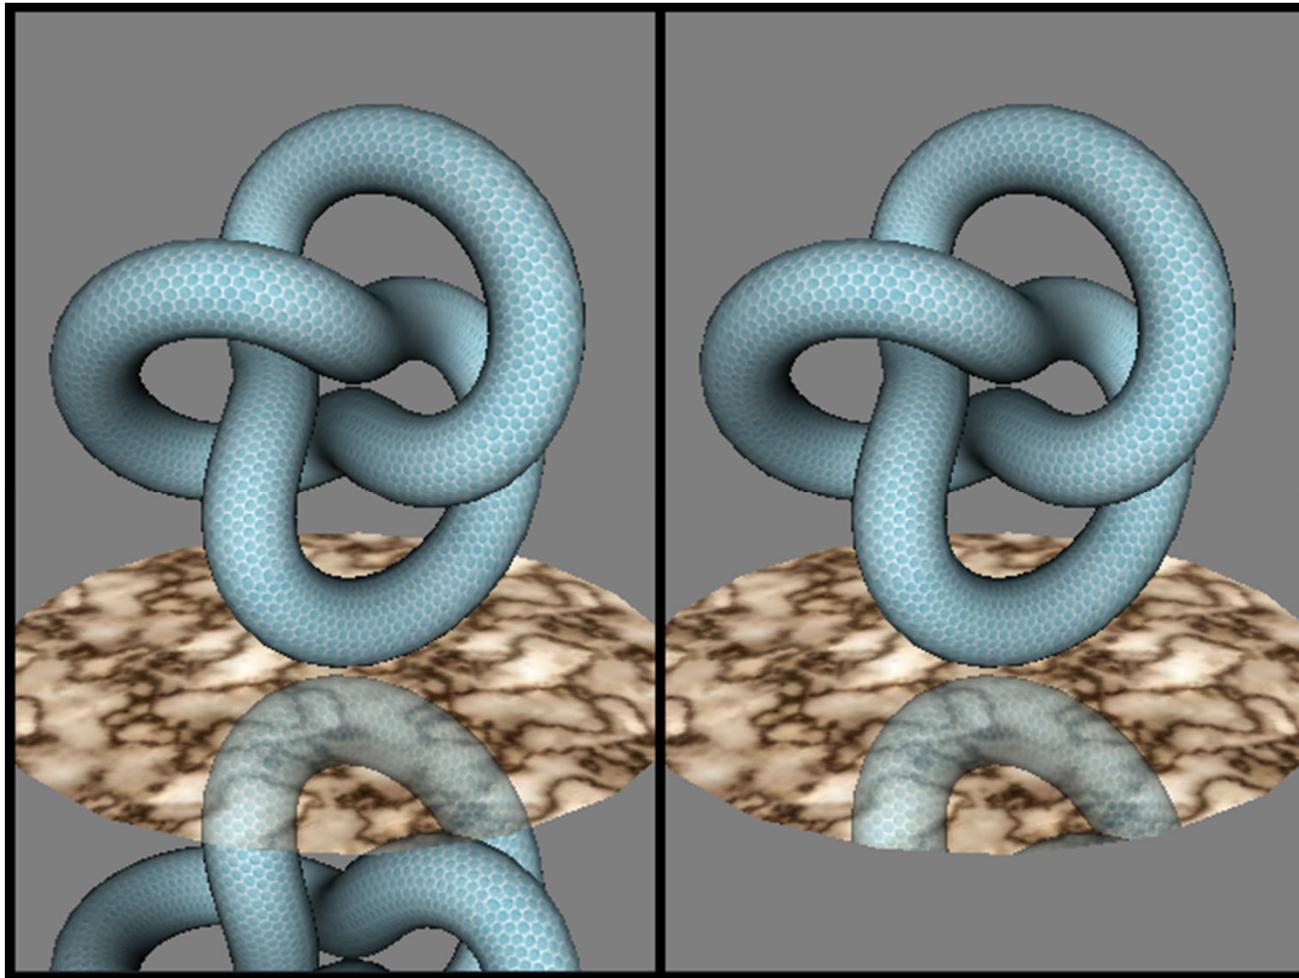


Planar Reflections in OpenGL

- Define a plane, e.g., z-plane
 - Scale(1, 1, -1)
- Adapt culling to mirroring
 - Cull front faces
 - Do *not* cull back faces
- Draw the object (mirrored)
- Draw the mirror (transparent)
- Draw the object (normal)



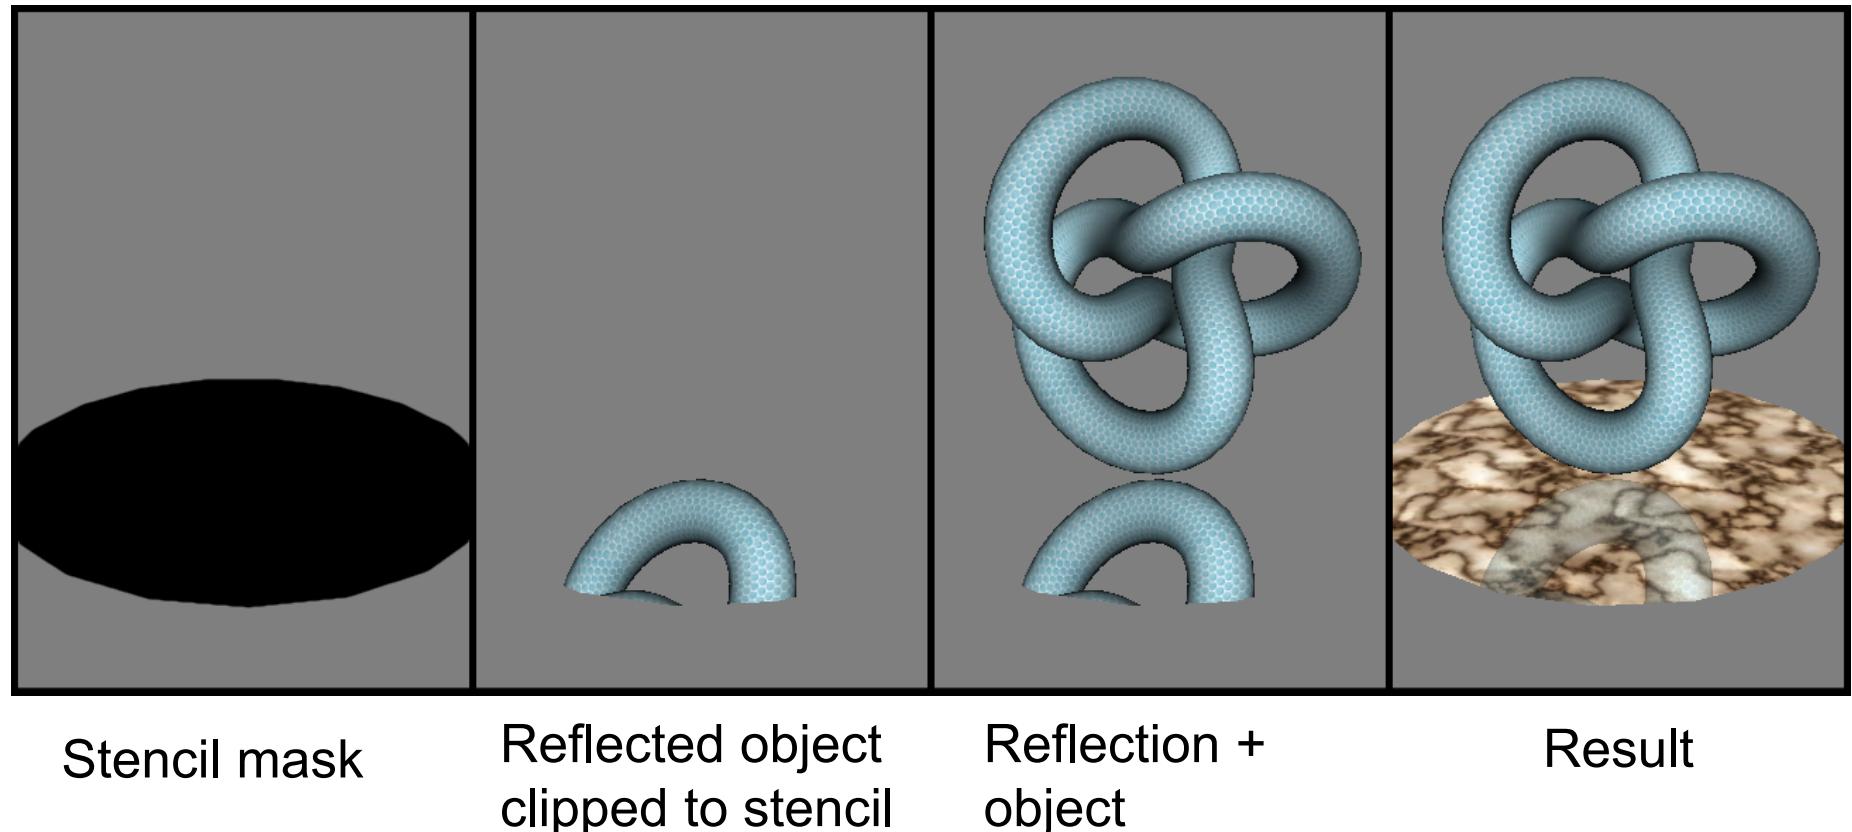
Problem



Object not restricted to mirror

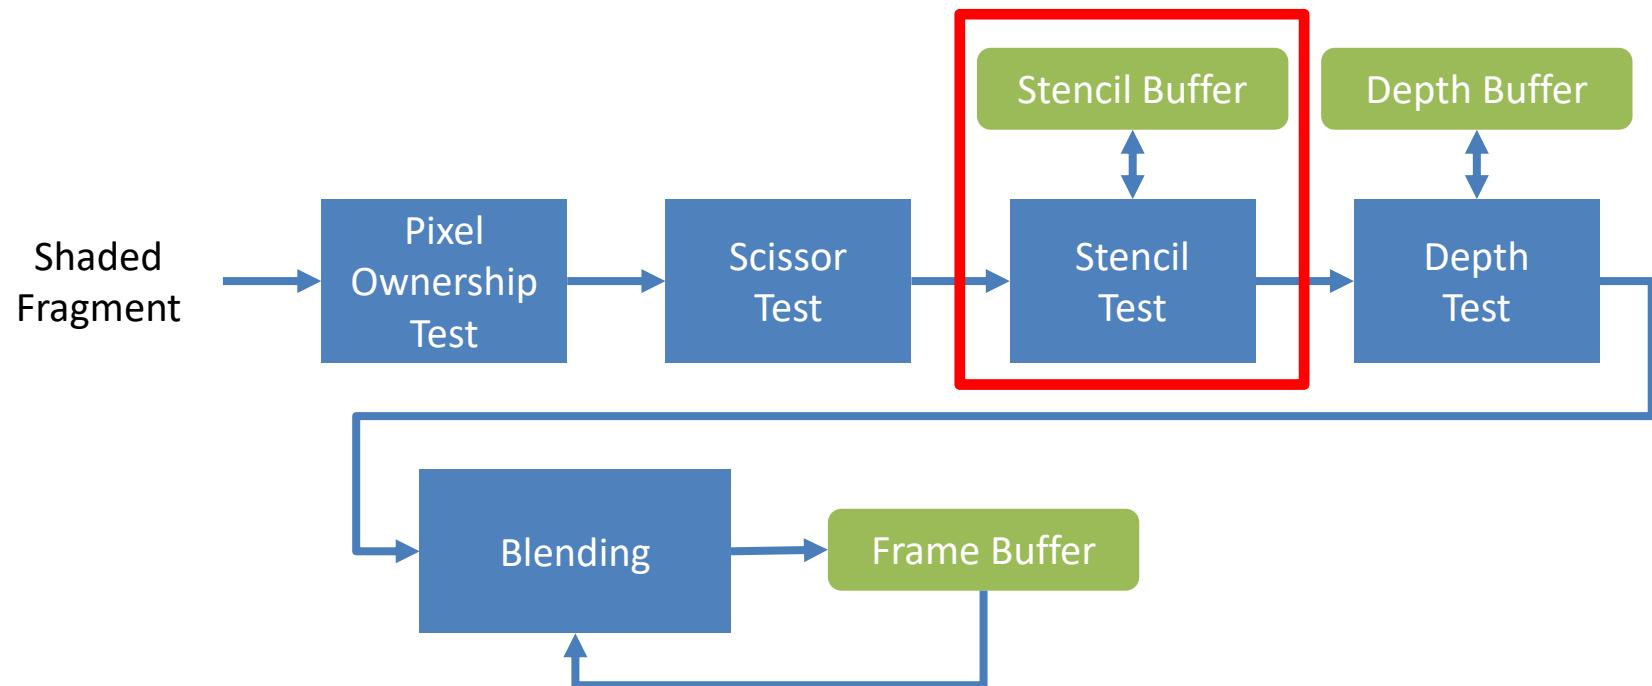
Solution: Stencil buffer

Reflection with Stencil Buffering



Fixed-Function Raster Operations

- When does the stencil test happen?



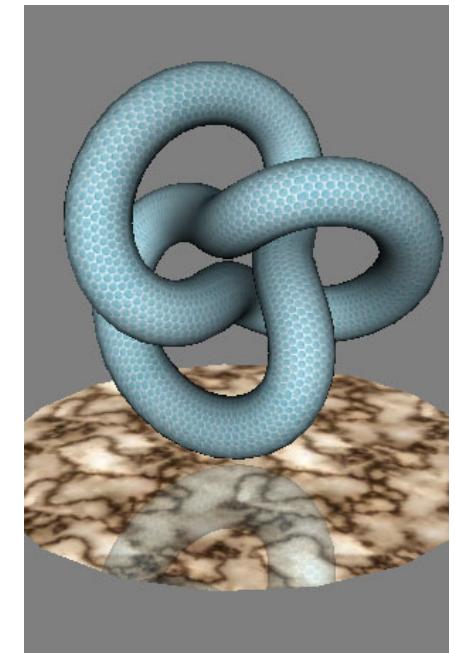
Stencil Buffer

- Allows to mask parts of the framebuffer
- 1 Bit Stencil Buffer ($\in \{0,1\}$)
- 8 Bit Stencil Buffer (256 states)
- Applications
 - Reflections
 - CSG (constructive solid geometry)
 - Shadow volumes
- Today:
 - Most applications do stenciling in shader code
 - Hardware stencil still very fast – example: Portal



Stencil Buffer Reflections

- Reset stencil buffer
 - Clear stencil buffer to zero
- Render mirror into stencil buffer
 - Enable stencil test
 - Set stencil operation to always write (on stencil-fail, depth-fail, depth-pass)
- Render mirrored object only in mirror
 - Set stencil operation to write if stencil bit set
- Render object normally



Advanced Reflections with Shaders

- Fade object color with increasing distance to reflector
- Fuzzy reflections: compute distorted texture coordinates



Caveats

- Reflection changes handedness
 - Adjust winding order of front faces when rendering reflected geometry
- If scene intersects mirror plane
 - Need to clip against plane to avoid artifacts



Image : Eric Lengyel

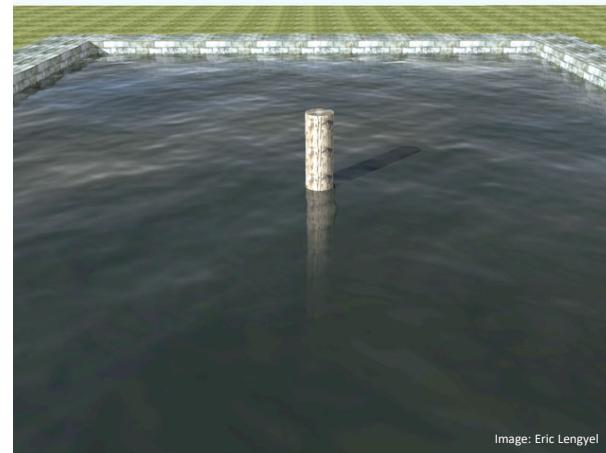
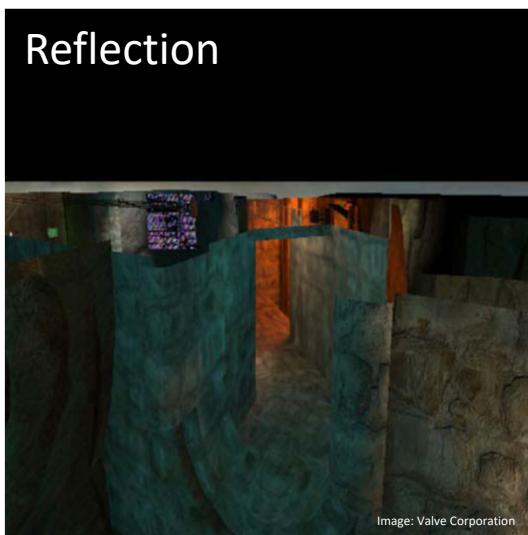


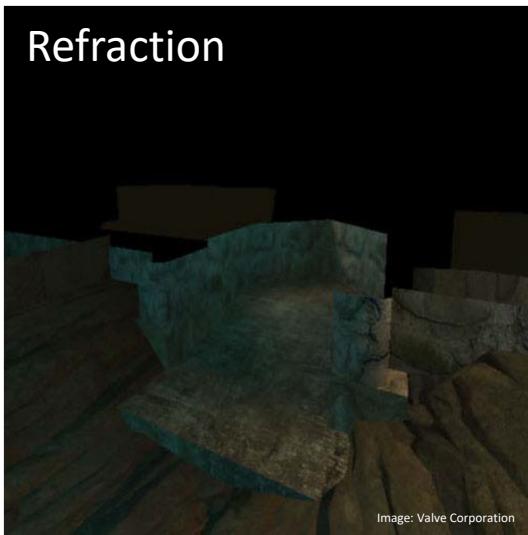
Image: Eric Lengyel

Shader-Based Approach

Reflection



Refraction



Dieter Schmalstieg

- Can also do refractions using shaders
- Render parts of the scene below the plane into additional texture
- Use fog to simulate scattering in water

Result



Semi-Global Illumination

Screen Space Reflections

Picture: Killing Floor 2 (Tripwire Interactive)



Screen Space Reflections (SSR)

- Fully dynamic raytraced reflections are very expensive (even with RTX)
- Idea:
 - Reflect view ray across normal
 - Find depth buffer intersection in screen space
- Result: cheap dynamic approximation of reflections



<https://lettier.github.io/3d-game-shaders-for-beginners/screen-space-reflection.html>

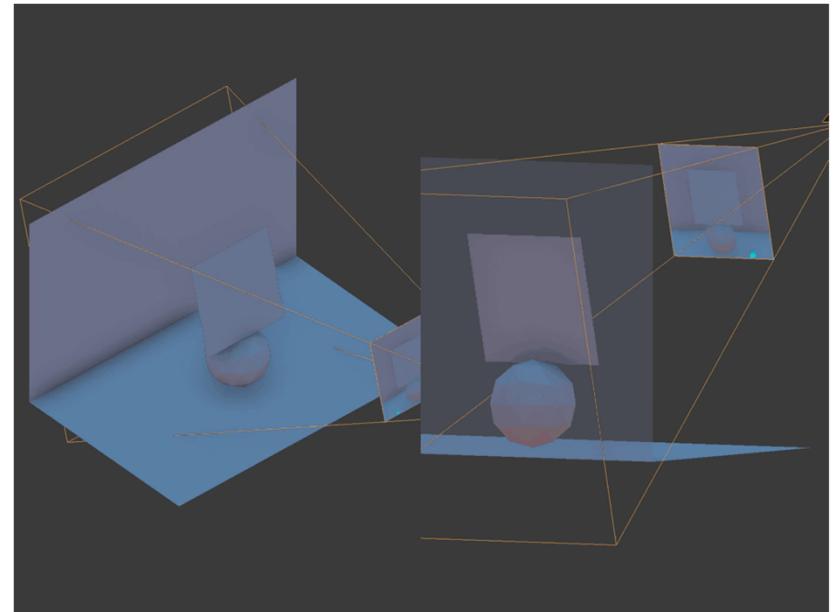
SSR: Reflection Map

- Compute a reflection map
 - Controls amount of reflection per fragment in screen space, depending on material
- Black pixels do not show SSR
- White pixels show full SSR



SSR: View Ray

- For each non-zero pixel in reflection map, reflect view ray across surface normal
- Along reflection vector
 - Choose a point
 - Project it
 - Compute reflection vector in screen space



<https://lettier.github.io/3d-game-shaders-for-beginners/screen-space-reflection.html>

SSR: Ray Marching

- From initial screen space position, **ray march** towards reflected screen space position
- If ray **intersects depth buffer**, we found our reflection point

<https://sakibsaikia.github.io/graphics/2016/12/26/Screen-Space-Reflection-in-Killing-Floor-2.html>



SSR: Ray Marching Code

- From the initial screen space position, we then **ray march** along towards the reflected screen space position
- If the ray **intersects the depth buffer**, we found our reflection point

```
#define MAX_REFLECTION_RAY_MARCH_STEP 0.02f
#define NUM_RAY_MARCH_SAMPLES 16

bool GetReflection(
    vec3 ScreenSpaceReflectionVec,
    vec3 ScreenSpacePos,
    out vec3 ReflectionColor)
{
    // Raymarch in the direction of the ScreenSpaceReflectionVec until
    // you get an intersection with your z buffer
    for (int RayStepIdx = 0; RayStepIdx<NUM_RAY_MARCH_SAMPLES; RayStepIdx++)
    {
        vec3 RaySample = (RayStepIdx * ReflectionRayMarchStep)
            * ScreenSpaceReflectionVec + ScreenSpacePos;

        float ZBufferVal = texture(DepthSampler, RaySample.xy).r;
        if (RaySample.z > ZBufferVal)
        {
            ReflectionColor = texture(SceneColorSampler, RaySample.xy).rgb;
            return true;
        }
    }
    return false;
}
```

<https://sakibsaikia.github.io/graphics/2016/12/26/Screen-Space-Reflection-in-Killing-Floor-2.html>

SSR: Intersection Code

- Fixed step size can lead to inaccurate results
- Use binary search between points just before and after depth intersection

```
#define NUM_BINARY_SEARCH_SAMPLES 6
// PrevRaySample: Sample just before the depth buffer intersection
// RaySample: Sample just after the depth buffer intersection
// DepthSampler: sampler2D to access the rendered depth texture

if (bFoundIntersection)
{
    vec4 MinRaySample = PrevRaySample;
    vec4 MaxRaySample = RaySample;
    vec4 MidRaySample;

    for (int i = 0; i < NUM_BINARY_SEARCH_SAMPLES; i++)
    {
        MidRaySample = mix(MinRaySample, MaxRaySample, 0.5);
        float ZBufferVal = texture(DepthSampler, MidRaySample.xy).r;

        if (MidRaySample.z > ZBufferVal)
            MaxRaySample = MidRaySample;
        else
            MinRaySample = MidRaySample;
    }
}
```

SSR: Limitations 1

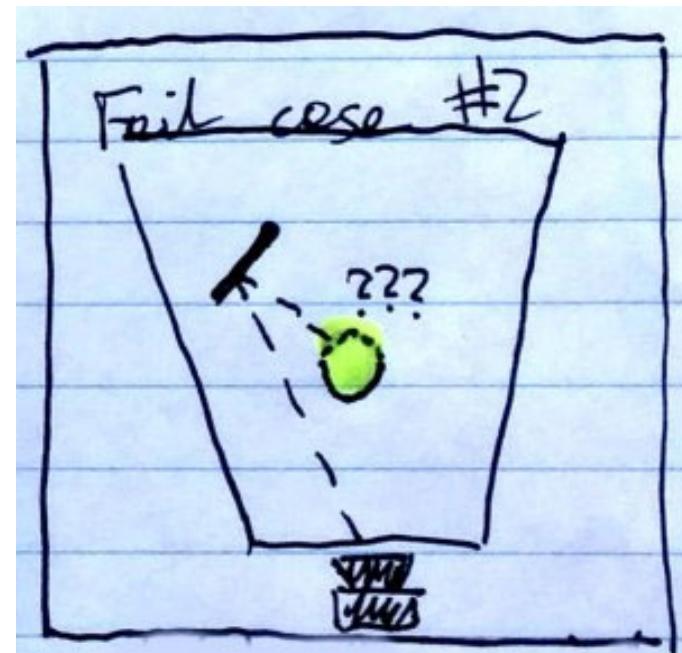
- Viewer facing reflections
- Reflections outside viewport



<https://bartwronski.com/2014/01/25/the-future-of-screenspace-reflections/>

SSR: Limitations 2

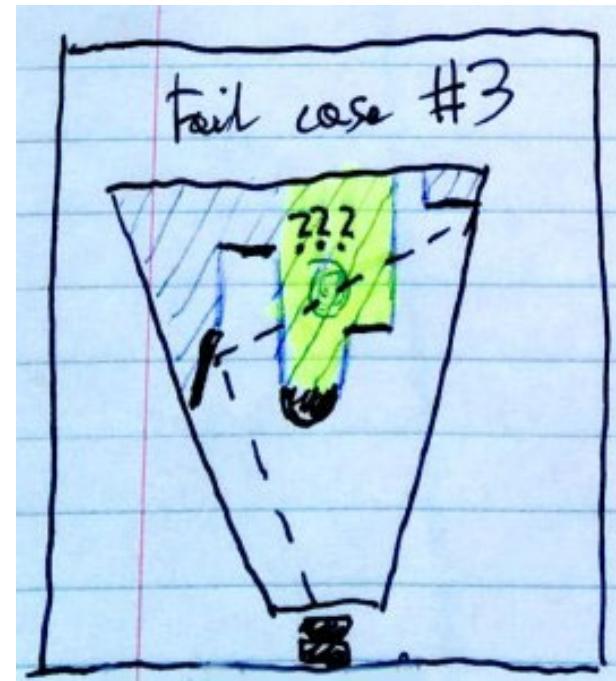
- Reflections of back faces



[https://bartwronski.com/2014/01/25/the
future-of-screenspace-reflections/](https://bartwronski.com/2014/01/25/the-future-of-screenspace-reflections/)

SSR: Limitations 3

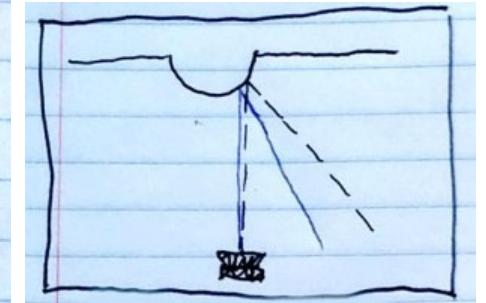
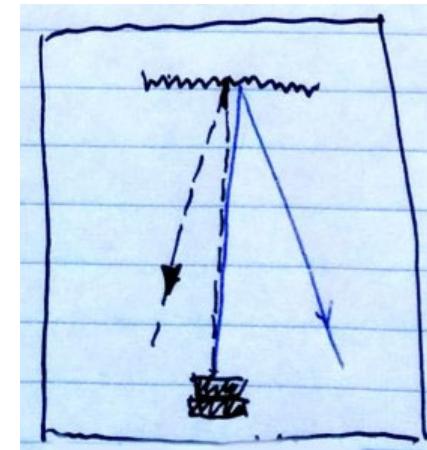
- Depth complexity



<https://bartwronski.com/2014/01/25/the-future-of-screenspace-reflections/>

SSR: Limitations 4

- Flickering
- Holes
- Temporal artifacts



<https://bartwronski.com/2014/01/25/the-future-of-screenspace-reflections/>

SSR: Further Improvements

- Blend with cubemaps or baked reflections
- Use temporal filtering to smoothen results
 - TAA is often exploited for this
- Use spatial filtering or blurring for rough surfaces
- Limit ray range in world space
 - For efficiency
 - To prevent far-away objects from flickering

Transparency

- Transparency blending is order dependent
- Blending must be back to front
- Need to sort scene elements by depth

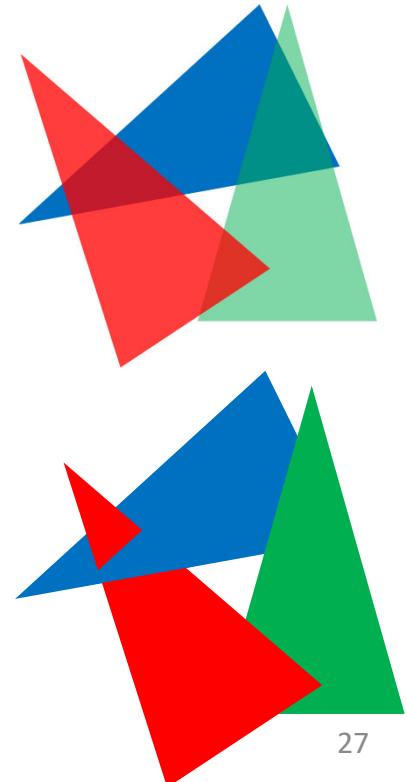


Methods for Sorting Fragments

- Explicit depth sorting
 - Write sorted order into draw buffer
- Depth peeling
 - Multi-pass rendering, once per *depth layer*
- Deep frame buffer
 - Write per-pixel linked lists
- Weighted blended order-independent rendering
 - Use depth as implicit weight (heuristic)

Explicit Depth Sorting

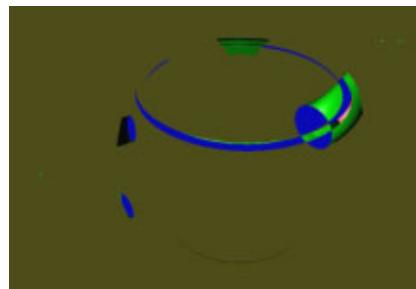
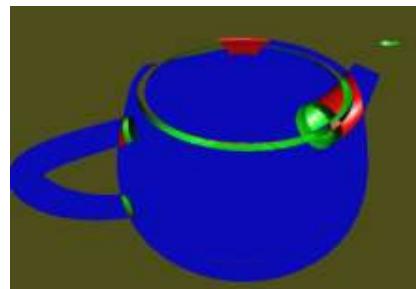
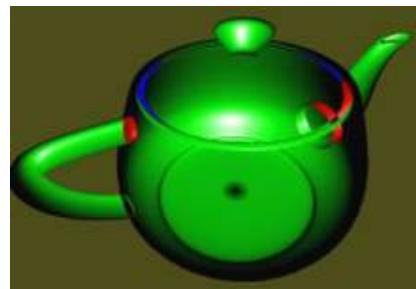
- Explicit depth sorting into draw buffer
- (+) Rendering with standard GPU primitive order
- (-) Writing sorted order to memory is expensive
- (-) Not always good enough
 - Cyclic overlap
 - Intersecting triangles
 - Can either split triangles (ugly, brittle), or
 - Can sort per fragments instead



Depth Peeling

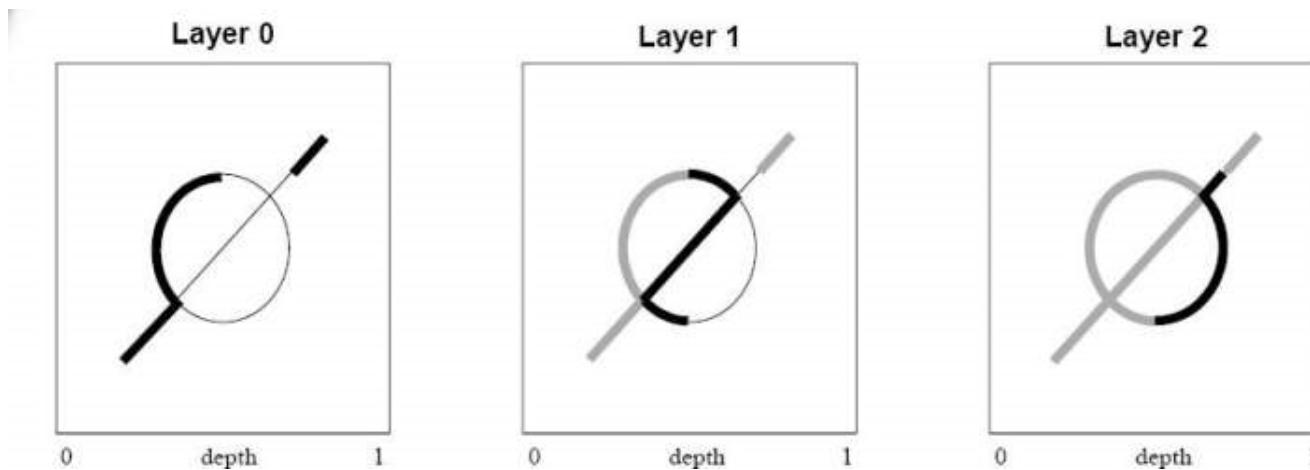
Interactive order independent transparency [Everitt'01]

- Multiple passes
 - First pass - find the front-most depth/color
 - Each successive pass - finds the depth/color for the next nearest fragment on a per-pixel basis
- Compare previous layer and current layer



Rendering Depth Peeling

- Render the scene n times
- Save the results in texture using render-to-texture/FBO
- Each pass forms a layer
- Starting from the second pass, use a shader for comparison with previous layer
- Compose n layers

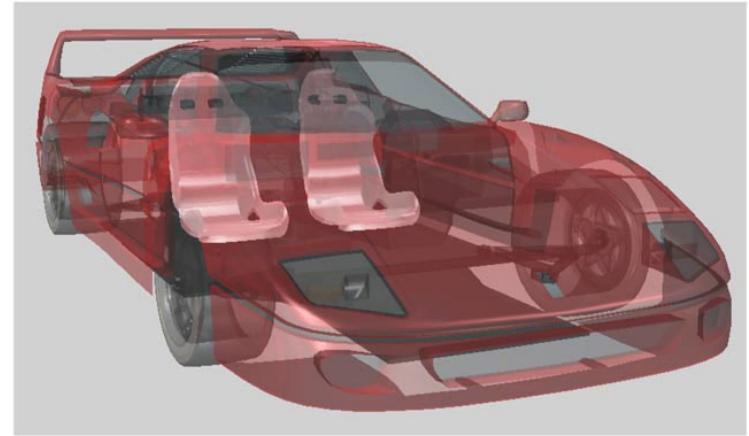


Dual Depth Peeling

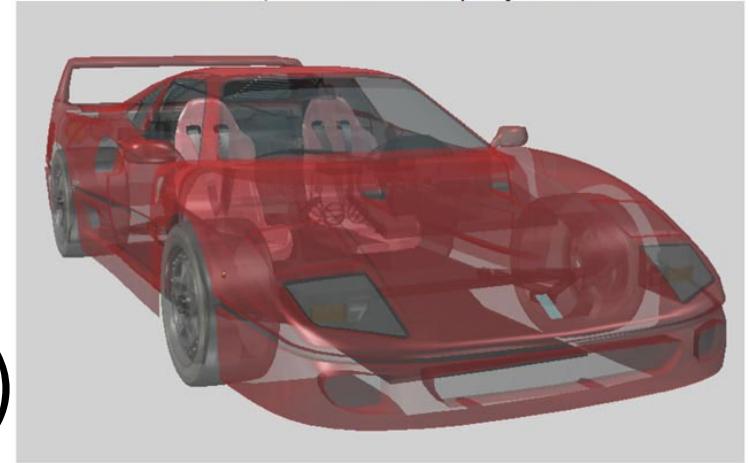
- Peel front and back in one pass
- Keep track of min/max depth in two textures
- Compare current depth with min/max values
 - If $<$ min write to front layer
 - If $>$ max write to back layer
- Use previous front/back layer depth values in current pass
- Decreases number of passes from n to $n/2+1$

Deep Framebuffer

- Depth peeling creates too much geometry overhead
 - $O(\text{triangles} * \text{overdraw})$
- Use pixel operations instead
 - Write every fragment to a per-pixel linked list
- Sort and blend the lists
- Historically called *accumulation buffer (A-buffer)*



Without OIT, note the incorrect depth of the seats



With OIT applied

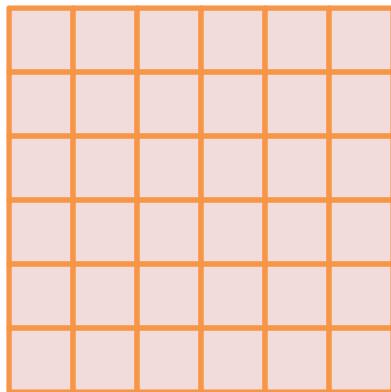
Per-Pixel Linked Lists

- Algorithm
 - Render scene and generate linked list per pixel
 - Sort list and compose fragments
- Requires read/write buffer
- Need two additional buffers
 - Fragment and link buffer
 - Start offset buffer

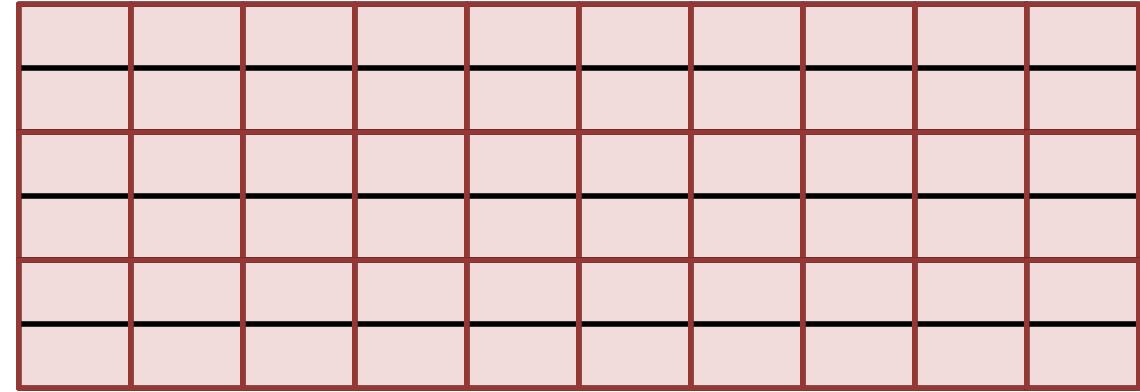
Fragment and Link Buffer

- Contains all fragment data produced during rasterization
- Must be large enough to store all fragments

Viewport



Fragment and Link Buffer



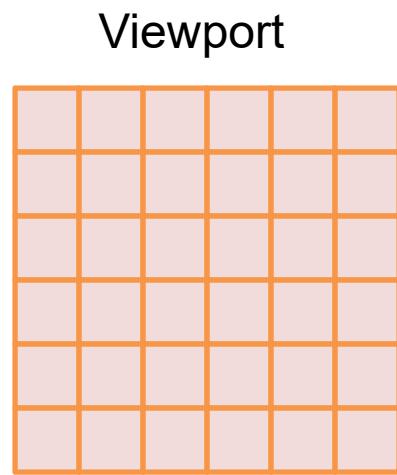
Start Offset Buffer

- Contains the offset of the last fragment written at every pixel location
- Screen-sized: (`width * height * sizeof(UINT32)`)
- Initialized to magic value
 - Magic value indicates end of list
(no more fragments stored)

-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1

Start Offset Buffer

Linked List Creation 1



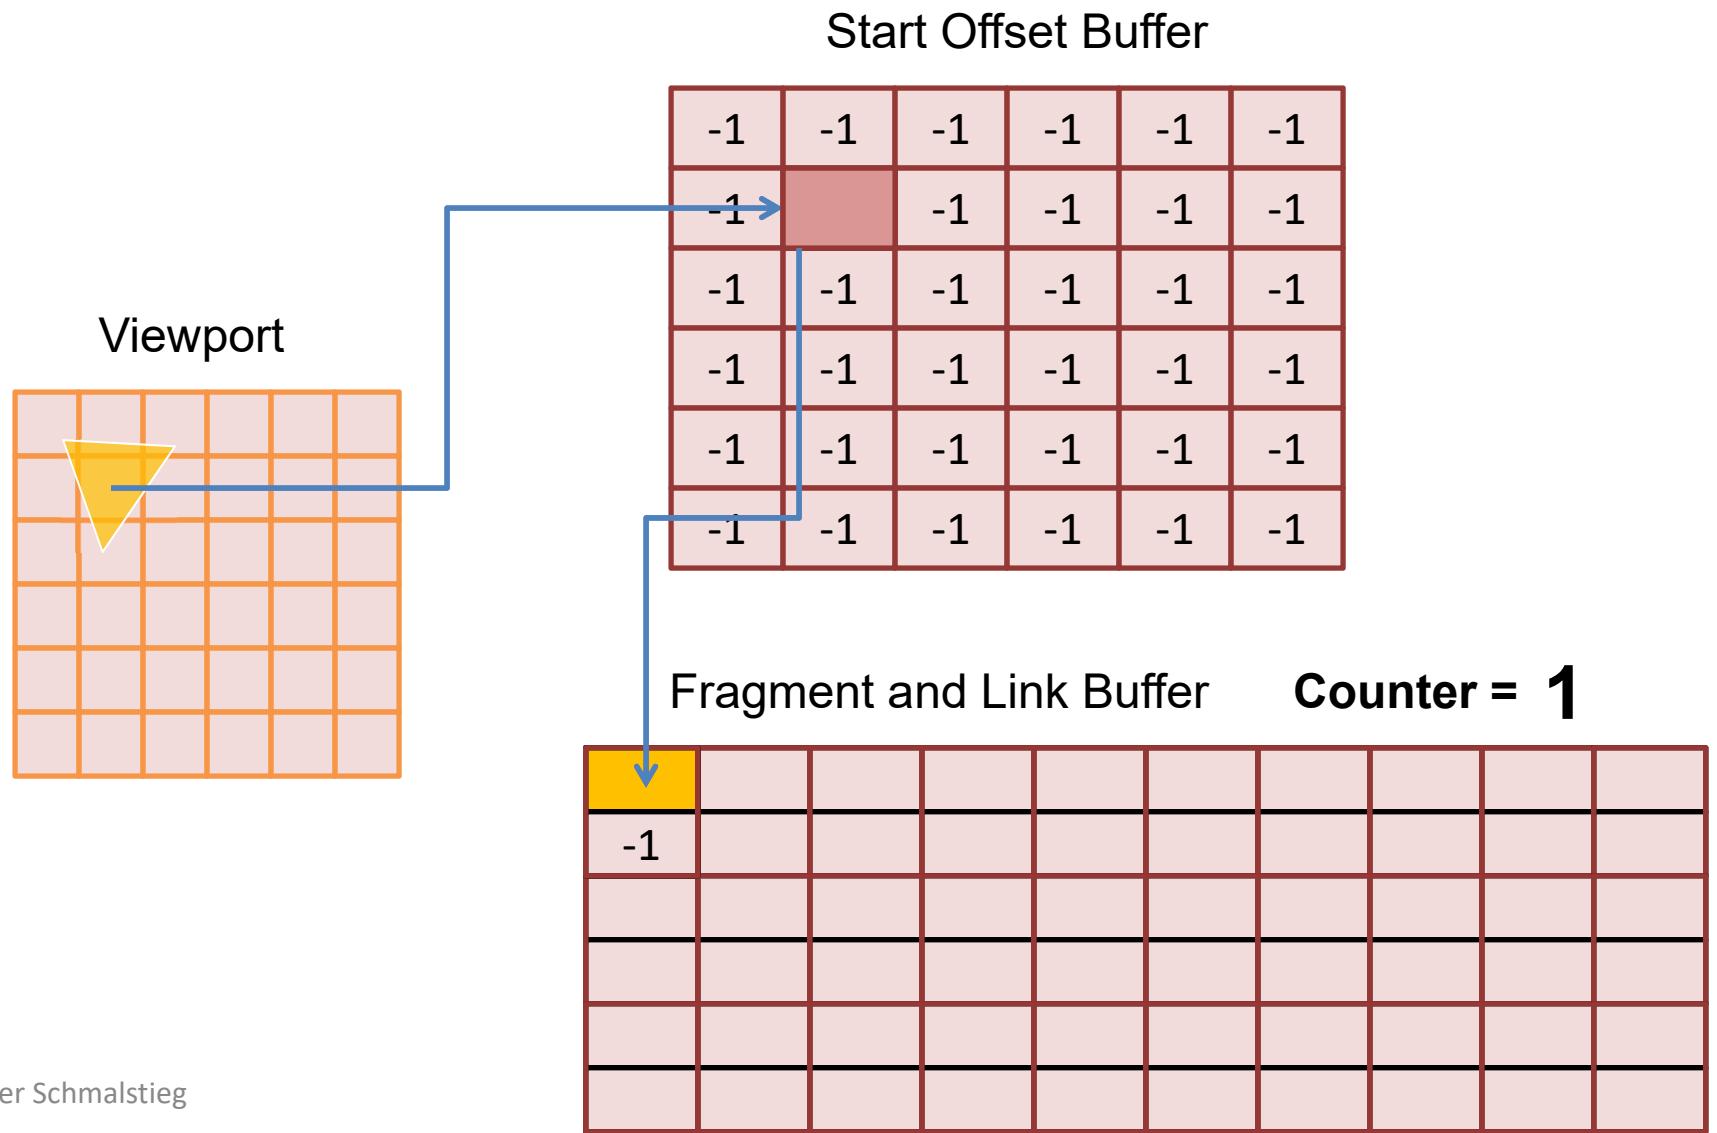
Start Offset Buffer

-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1

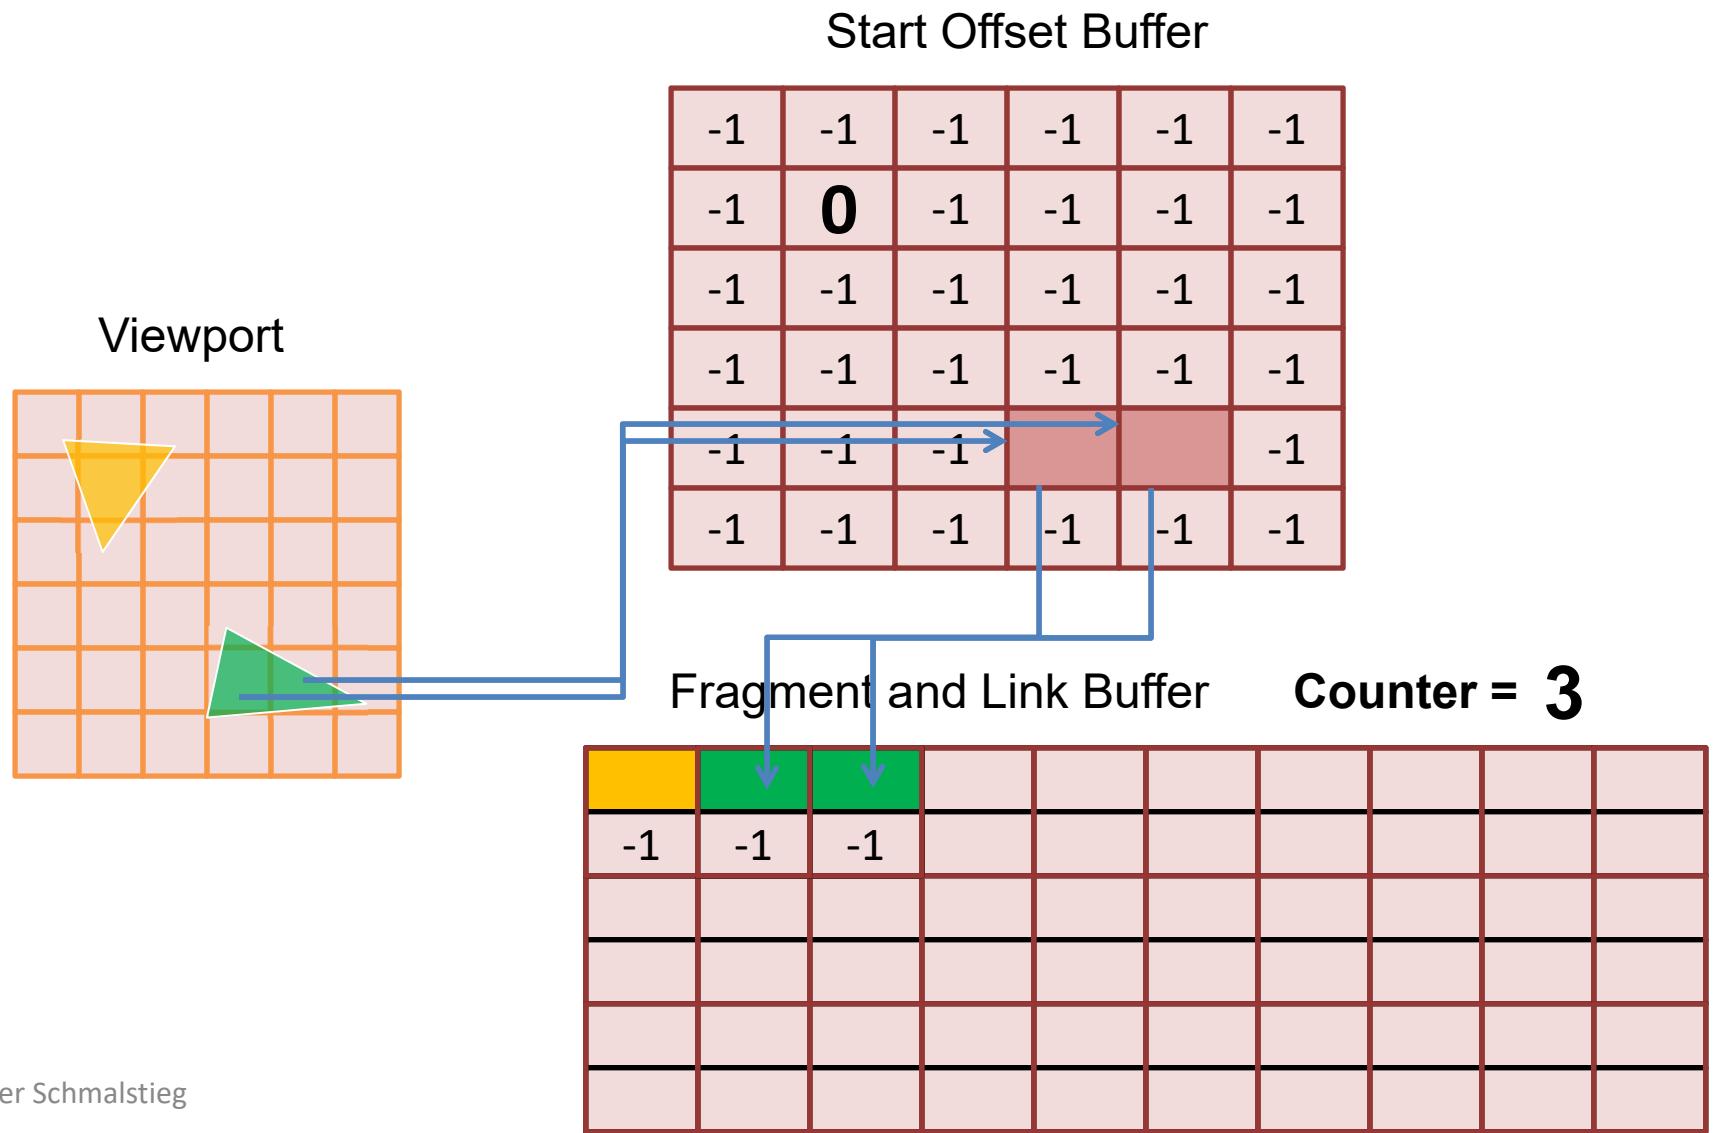
Fragment and Link Buffer

Counter = 1

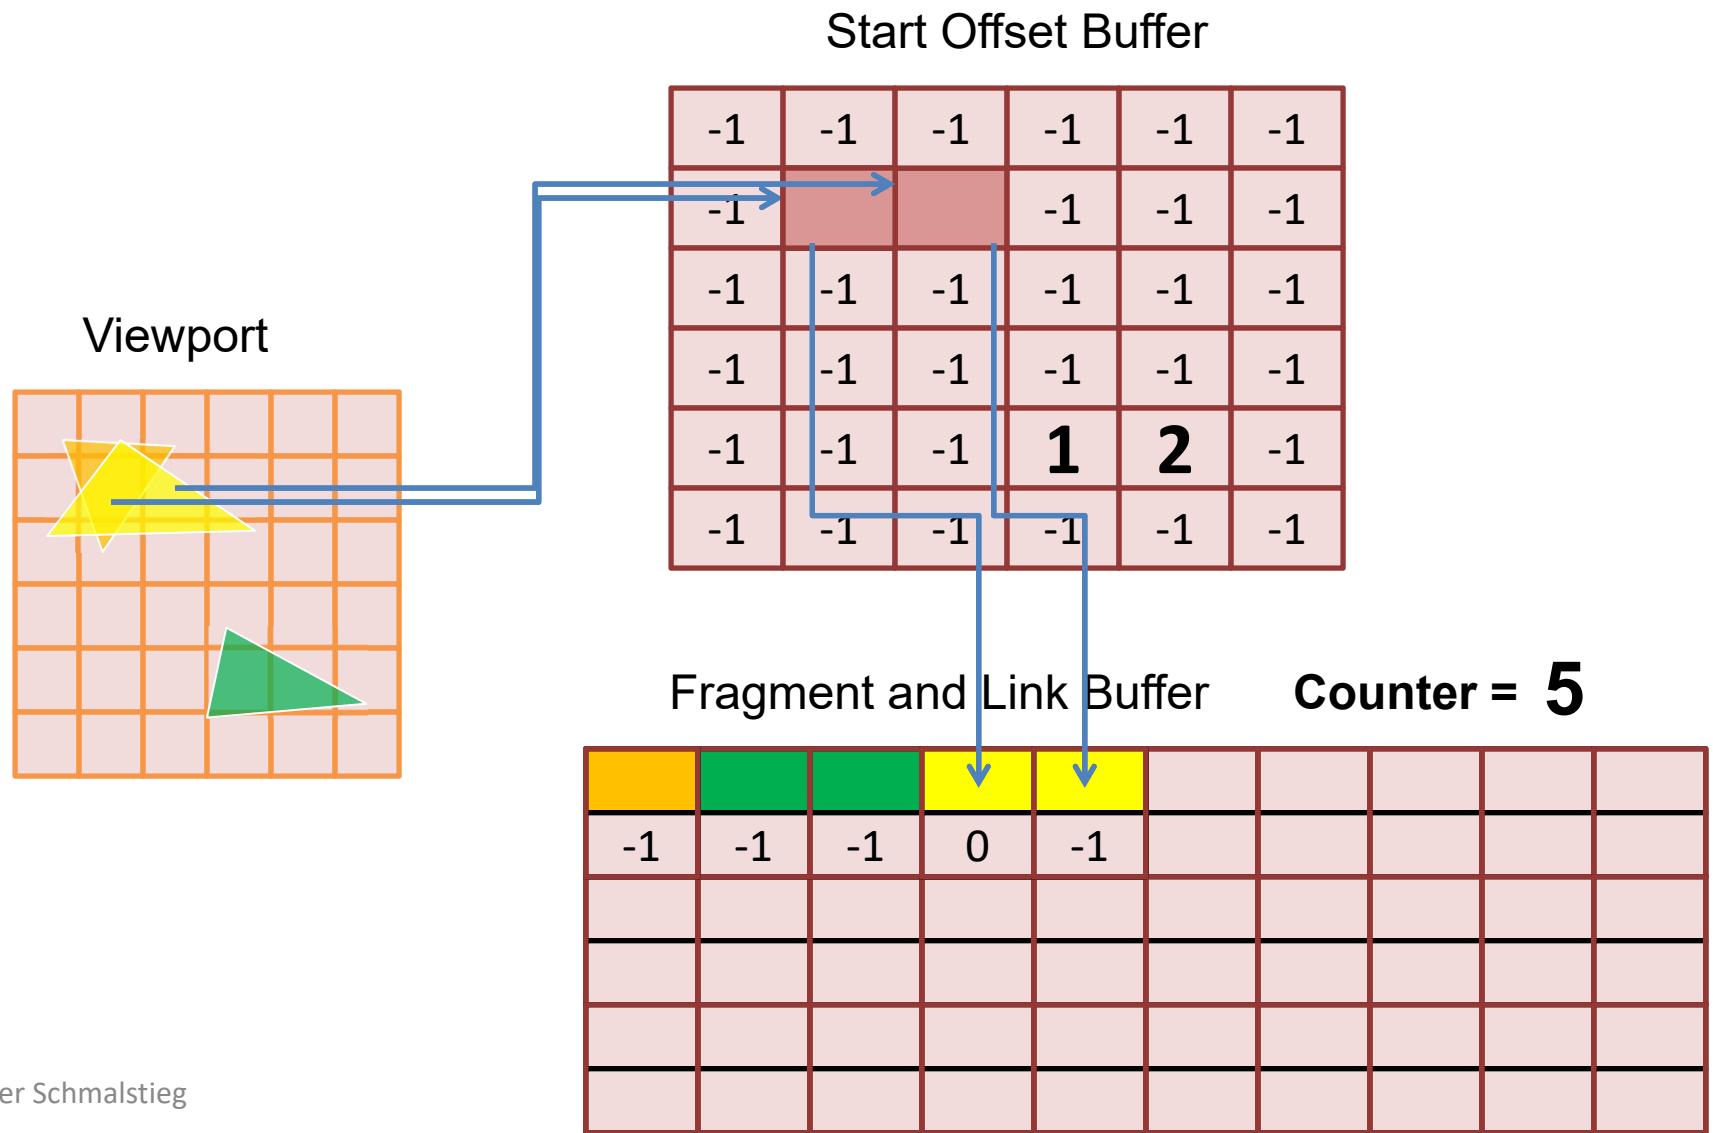
Linked List Creation 2



Linked List Creation 3



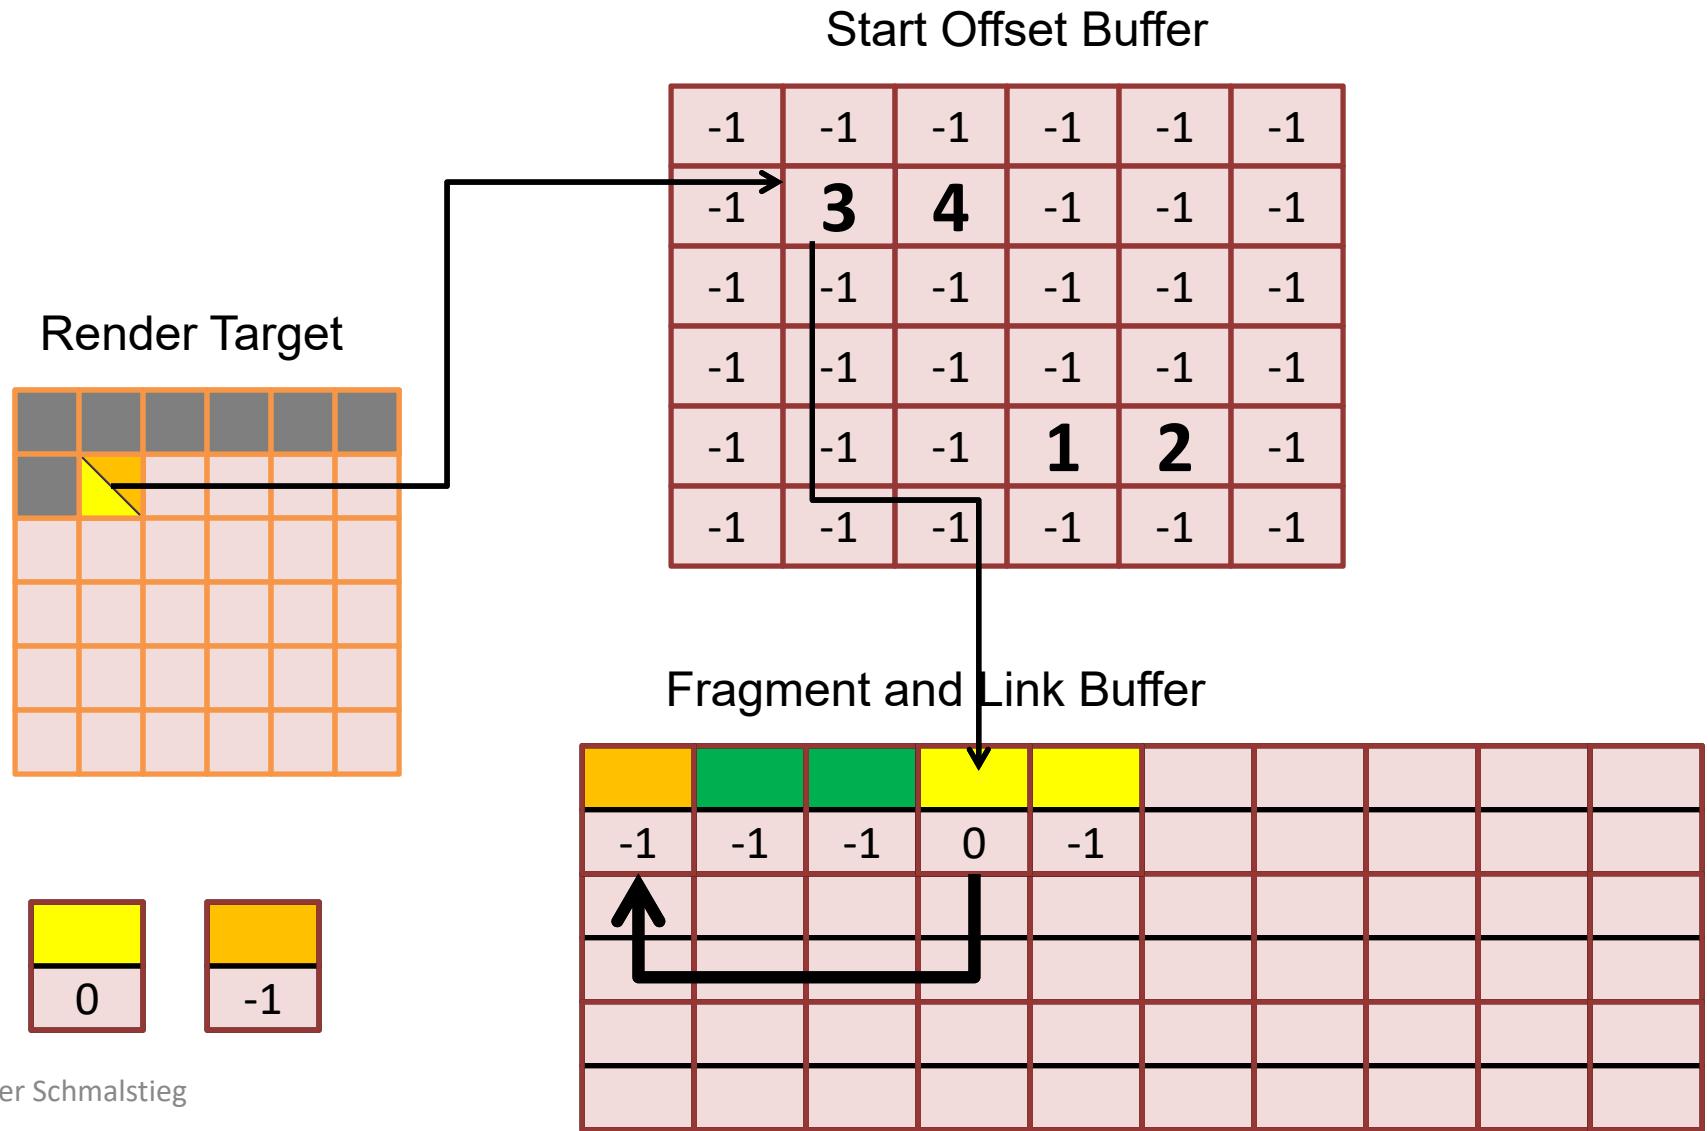
Linked List Creation 4



Linked List Traversal

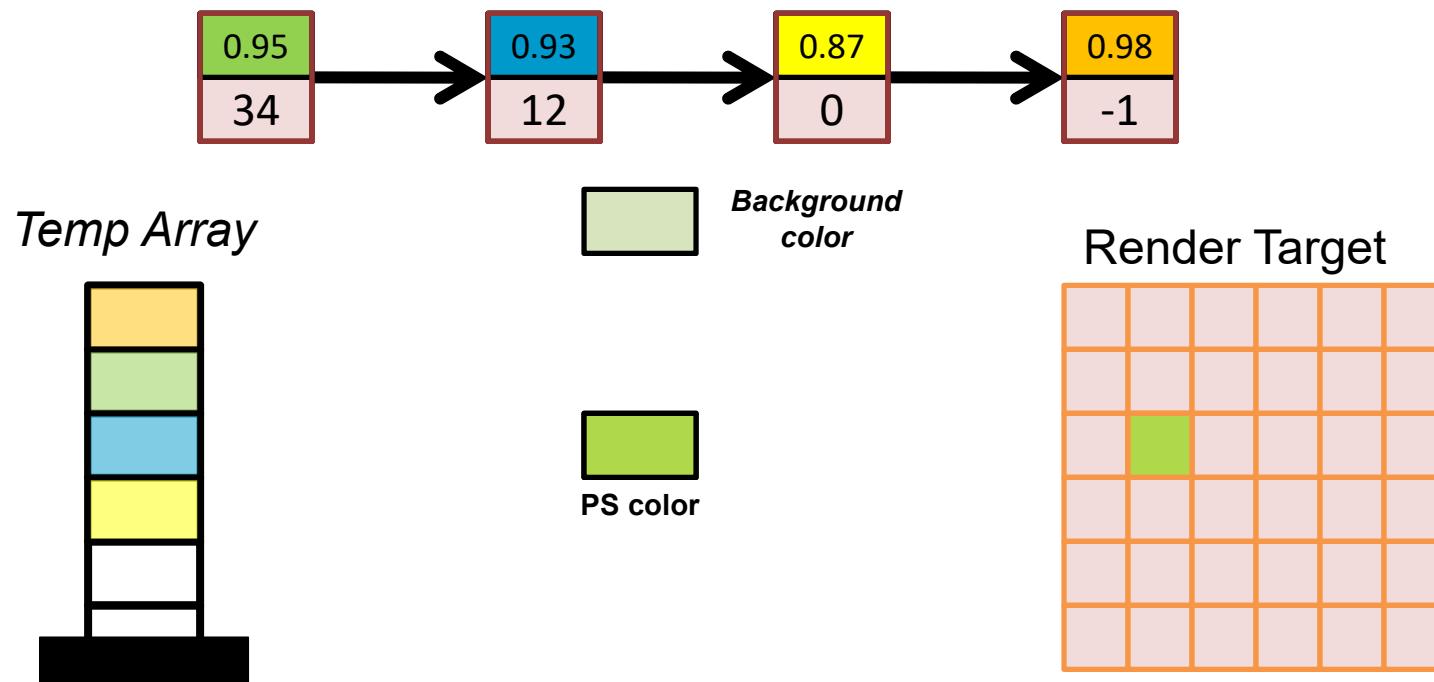
- Render a fullscreen quad (or compute shader)
- For each pixel
 - Parse the linked list
 - Retrieve fragments for this screen position
- Process list of fragments (Sort and Blend)

Rendering from Linked List



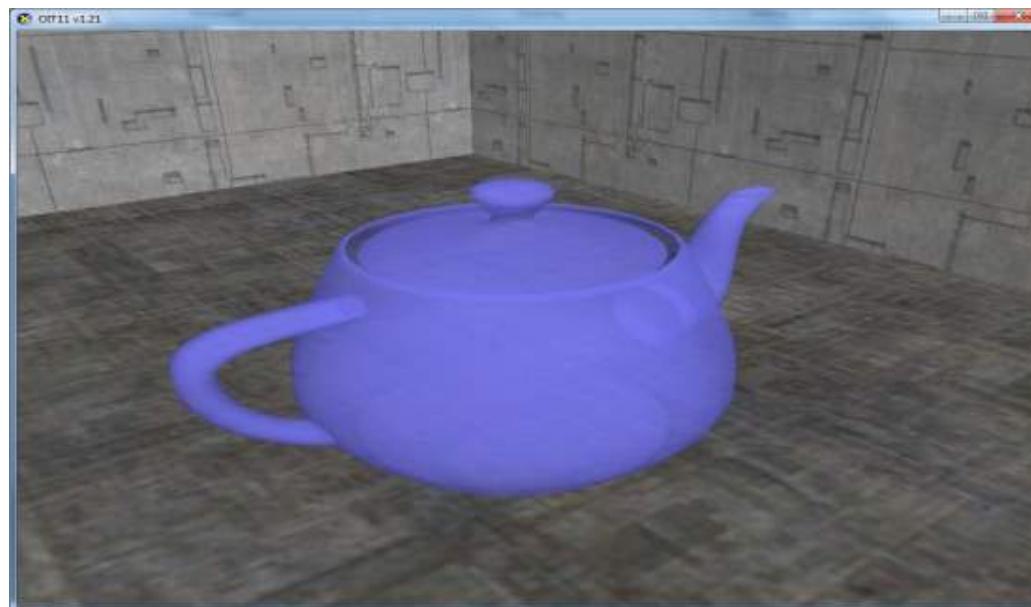
Sorting and Blending

- Blend fragments back to front in PS
- Must sort by depth value in temp.array



Performance Comparison

	Teapot	Dragon
Linked List	743 fps	338 fps
Depth Peeling	579 fps	45 fps
Dual Depth Peeling	---	94 fps



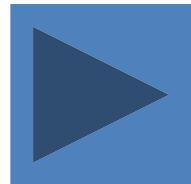
Dieter Schmalstieg



Semi-Global Illumination

Result

- 602K scene triangles VIDEO
 - 254K transparent triangles

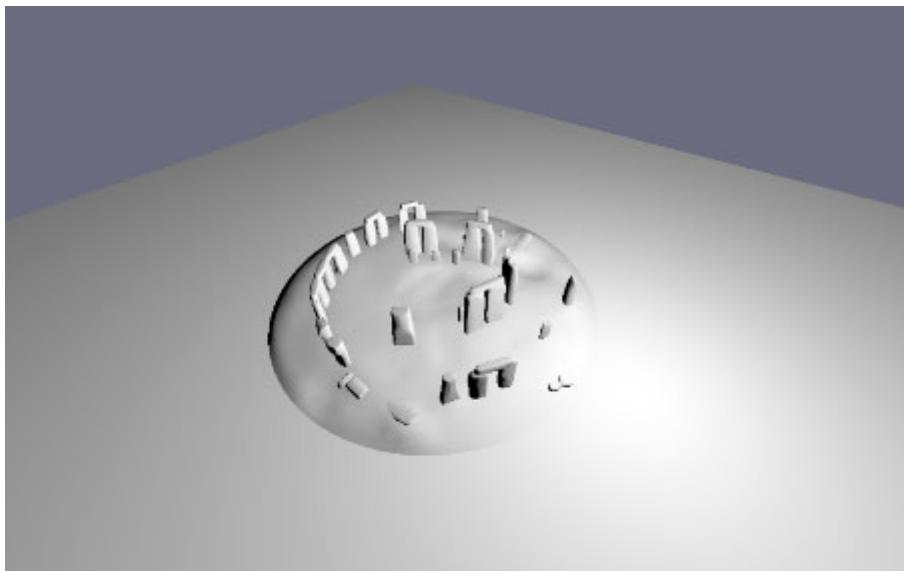


<https://youtu.be/SYrHi4jF4dU>

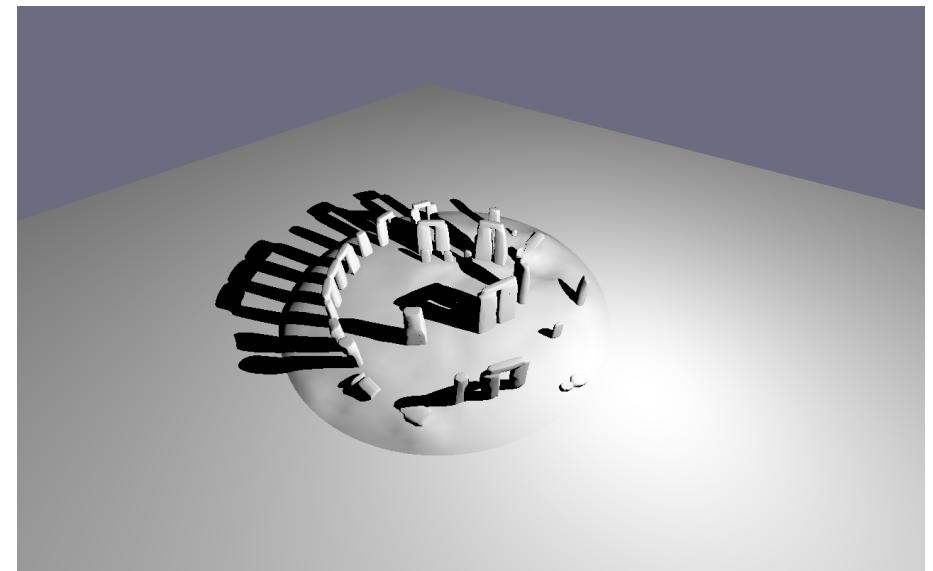


Shadows

- Light blocked by other objects
- Important visual cue
 - Relative location
 - Position of light



Dieter Schmalstieg



Semi-Global Illumination

Shadows for Visual Realism

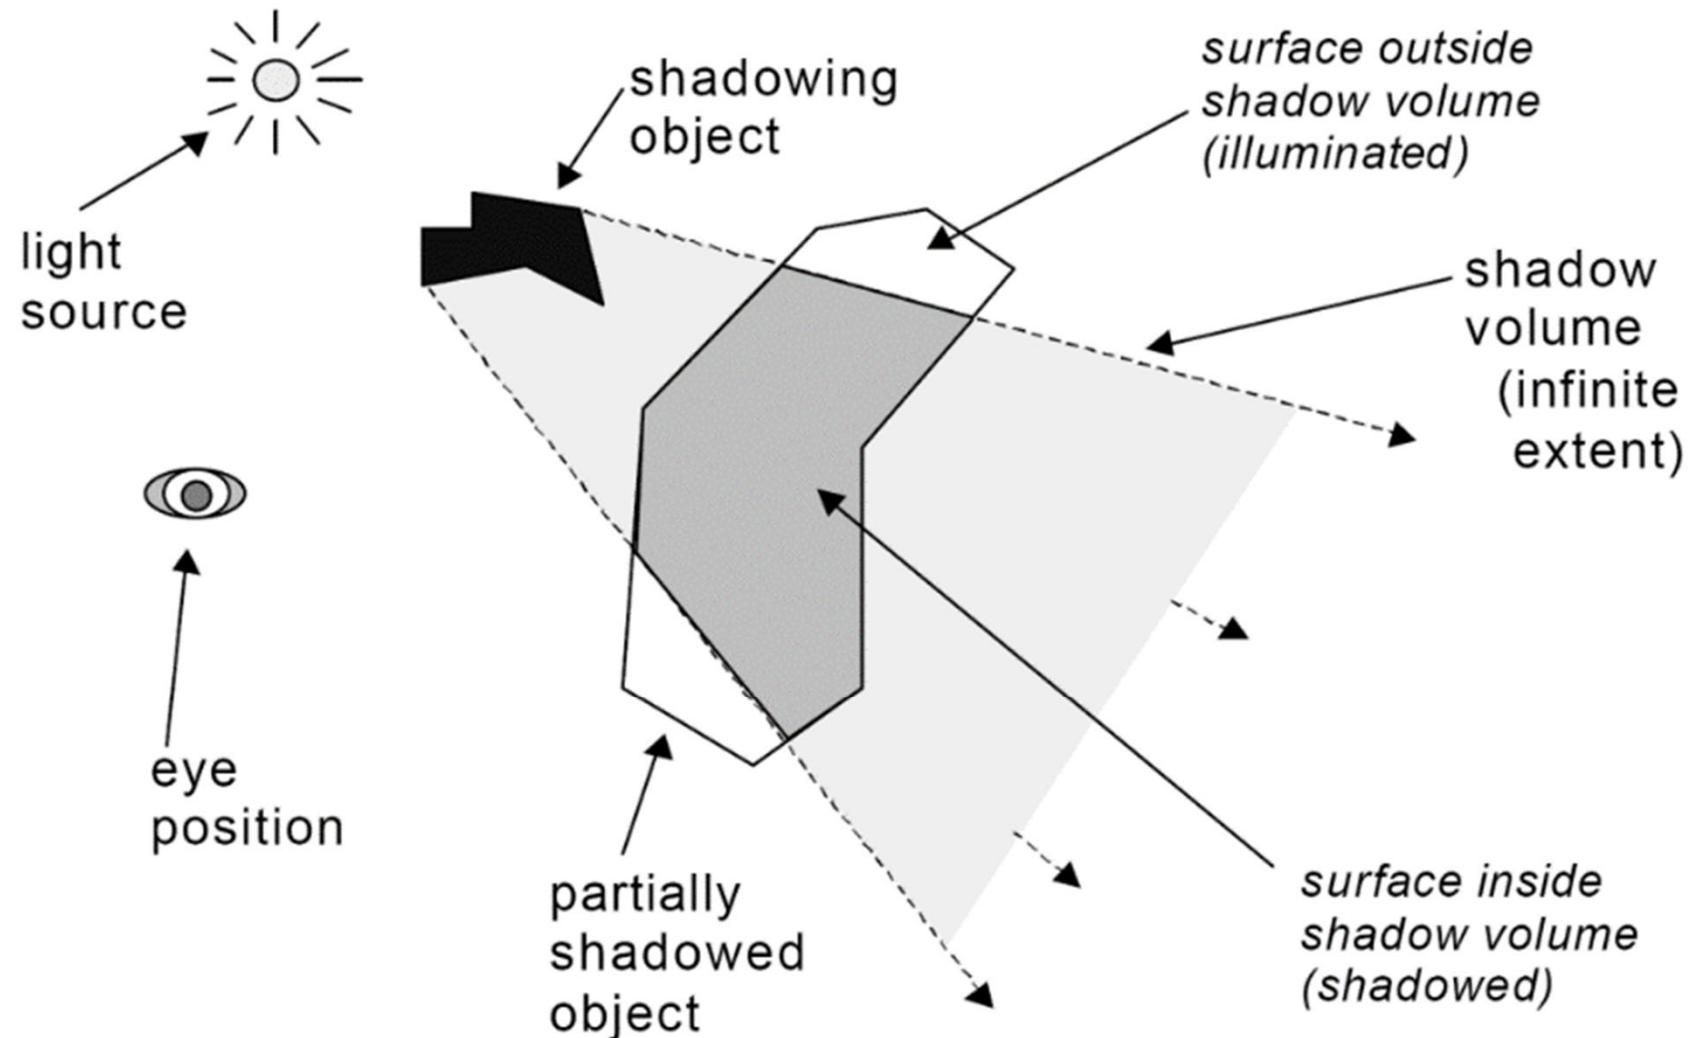


Objects look like they are “floating” → Shadows can fix that!

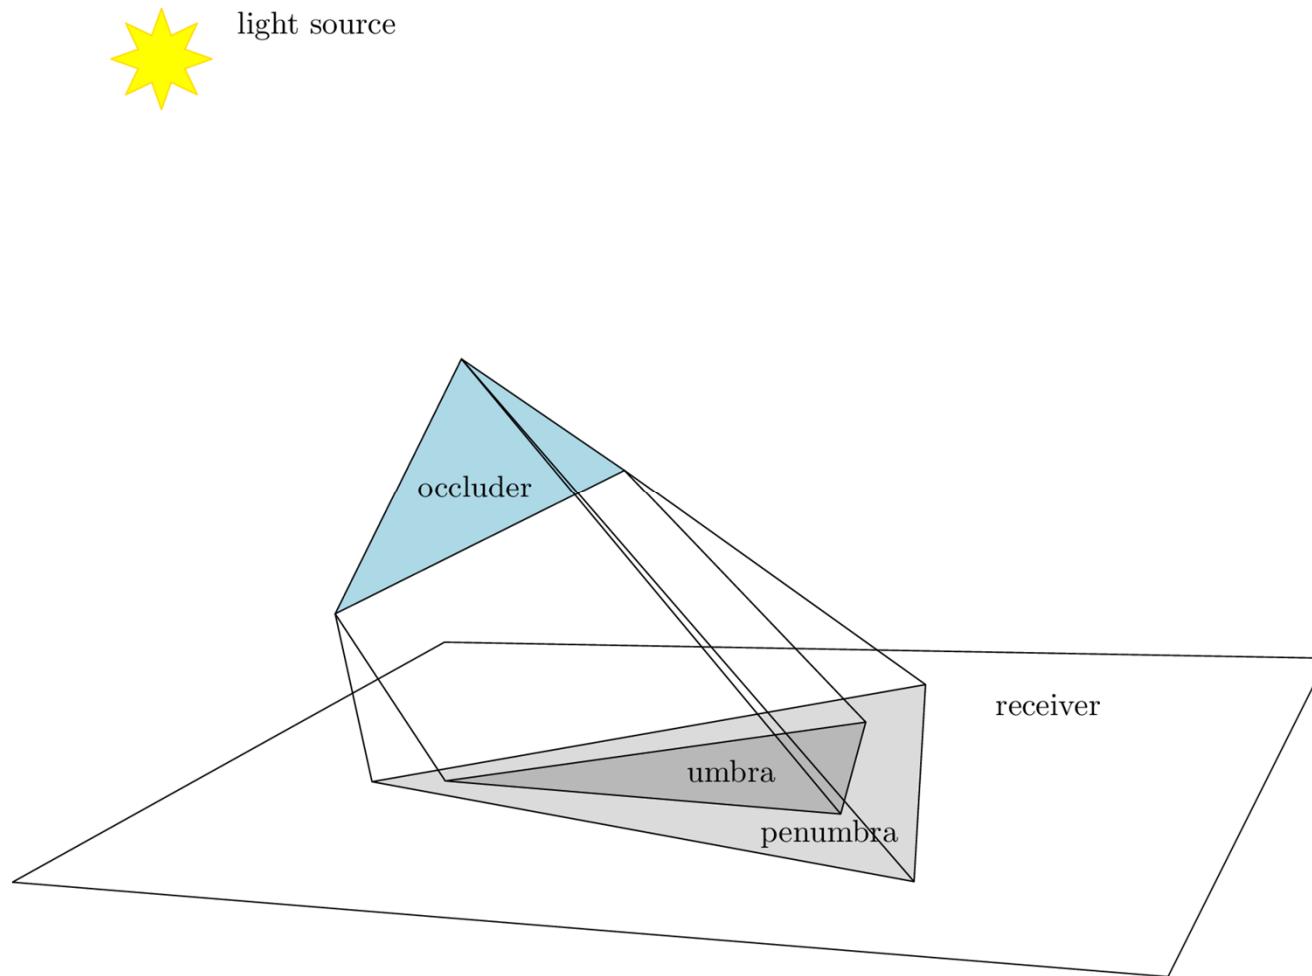
Shadows are Non-Local

- Shadows are a non-local effect between
 - Light source
 - Object currently being rendered (“receiver”)
 - Object blocking light from reaching rendered object (“blocker”, “occluder”, “caster”)
- Shading can be seen as pseudo self-shadowing

Shadow Geometry

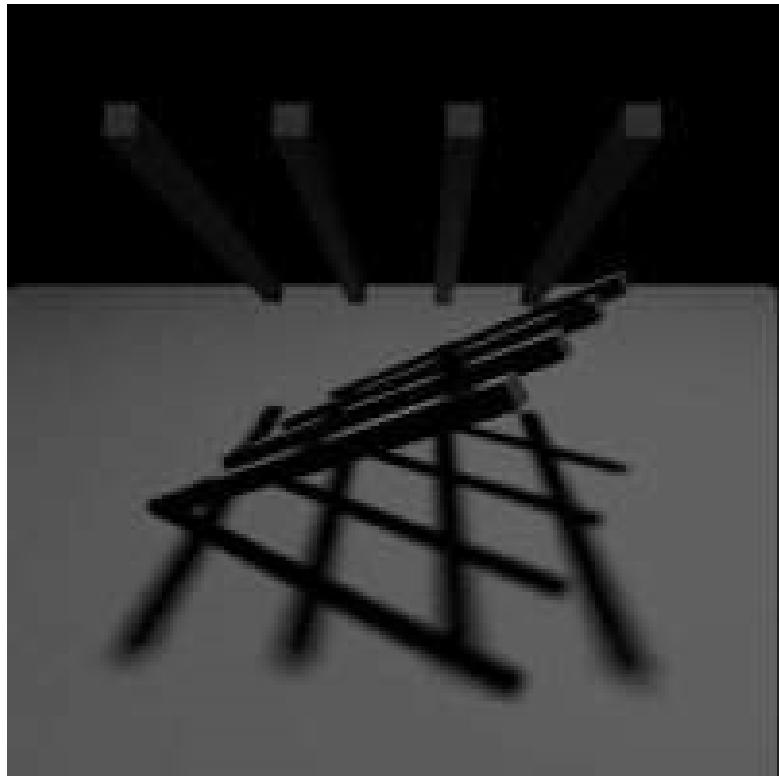


Umbra and Penumbra



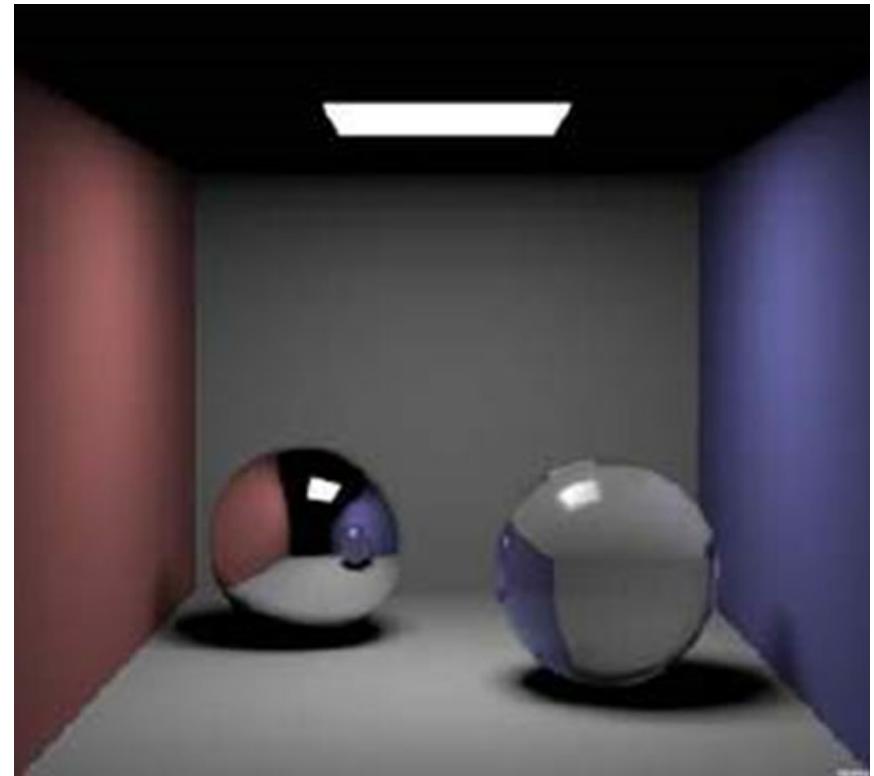
Global Illumination

Shadows as byproduct of global light transport



Radiosity

Dieter Schmalstieg



Raytracing

Semi-Global Illumination

49

The Problem

- Shadows are a *global* illumination (GI) problem
- GI is hard, especially in online rendering
 - Global information not available during rendering
- Address problem with two-pass techniques
 1. Compute and store (global) shadow information
 2. Render scene using that information

Real-Time Shadows

- Static shadows
 - Light maps - see also previous lectures
 - Updating in real-time is problematic
- Approximate shadows
- Planar projected shadows
- Shadow Maps
- Stencil Shadows

Static Shadows

- Glue to surface whatever we want
- Idea: incorporate shadows into light maps
 - For each texel, cast ray to each light source
- Bake soft shadows in light maps
 - Not by texture filtering alone
 - Must also sample area light sources

Static Soft Shadow Example

1 sample

no filtering



filtering

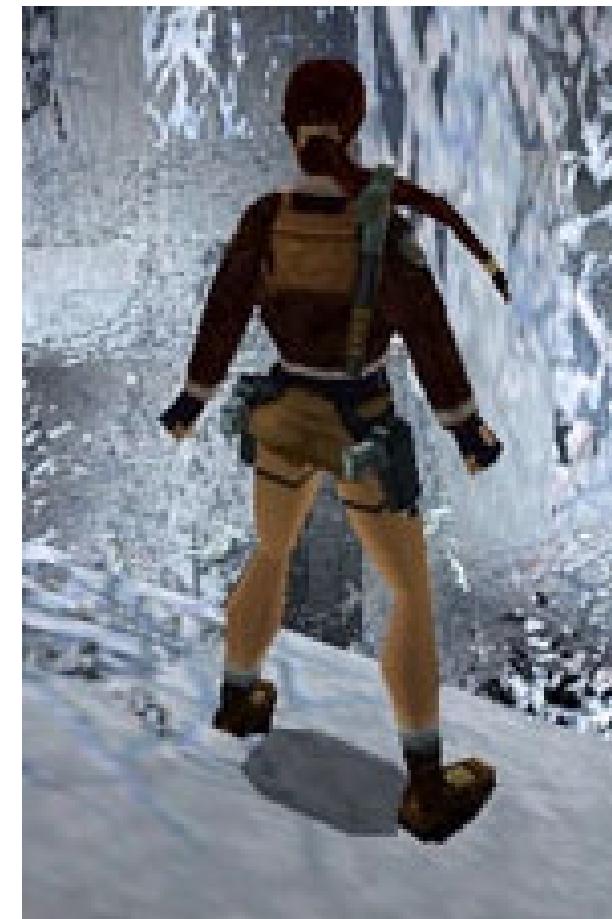


n samples



Approximate Shadows 1

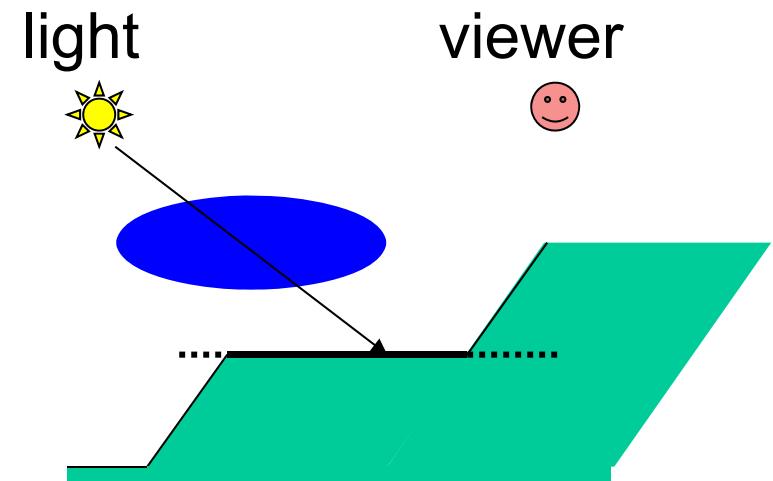
- Hand-drawn approximate geometry
 - Perceptual studies suggest:
shape not so important
 - Minimal cost



You can tell how young Lara was by counting the vertices of her shadow...

Approximate Shadows 2

- Dark polygon (maybe with texture)
 - Cast ray from light source through object center
 - Blend polygon into frame buffer at location of hit
 - May apply additional rotation/scale/translation to incorporate distance and receiver orientation
- Problem with z-quantization



Blend at hit polygon + z-test equal \rightarrow z-buffer quantization errors!

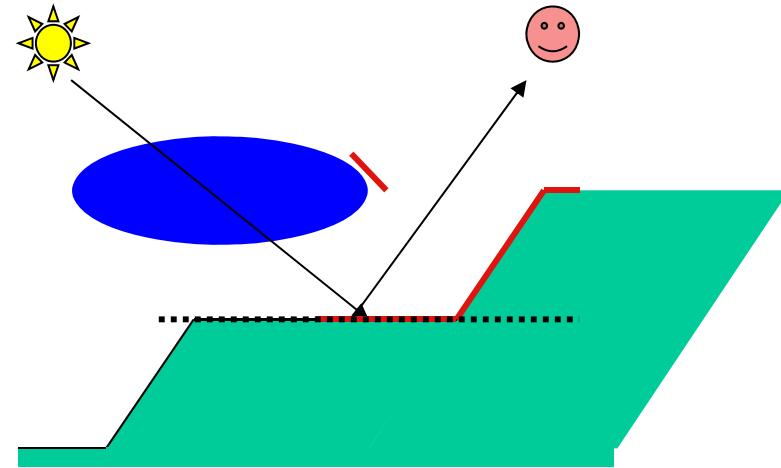
Approximate Shadows 3



light ☼

viewer ☺

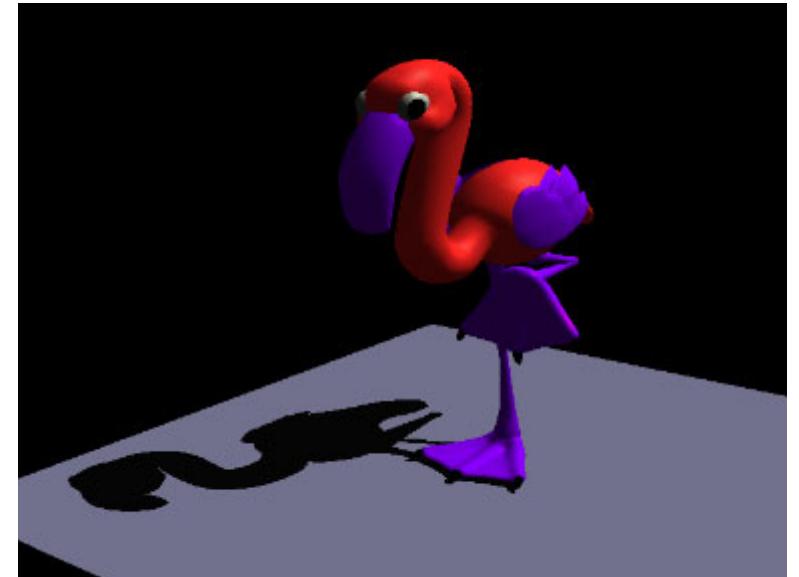
Elevate above hit polygon + Z-test less or equal
→ shadow too big → may appear floating



No z-test, only one eye ray → shadow too big, maybe in wrong place

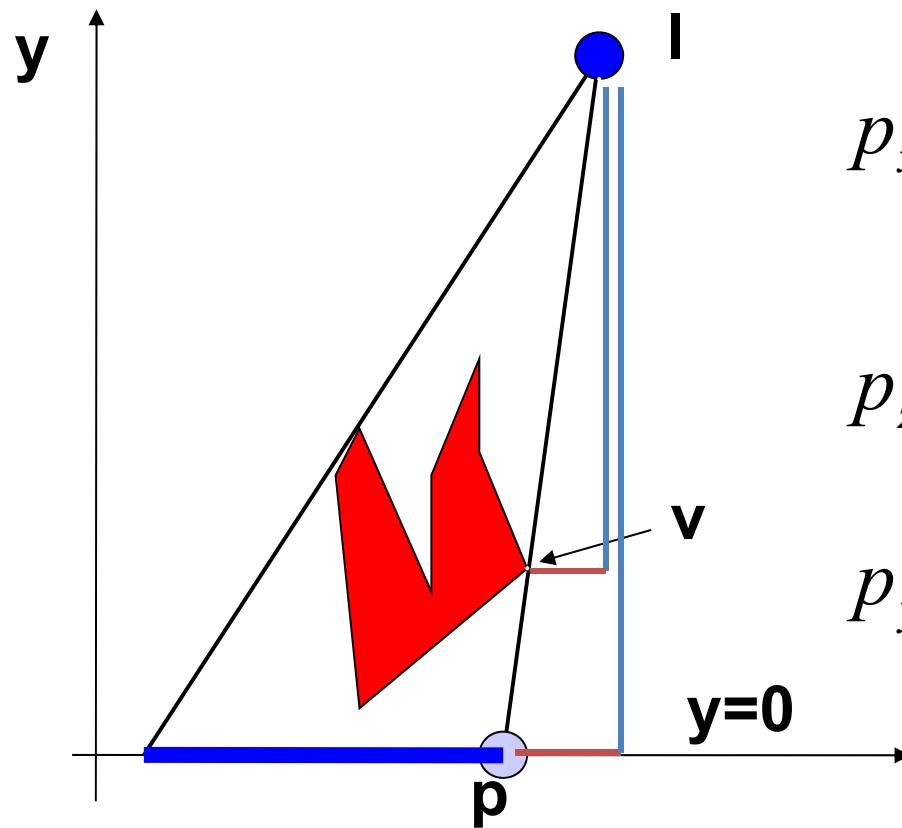
Planar Projected Shadows

- Very simple trick [Jim Blinn, 1988]
- Supports arbitrary shadow caster
- But only planar shadow receiver
- Method
 - ModelView matrix transforms object into plane
 - Object is drawn “flat”
 - Darken surface underneath using framebuffer blend



Projection Onto XZ-Plane

- $\mathbf{p} = \mathbf{M} \mathbf{v}$
- Similar triangles



$$\frac{p_x - l_x}{v_x - l_x} = \frac{l_y}{l_y - v_y}$$

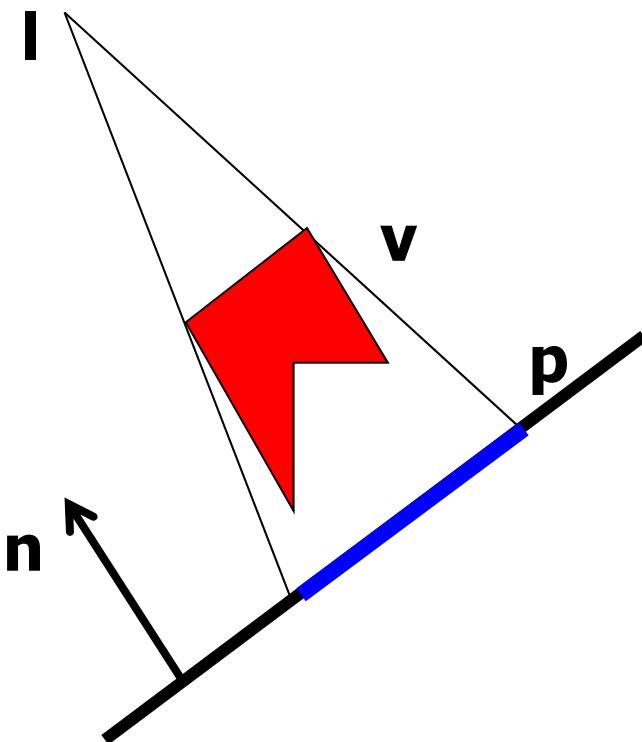
$$p_x = \frac{l_y v_x - l_x v_y}{l_y - v_y}$$

$$p_z = \frac{l_y v_z - l_z v_y}{l_y - v_y}$$

$$p_y = 0$$

$$\mathbf{M} = \begin{pmatrix} l_y & -l_x & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & -l_z & l_y & 0 \\ 0 & -1 & 0 & l_y \end{pmatrix}$$

Projection Onto General Plane



$$\text{ray: } \mathbf{p} = \mathbf{l} + \alpha(\mathbf{v} - \mathbf{l}) \quad \dots \dots \alpha = (p - l) / (v - l)$$

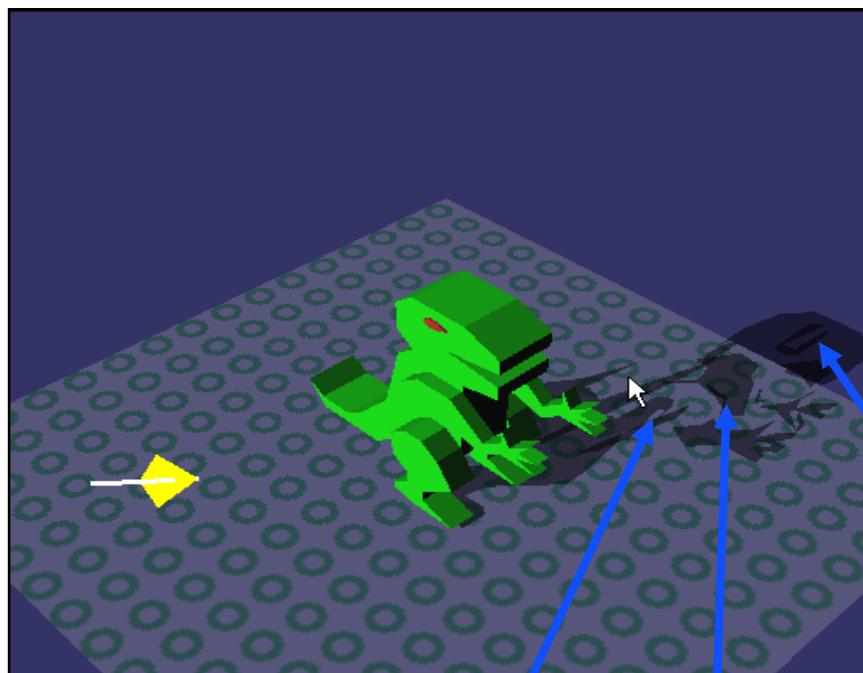
$$\text{plane: } \mathbf{n} \cdot \mathbf{p} + d = 0 \quad \dots \dots \begin{aligned} p &= -d/n \\ \alpha &= (-d/n - l) / (v - l) = -d/(n(v - l)) - l/(v - l) \end{aligned}$$

$$\mathbf{p} = \mathbf{l} - \frac{d + \mathbf{n} \cdot \mathbf{l}}{\mathbf{n} \cdot (\mathbf{v} - \mathbf{l})} (\mathbf{v} - \mathbf{l})$$

$$\mathbf{p} = \mathbf{M} \mathbf{v} \rightarrow \mathbf{M} = \begin{pmatrix} \mathbf{n} \cdot \mathbf{l} + d - l_x n_x & -l_x n_y & -l_x n_z & -l_x d \\ -l_y n_x & \mathbf{n} \cdot \mathbf{l} + d - l_y n_y & -l_y n_z & -l_y d \\ -l_z n_x & -l_z n_y & \mathbf{n} \cdot \mathbf{l} + d - l_z n_z & -l_z d \\ -n_x & -n_y & -n_z & \mathbf{n} \cdot \mathbf{l} \end{pmatrix}$$

Projected Shadows - Artifacts

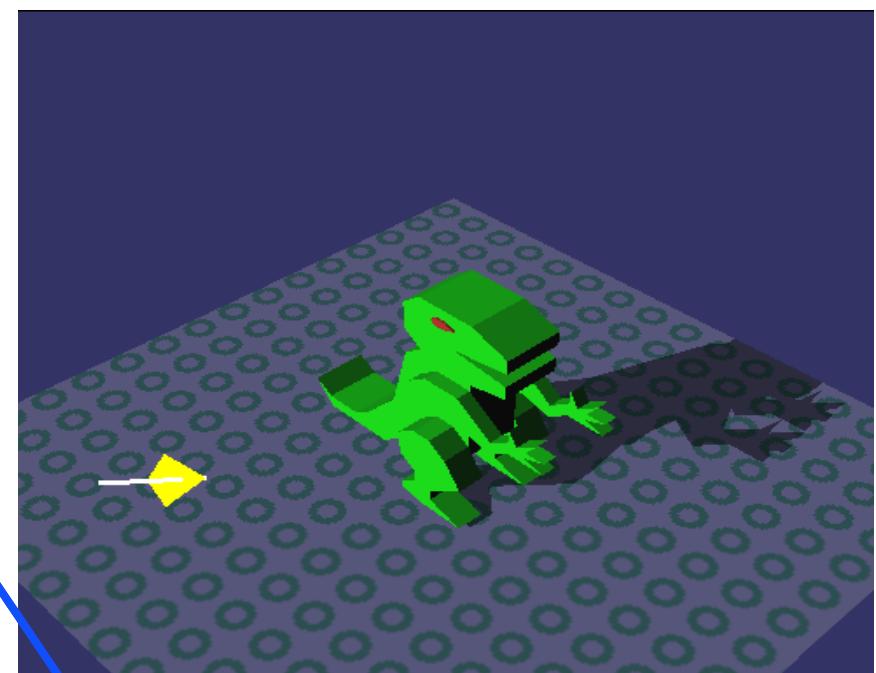
Bad



Z-fighting

double blending

Good



extends off ground

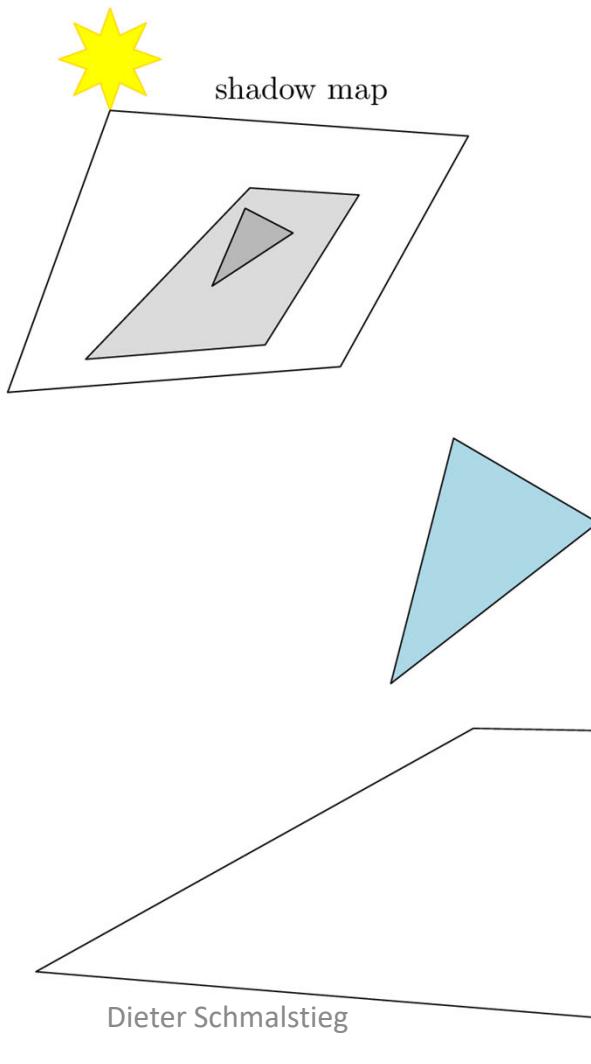
Projected Shadows with Stencil

- Clear stencil buffer to 0
- Draw receiver plane
 - Write 1 to stencil buffer
- Draw 3D object
- Draw shadow
 - Where stencil is 1
 - Write 0 to stencil buffer
 - No z-fighting problem

Projected Shadows - Summary

- Easy to implement
 - GLQuake first game to implement it
- Only practical for very few, large receivers
- No self-shadowing
- Possible remaining artifacts: wrong shadows
 - Objects behind light source
 - Objects behind receiver

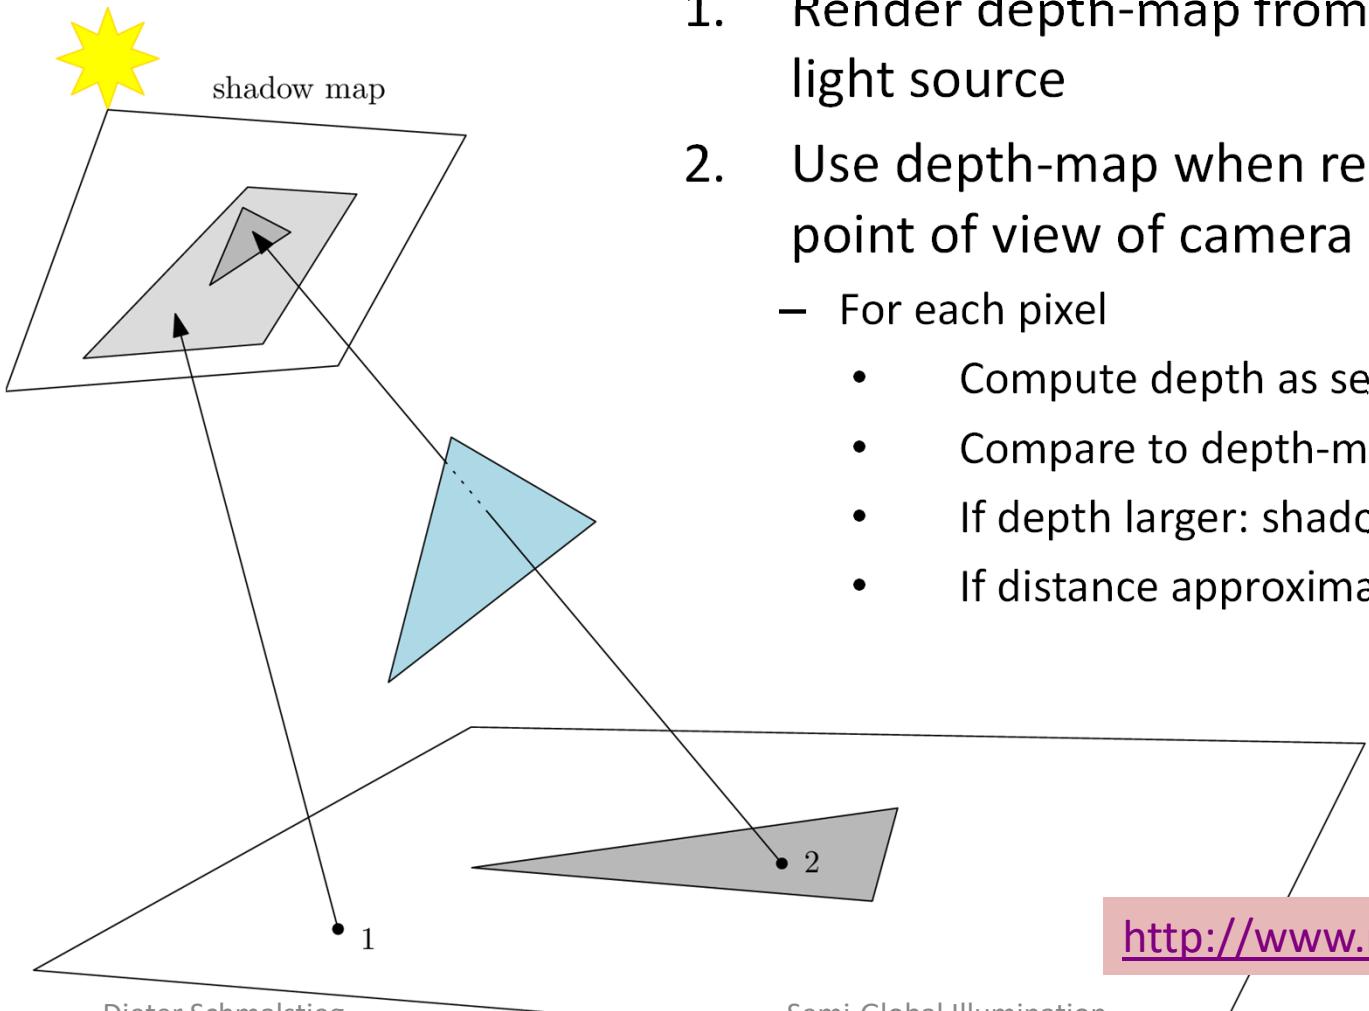
Shadow Mapping Algorithm 1



Two-step algorithm:

1. Render depth-map from point light source of view
2. Use depth-map when rendering scene from point of view of camera

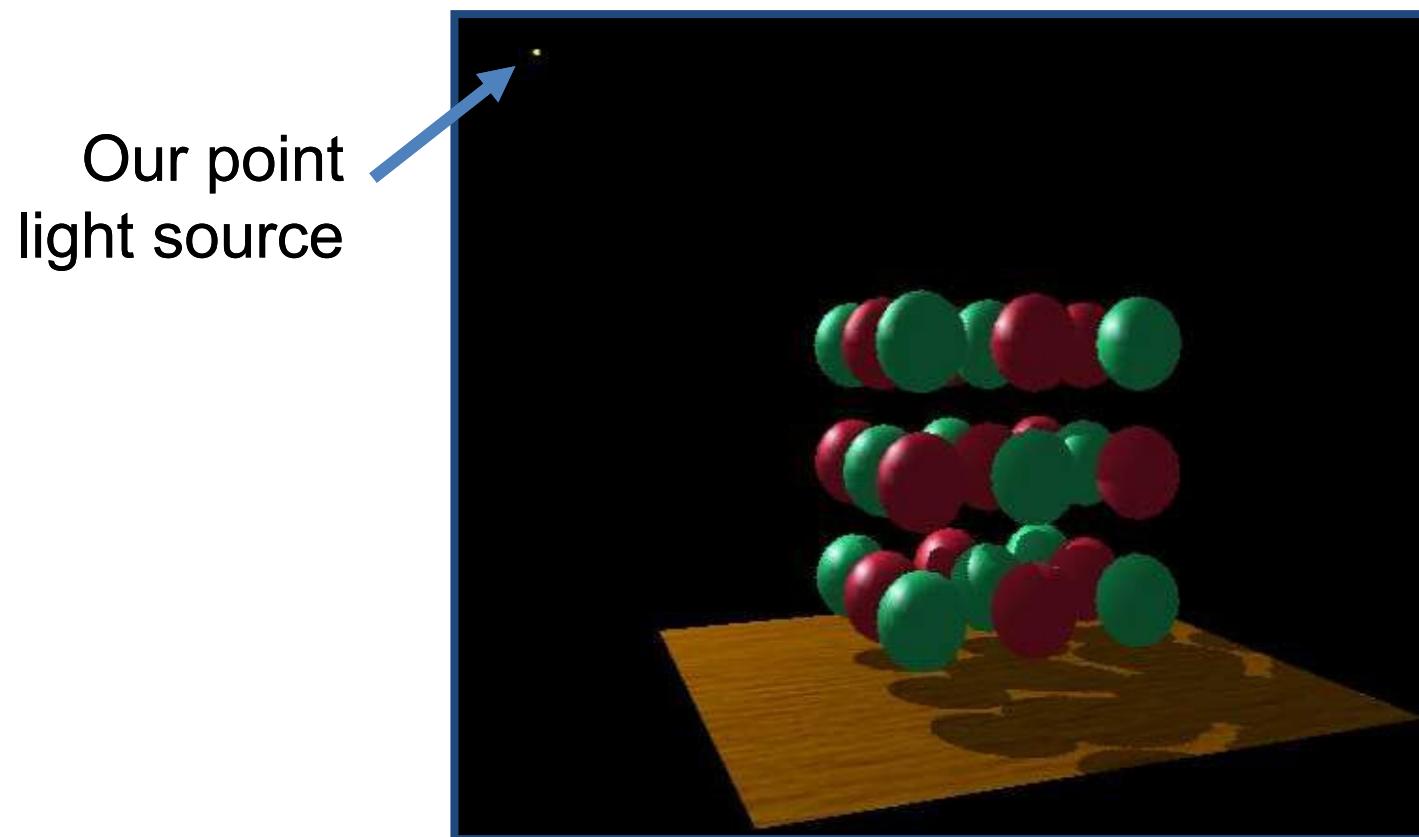
Shadow Mapping Algorithm 2



1. Render depth-map from point of view of the light source
2. Use depth-map when rendering scene from point of view of camera
 - For each pixel
 - Compute depth as seen from light source
 - Compare to depth-map
 - If depth larger: shadow
 - If distance approximately equal: light

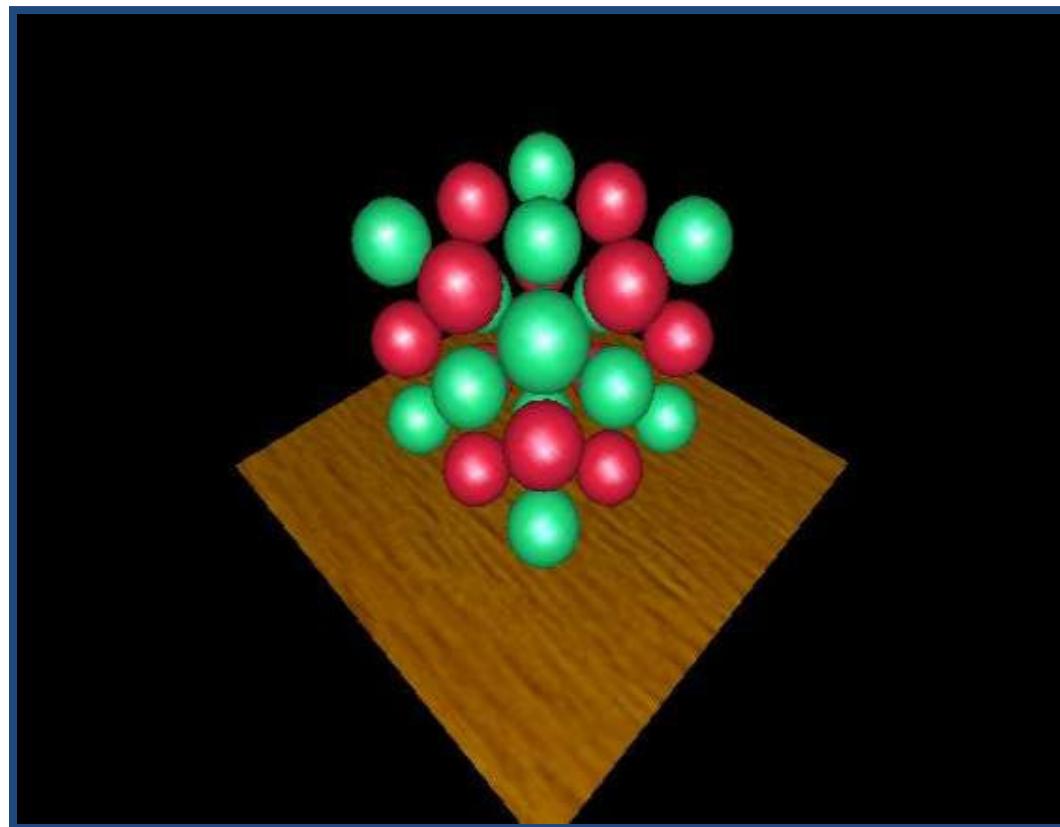
<http://www.nutty.ca/webgl/shadows/>

Shadow Mapping - Point Light

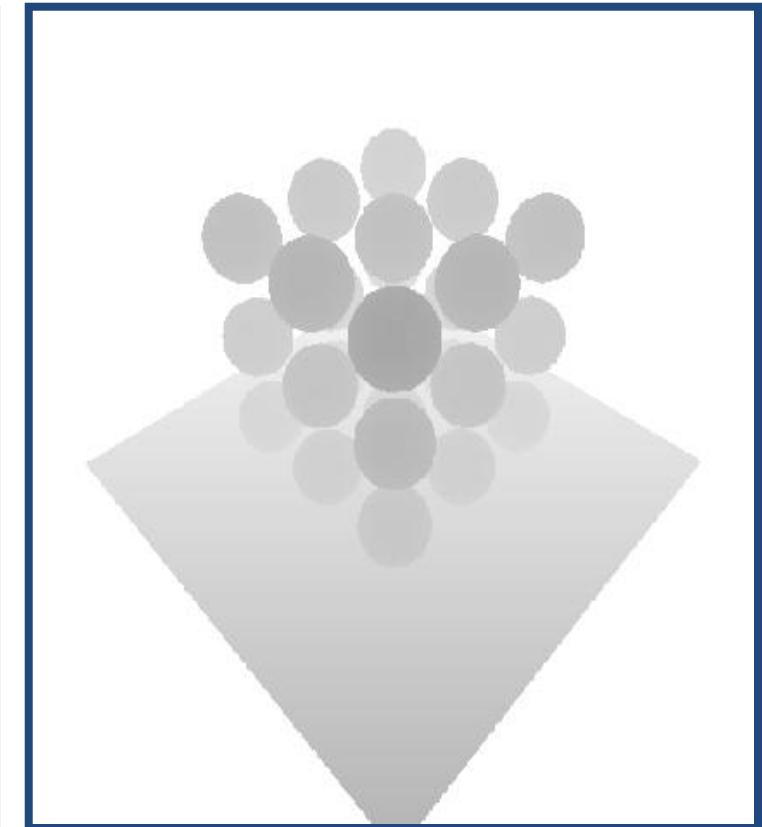


Shadow Map as Depth Map

- Rendering the Depth-Map



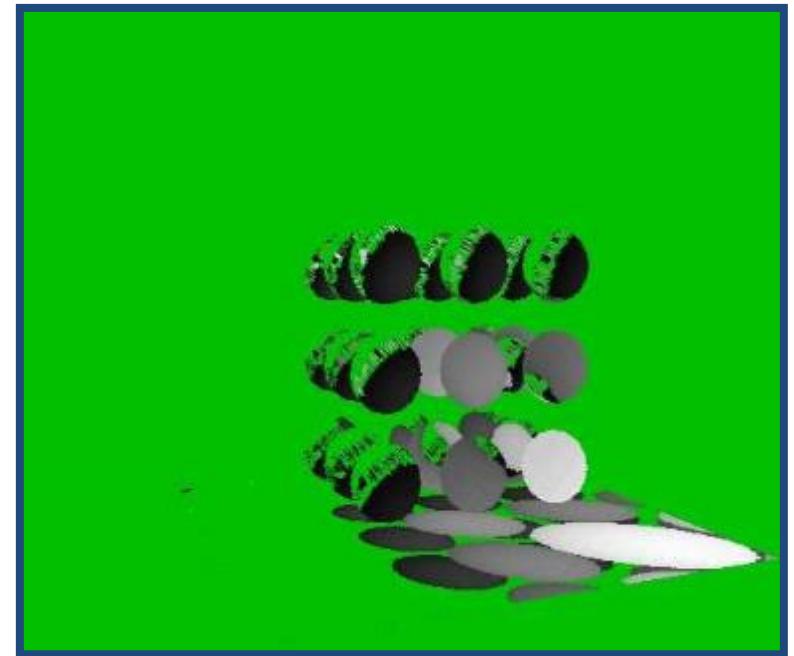
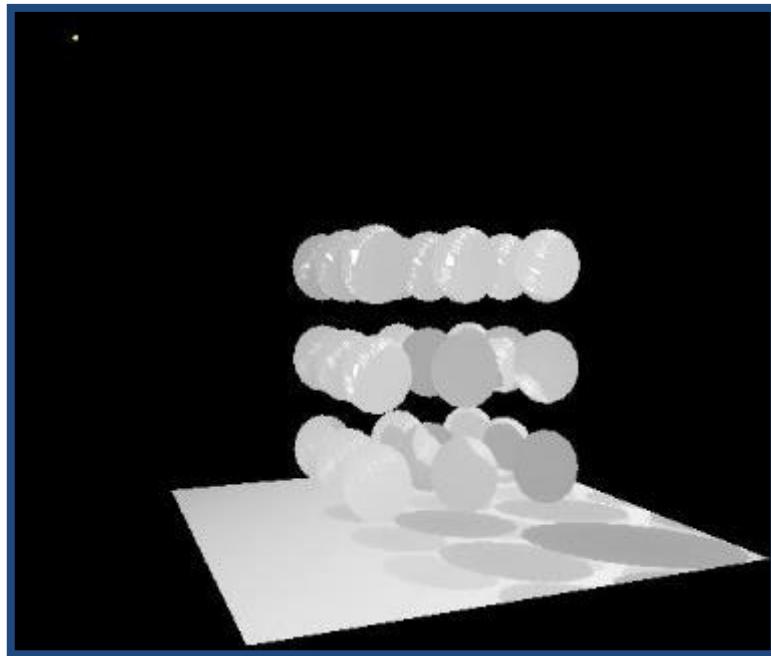
Dieter Schmalstieg



Semi-Global Illumination

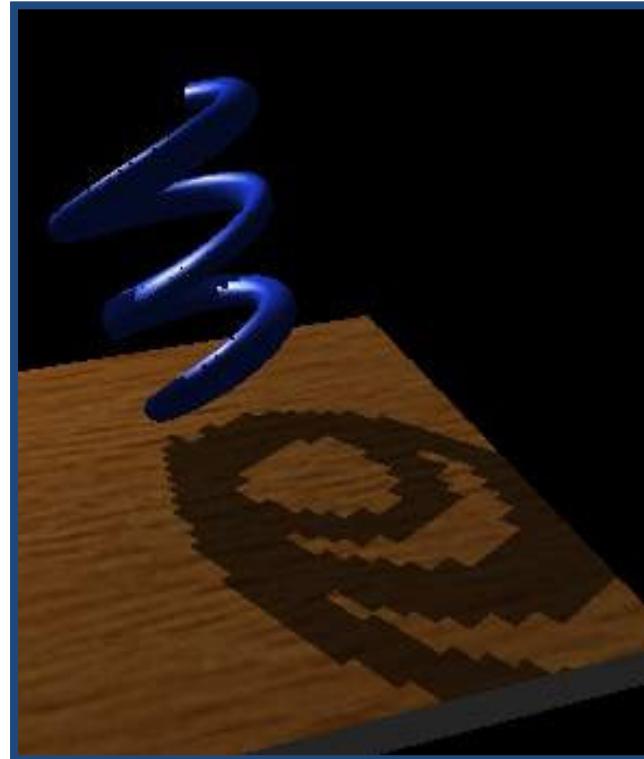
Shadow Mapping - Depth Test

- Green: depth approximately equal
- Non-green: shadow

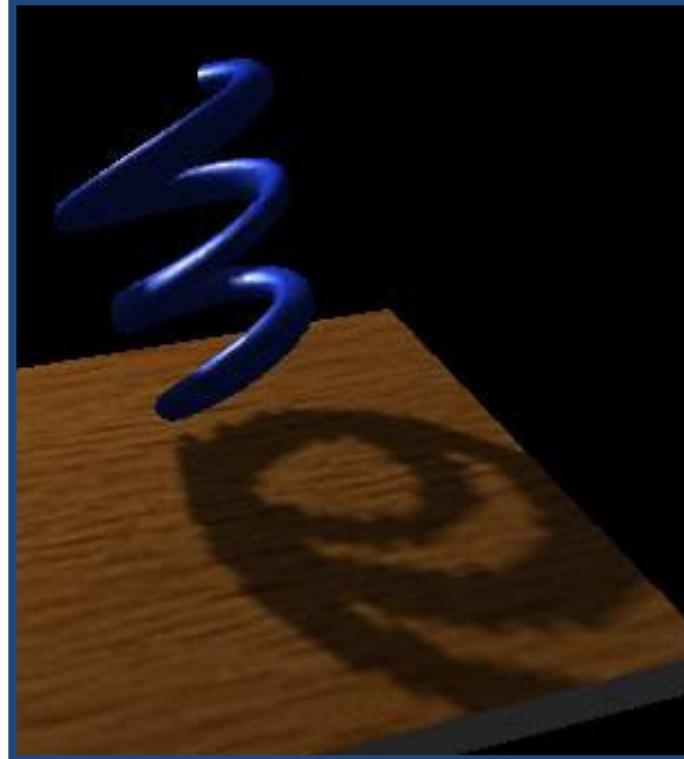


Aliasing

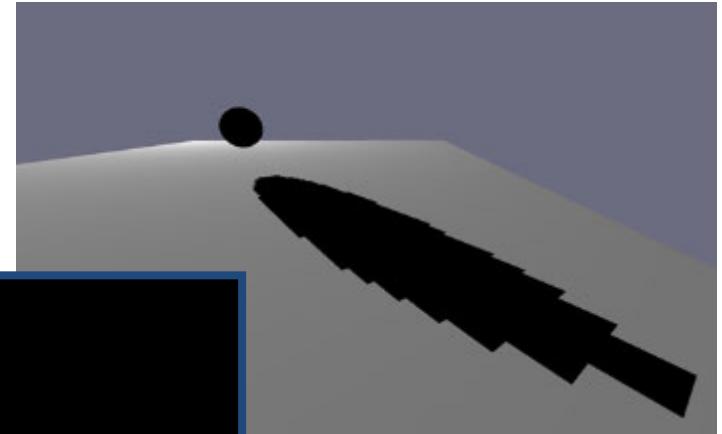
- Two components
 - Perspective aliasing
 - Projection aliasing



Dieter Schmalstieg

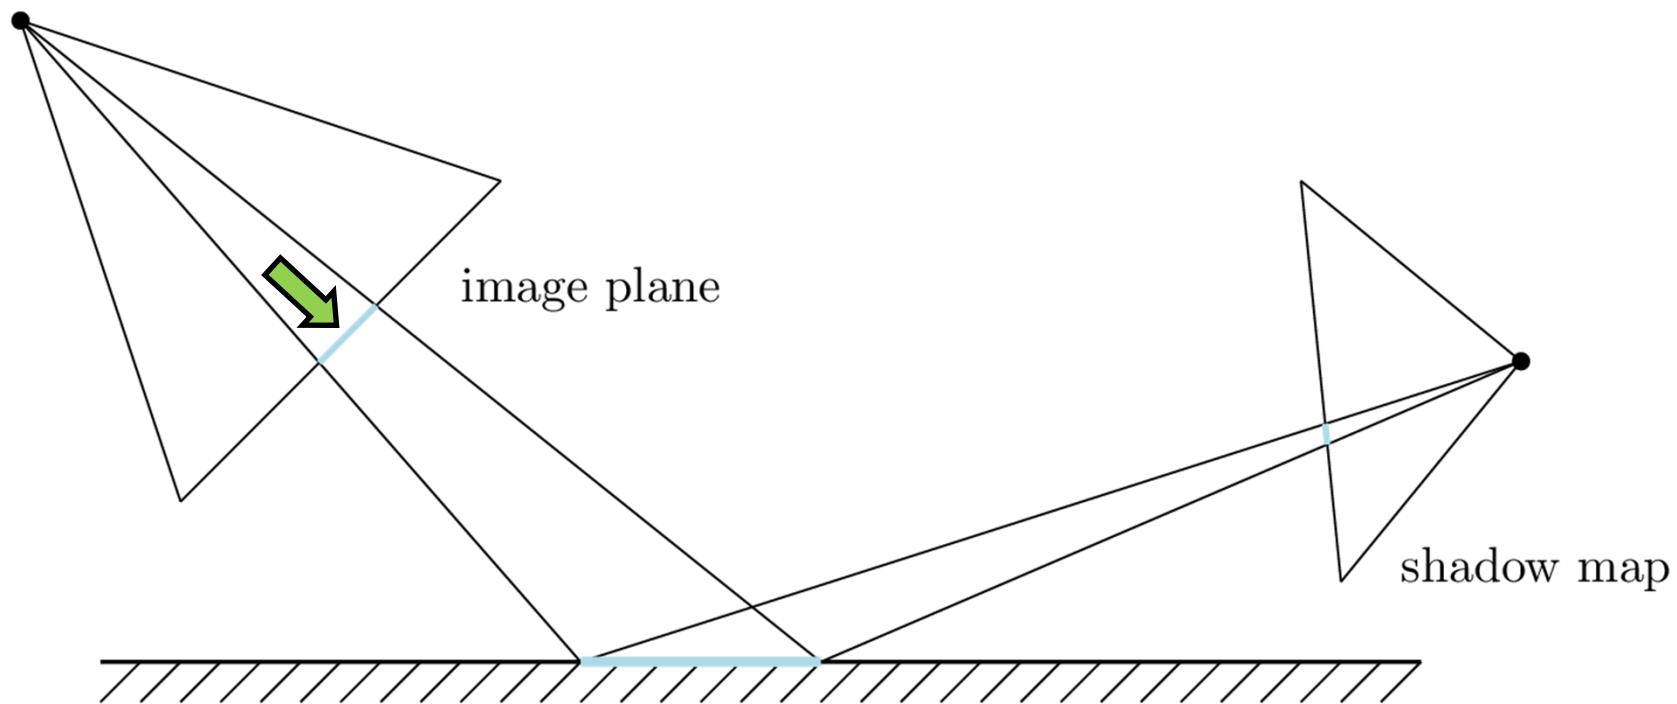


Semi-Global Illumination



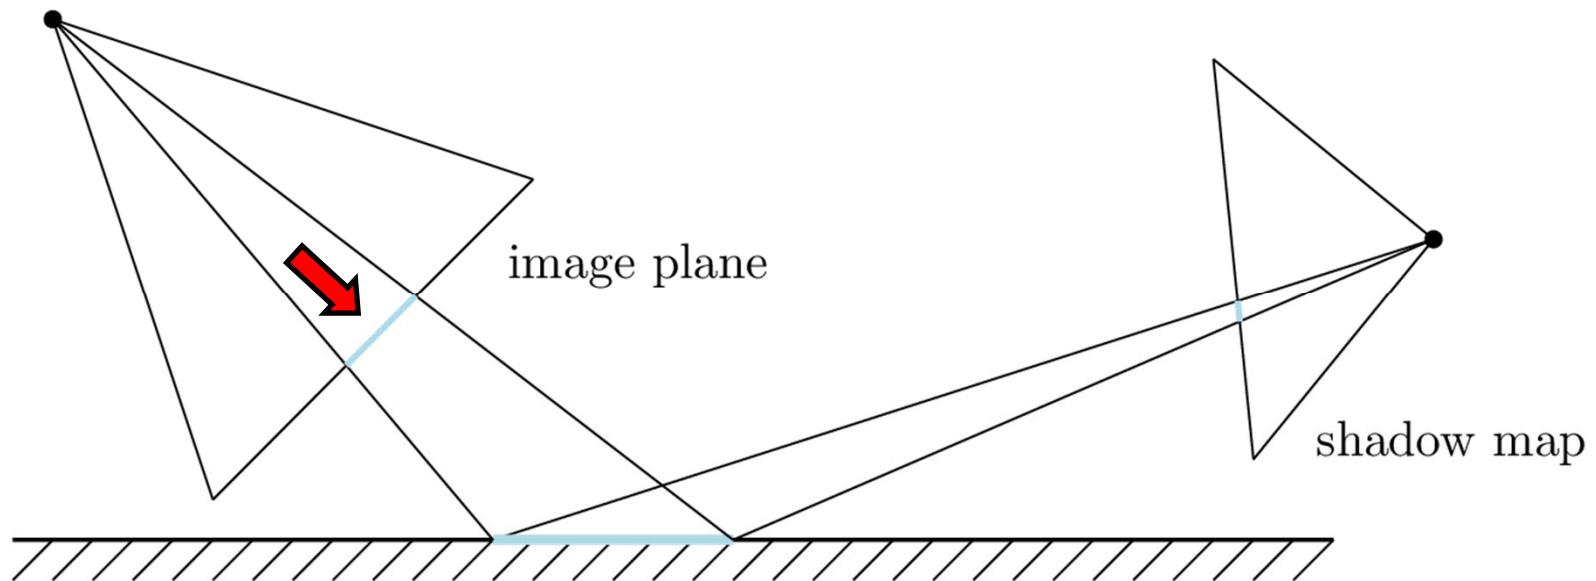
Perspective Aliasing 1

- Texel resolution in image plane is okay



Perspective Aliasing 2

- If camera moves closer,
texel resolution in image space gets worse



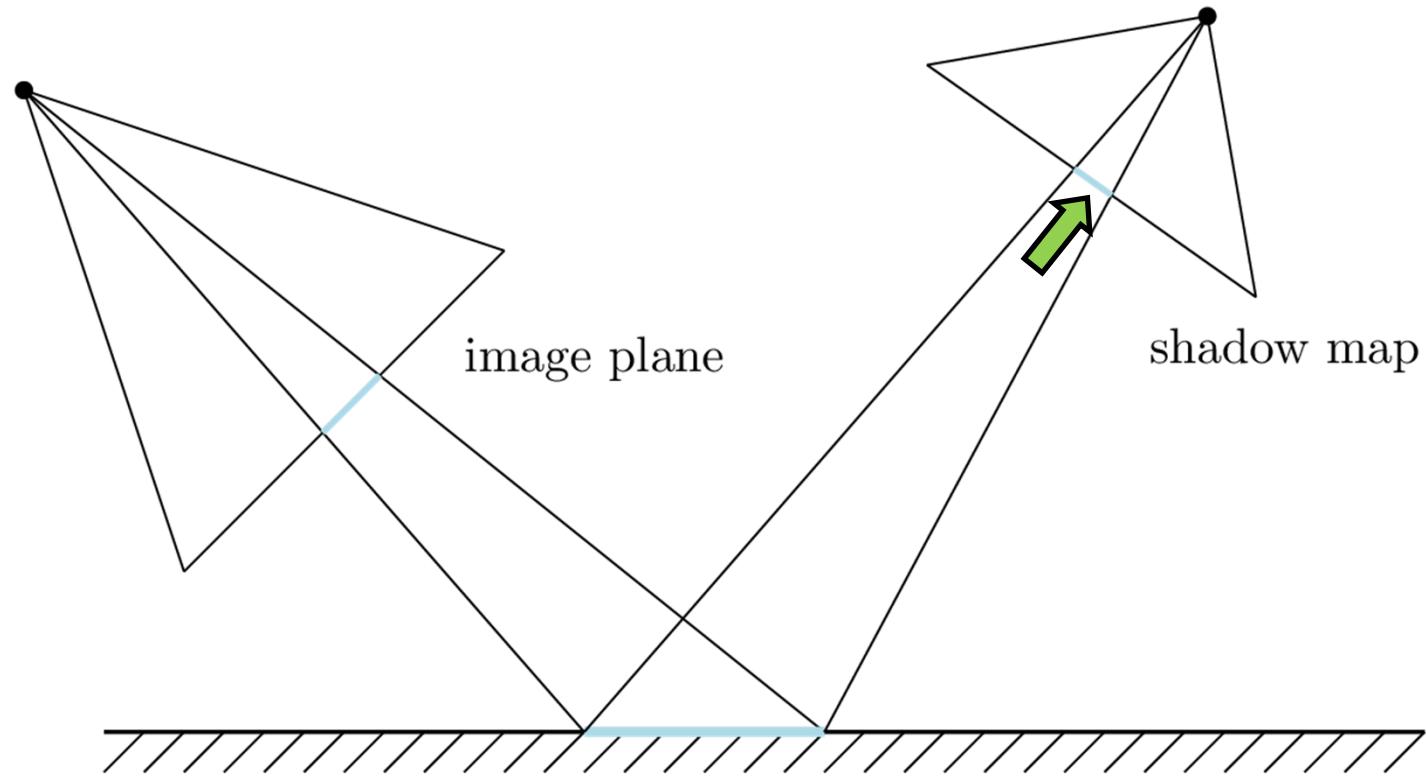
Example: Perspective Aliasing

Shadow texels large and blocky



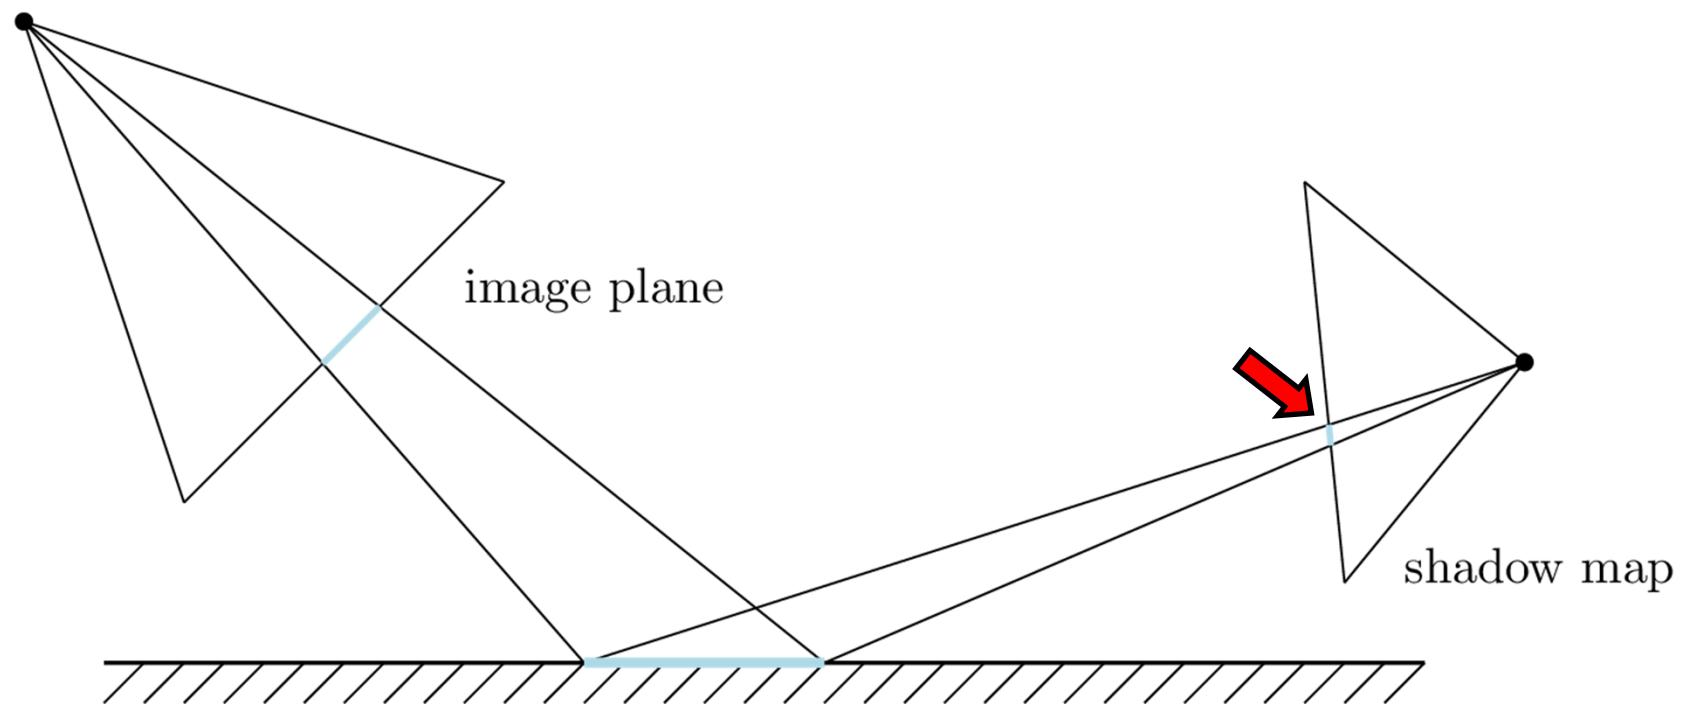
Projection Aliasing 1

- Shadow map resolution okay



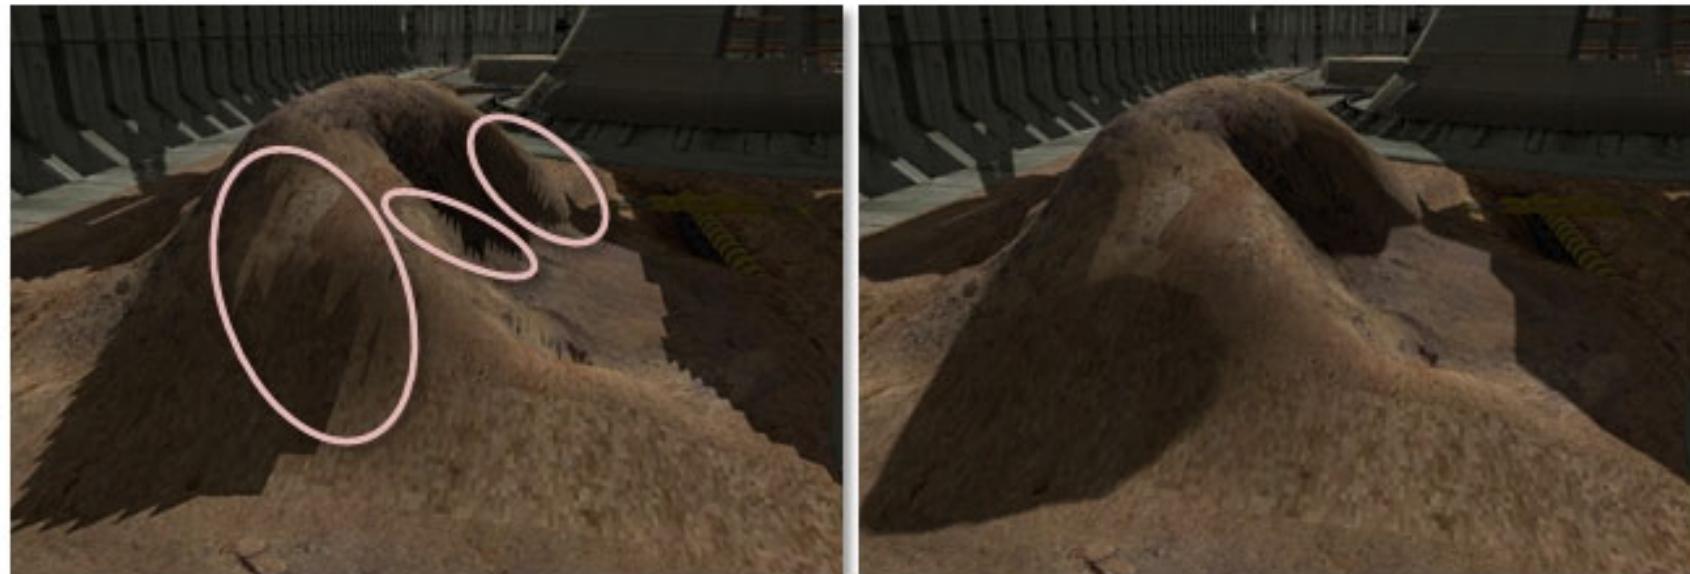
Projection Aliasing 2

- If angle to light source gets larger, shadow map resolution gets worse



Example: Projection Aliasing

Shadow texels stretched

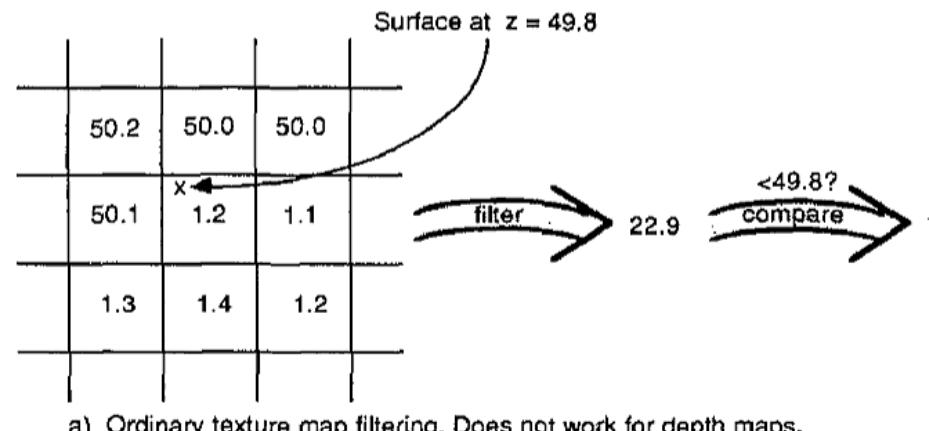


How to Reduce Aliasing?

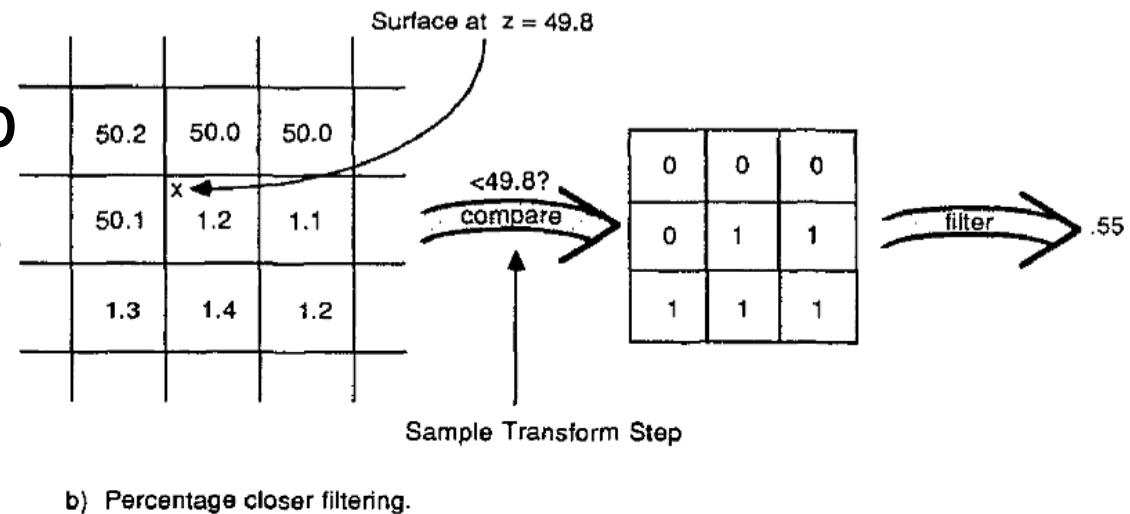
- Increase shadow map resolution
 - Not practical in general
- Practical approaches
 - Anti-aliasing by filtering
 - Optimize sample distribution in shadow map

Percentage Closer Filtering

- Normal bilinear filtering cannot be used on depth

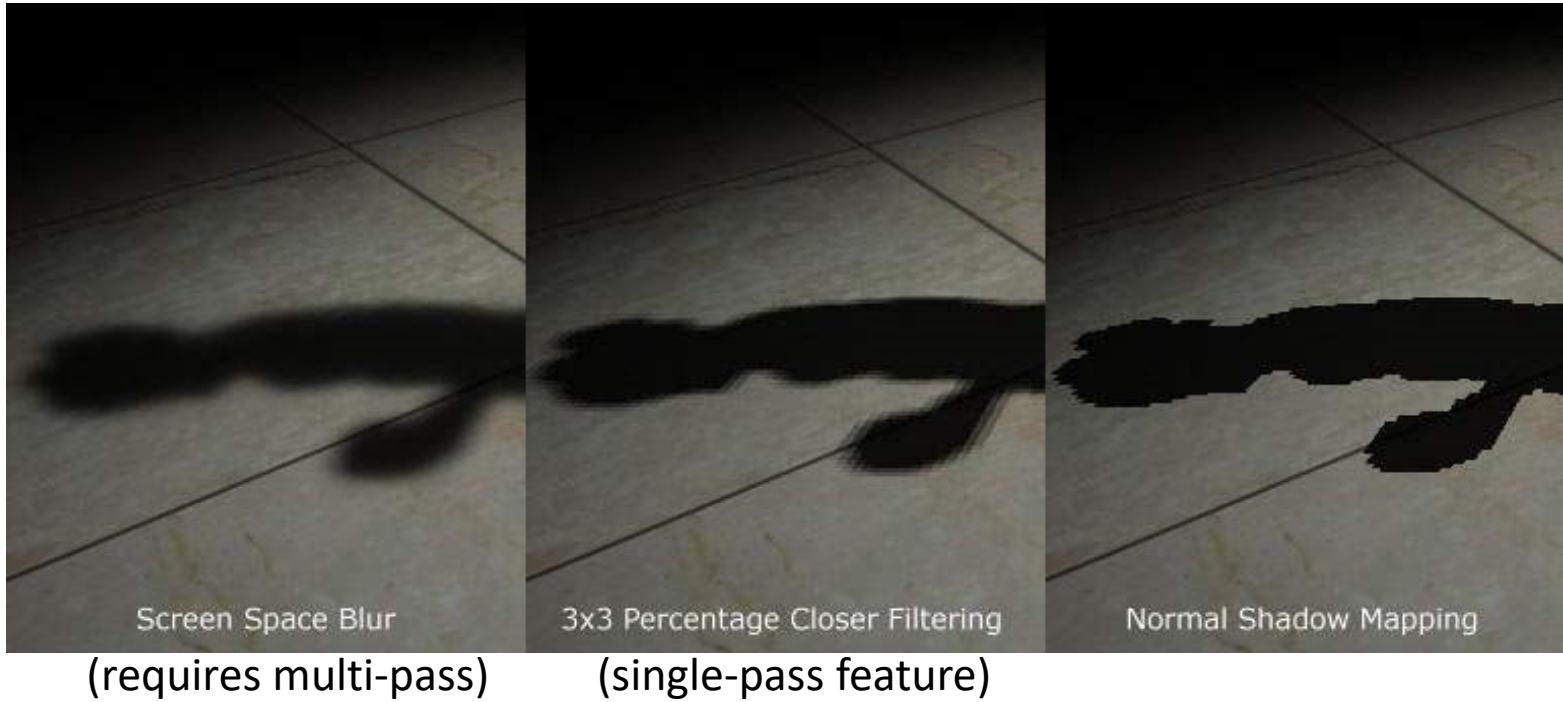


- Must filter lookup result, not depth!



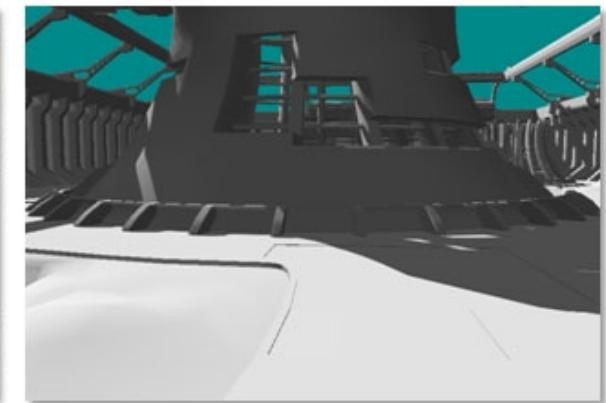
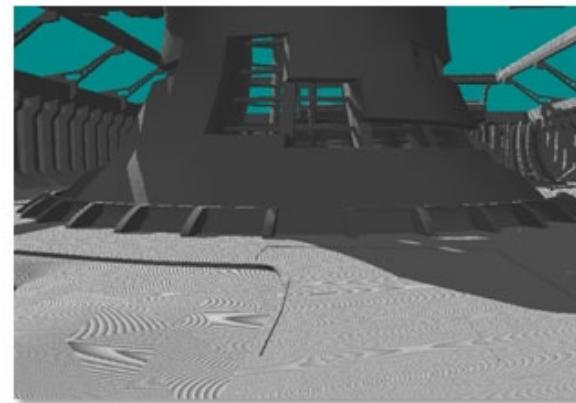
PCF Example

- <https://www.youtube.com/watch?v=qA920lufMxU>



Depth Buffer Precision Problems

- Shadow Acne
 - Depth offset too small

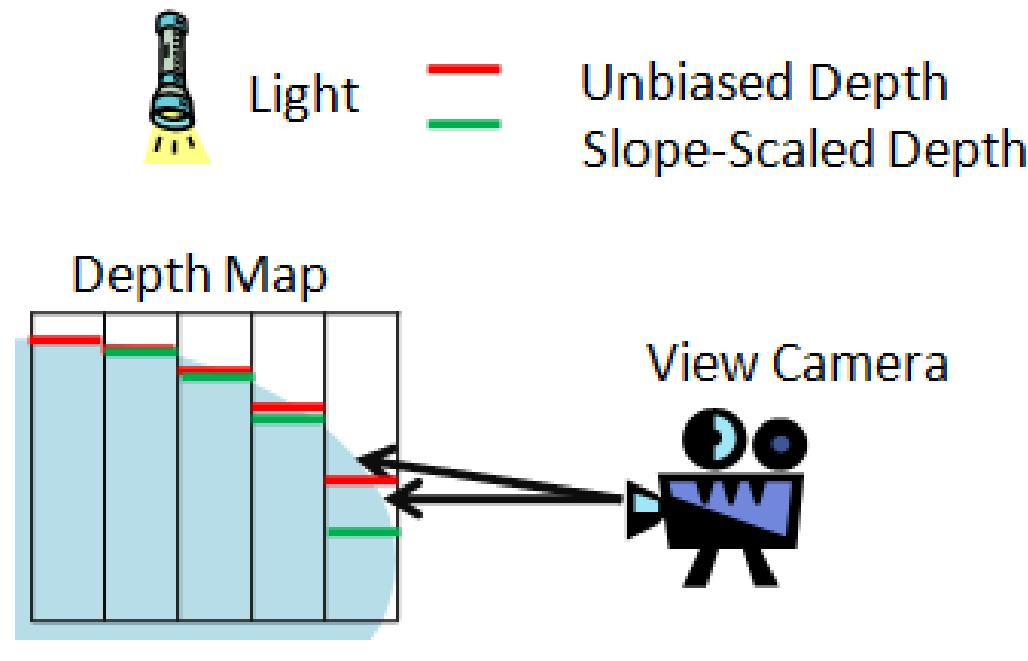


- Peter Panning
 - Shadow offset too large



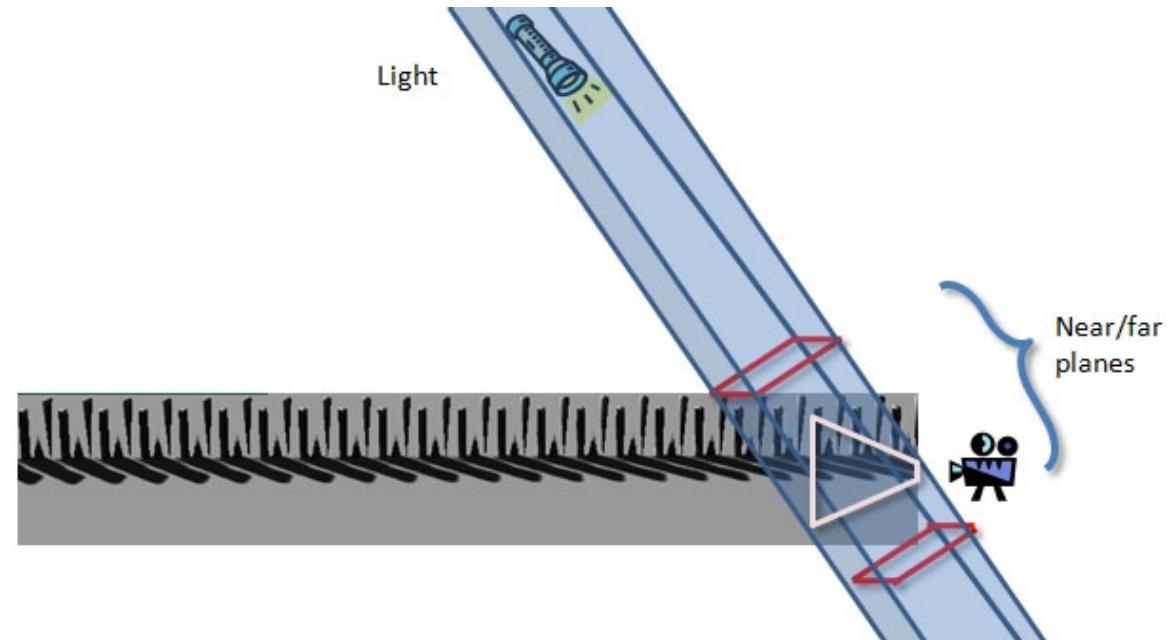
Slope Scale Bias

- Depth offset is modified by adding a bias
- Bias a polygon based on its slope with respect to the view direction

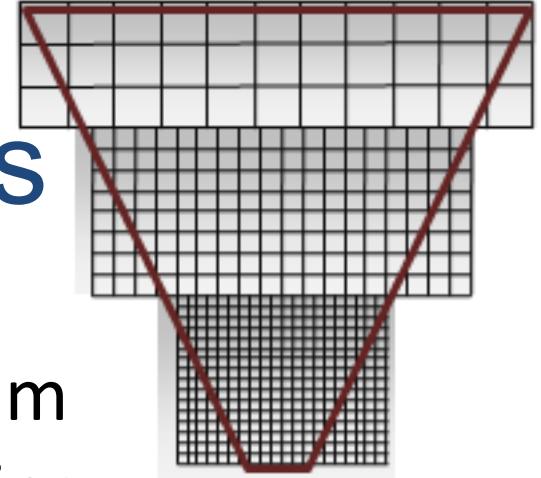


Optimize Shadow Frustum

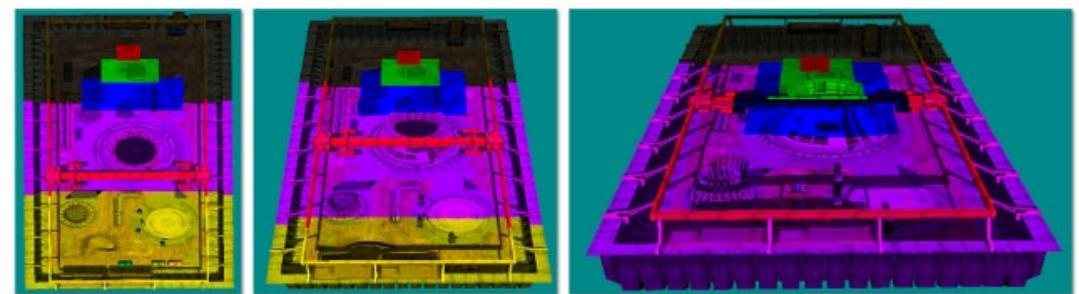
- Optimal use of depth precision: far-near → min
- Choose near/far planes based on intersection of light frustum and scene bbox



Cascaded Shadow Maps

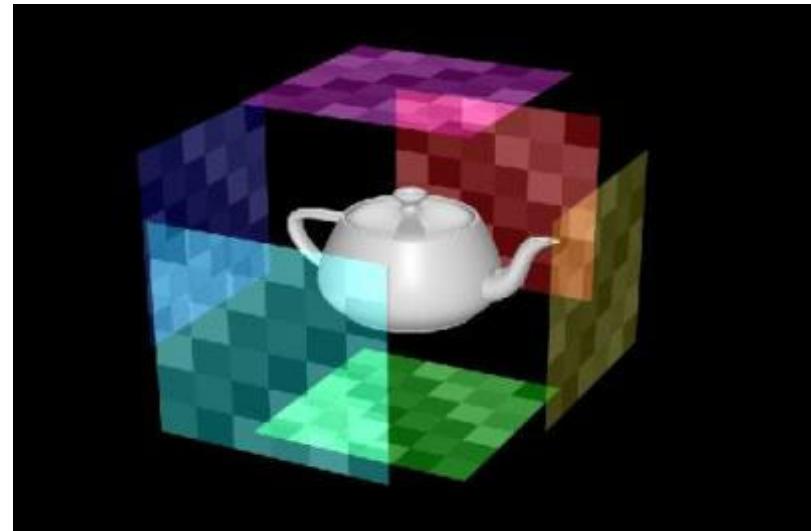


- Partition view frustum
- Generate one cascade per sub-frustum
- Simultaneous multi-viewport rendering
 - One viewport created per cascade
 - Each viewport renders to separate portion of the depth buffer
- Cascades must overlap (for blending)
- Cascade overlap increases as light direction becomes parallel with camera direction

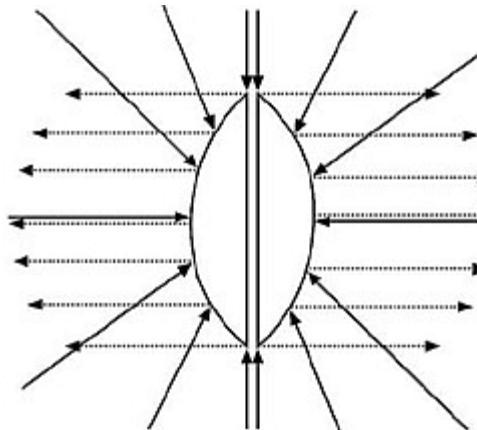
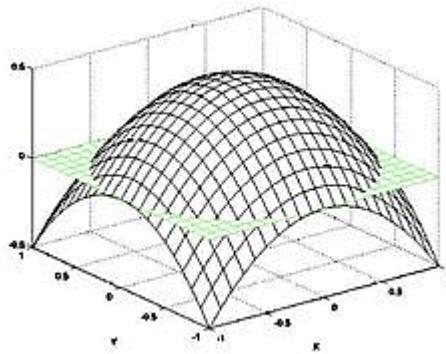


Omni-directional Shadow Maps

- Shadow maps limited to spot-light
- Cannot have a 360° “viewing frustum”
- Use six shadow maps for omni-directional light source
- Expensive!



Dual Paraboloid Shadow Map



- Paraboloid maps directions inside one half-sphere to a disc
- Arrange two such discs inside a shadow map (or environment map)
- Supported as native texture coordinate mode
- Can be generated in two passes (instead of six)

<https://www.youtube.com/watch?v=xoTMdEoMIhQ>

Shadow Maps Summary

- Advantages
 - Fast – only one additional pass
 - Independent of scene complexity
(no additional shadow polygons!)
 - Self shadowing (but beware bias)
 - Can sometimes reuse depth map
- Disadvantages
 - Problematic for omnidirectional lights
 - Biasing tweak (light leaks, surface acne)
 - Jagged edges (aliasing)

Shadow Volumes

[Frank Crow, 1977]

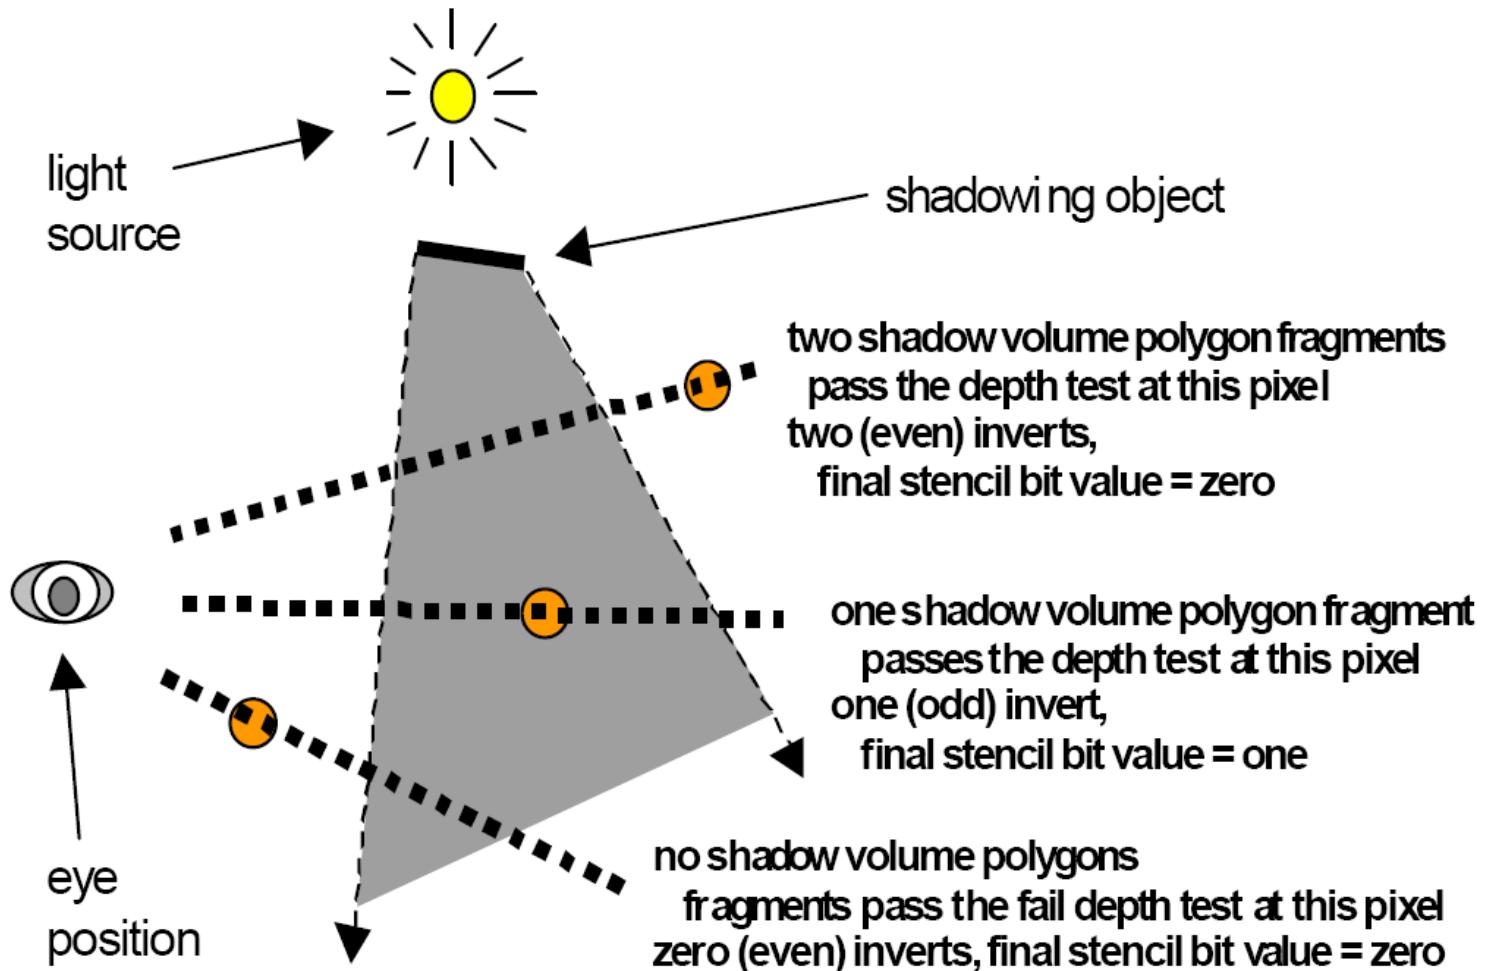
- Complex, but can be implemented efficiently using stencil buffer
- No aliasing
- Method
 - Intersect view rays with shadow volume
 - Count number of intersections, until receiver is hit



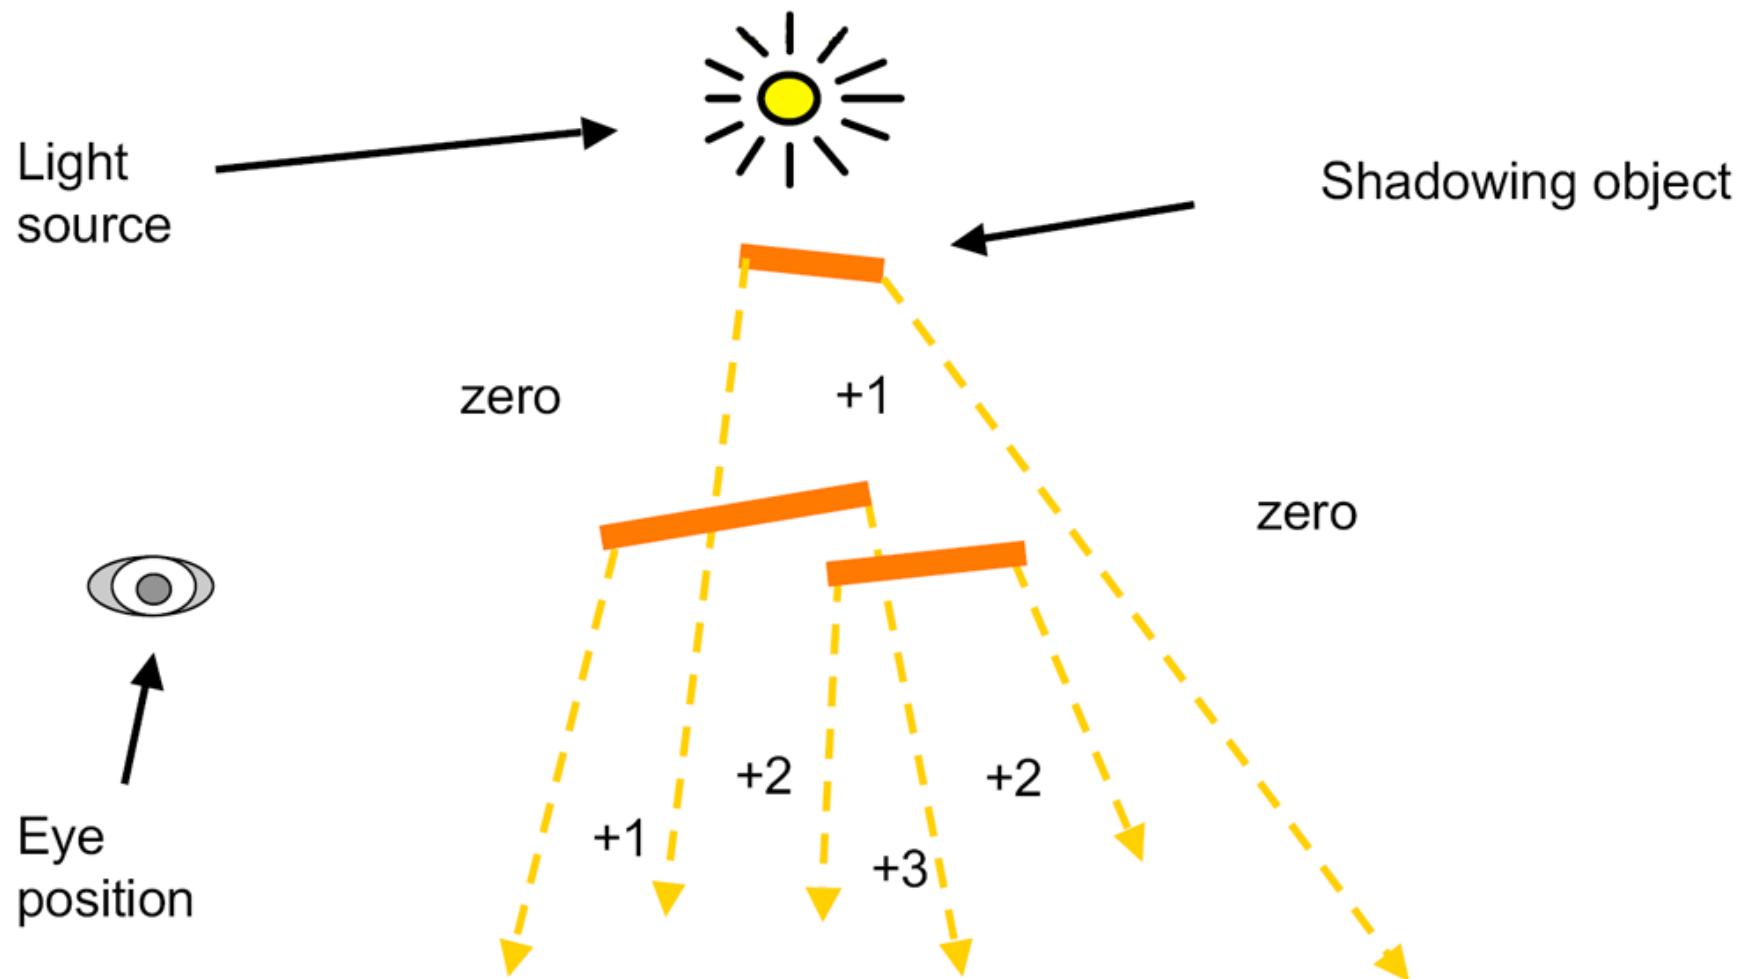
Example: Doom 3



Shadow Volumes

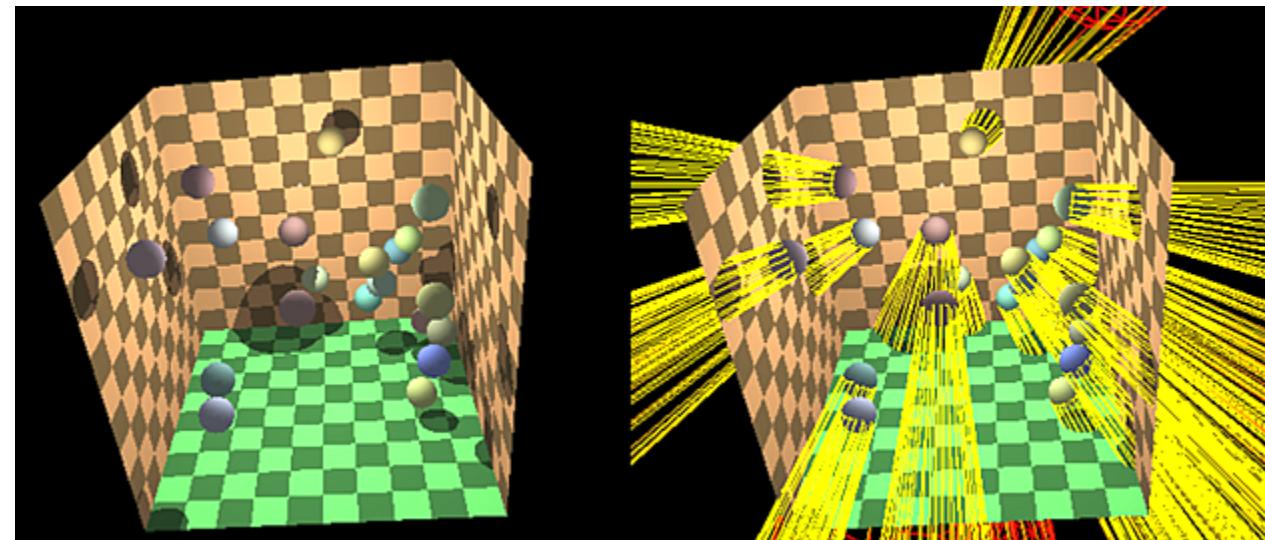


Counting Scheme



Shadow Volume Problems

- Camera inside shadow volume
 - Treat as special case or use “depth-fail” method
- Shadow volume intersects near-plane
 - Solution: render “front-caps”
- Objects must be manifold
- High amount of overdraw
- Fill-rate bound



Shadow Volume Geometry

- Closed polyhedron with 3 sets of polygons
 - Light cap
 - Object polygons facing the light
 - Dark cap
 - Object polygons facing away from the light
 - Projected to infinity (with $w=0$)
 - Sides
 - Actual extruded object edges
 - Which edges?

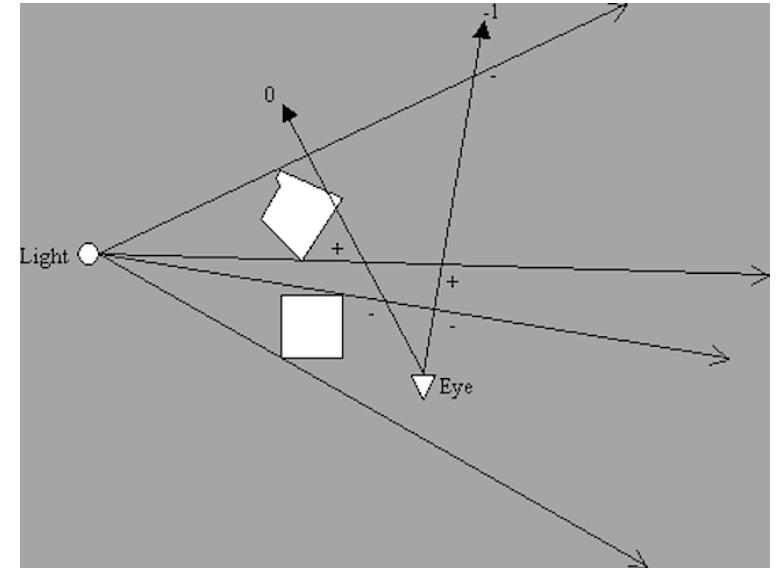
Silhouette Detection

- Classify polygons into frontfaces and backfaces
- Edge shared by front + back face is on silhouette
- Extrude silhouette edges away from light source
 - In vertex shader
- If necessary: add front-cap and back-cap



Stencil Shadow Volumes Algorithm

- Fill depth-buffer with scene
- Disable depth writes
- Render front-faces of stencil volumes with stencil increment on depth test pass
- Counts shadows in front of objects
- Render back-faces of stencil volumes with stencil decrement on depth test pass
- All lit surfaces have stencil value of 0
- Incorrect if the camera is inside a shadow volume!



Shadow Algorithms Compared

Shadow Volumes

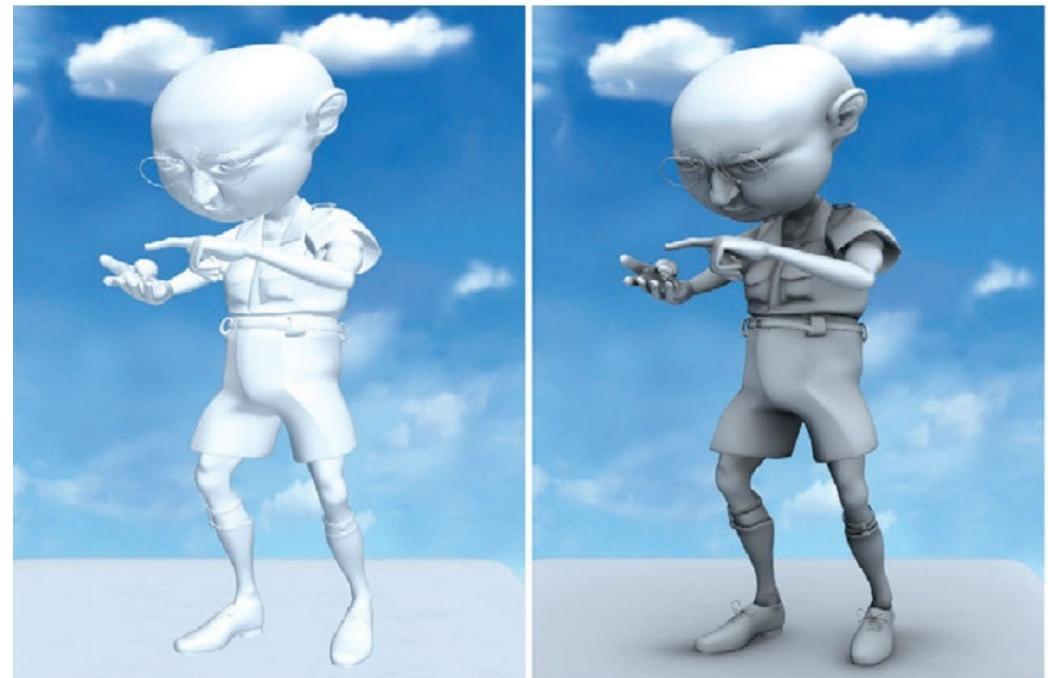
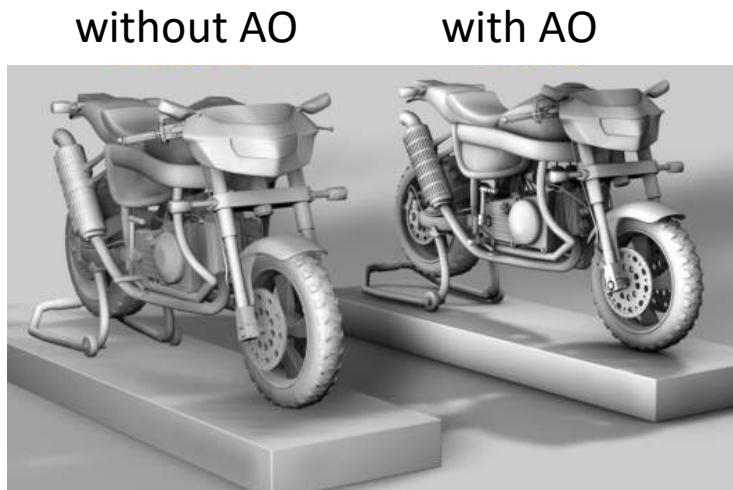
- Needs geometry
 - Closed manifolds ☺
- Cost hard to predict
 - Geometric complexity of shadow volume ☹
- Shadow rendering is bandwidth intensive ☹
 - Stencil buffer techniques
 - Additional pass
- No sampling artifacts ☺

Shadow Mapping

- Only depth needed
 - Everything one can render can cast a shadow ☺
- Predictable cost
 - Rendering the scene once ☺
- Shadow rendering is simple texture lookup ☺
 - Hardware Support
- Sampling artifacts ☹

Shadowing with Environmental Lighting

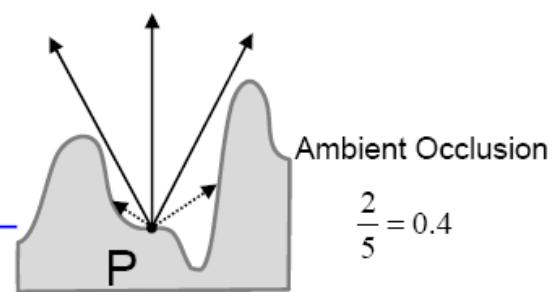
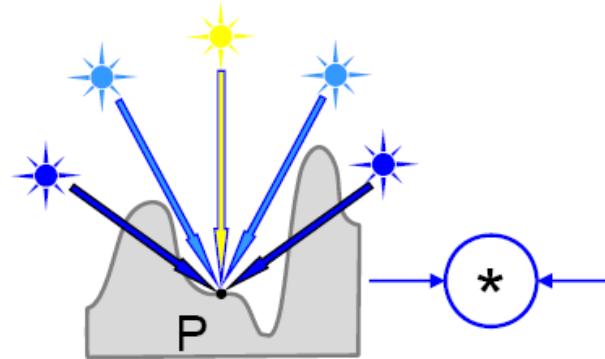
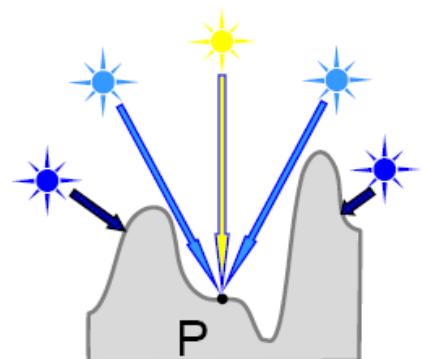
- So far: shading with known light position
- Now: shadowing by ambient light
- Independent of light/object orientation
- Darken surfaces which are partially visible in the environment
- Adds depth and contrast



Ambient Occlusion

- Compute (per vertex)
 - Mean visibility $V_{AO} = 1 - AO$
 - Bent normal vector pointing into direction of average visibility
 - Used for advanced lighting instead of surface normal
- Weigh ordinary shading using mean visibility

$$I = (1 - AO)(k_a I_a + k_d I_d \max(\mathbf{n} \cdot \mathbf{l}, 0))$$



Mean Visibility

- Pre-computation = raycasting to find self-occlusions
- For each vertex
 - Cast n rays into the half sphere oriented by the normal
 - Count blocked rays m
 - Mean visibility $V_{AO} = 1 - m/n$
- Account for Lambert's law
 - Apply cosine distribution around normal vector
- To account for other nearby objects
 - Use ordinary raycasting with maximal distance
- Problem: raycasting from every point is expensive

Ambient Occlusion



(a) Gouraud shading



(b) Ambient Occlusion

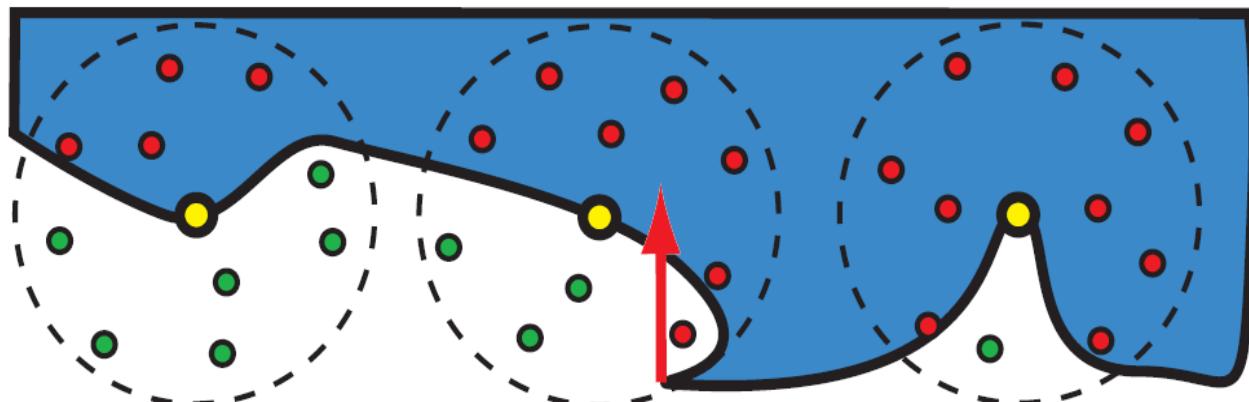


(c) Ground truth

<https://www.youtube.com/watch?v=9Fe1nYnvmiA>

Screen Space Ambient Occlusion

- Z-buffer = approximation of surrounding
- Compute ambient occlusion from neighboring pixels
- Use random sampling of z-buffer
 - V_{AO} = Ratio occluded/unoccluded samples
 - L_{in} = incoming radiance
 - ω_i = incoming (=light source) direction
 - θ_i = angle to normal



Questions?

