



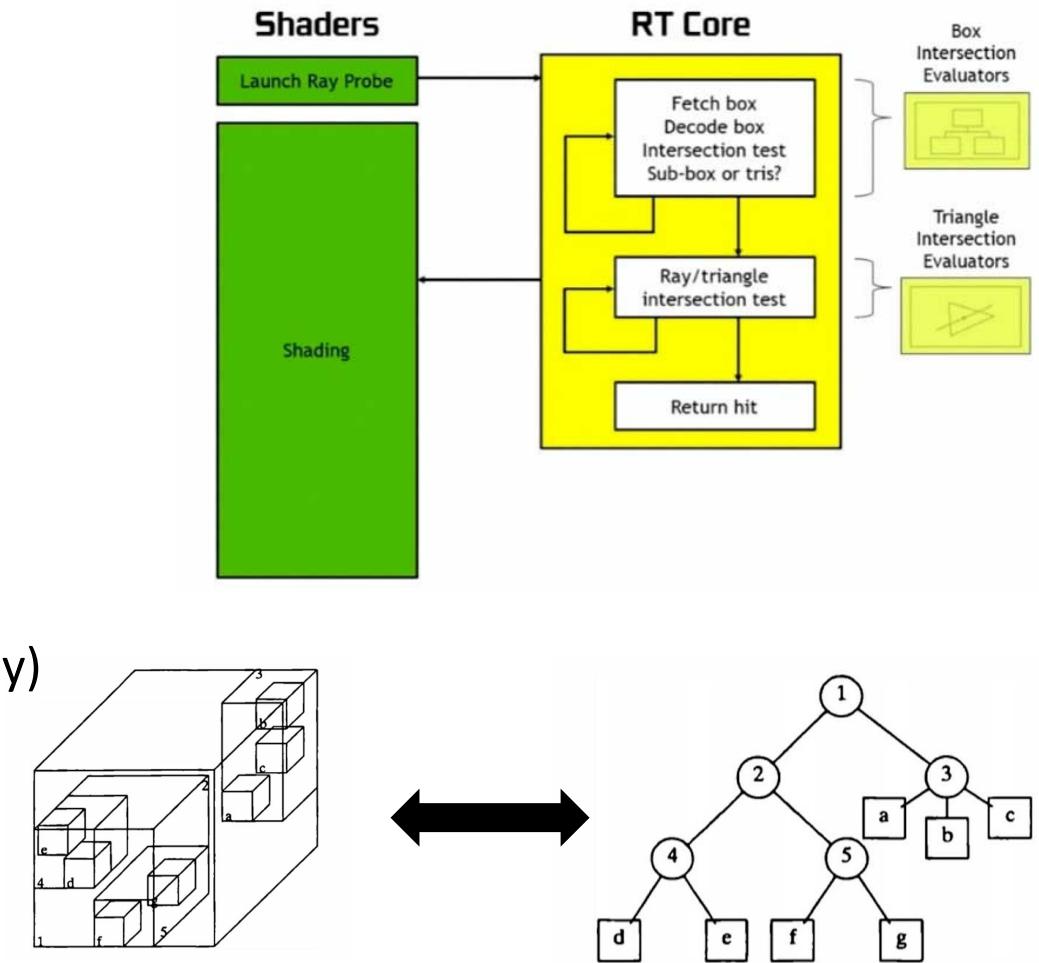
GPU Raytracing

Dieter Schmalstieg



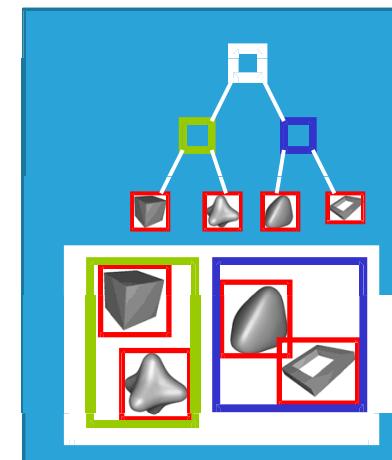
GPU Hardware

- Shader cores
 - General purpose code
- Fixed function units
 - Rasterizer
 - Texture sampler
 - Tensor core (matrix-matrix multiply)
 - ...
- RT cores (**new**)
 - Intersection ray/box
 - Intersection ray/triangle



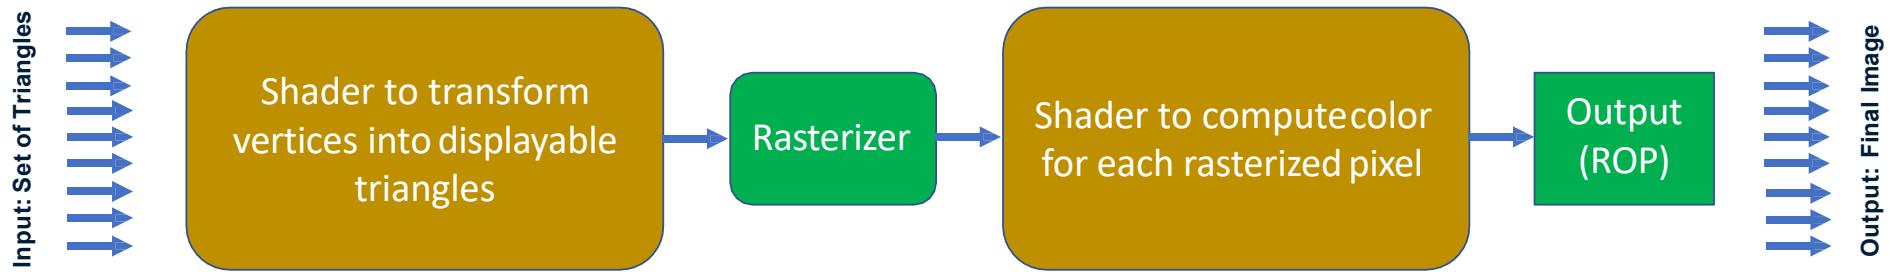
Bounding Volume Hierarchy

- Most important spatial data structure
- Common bounding volume types
 - Spheres
 - Axis-aligned bounding box (AABB)
 - Oriented bounding box (OBB)
- Encloses the bounded object
- Hierarchical data structure
 - Tree (binary or n-ary)
 - Leaves: store objects
 - Interior nodes: stores bounding volume enclosing all contained bounding volumes

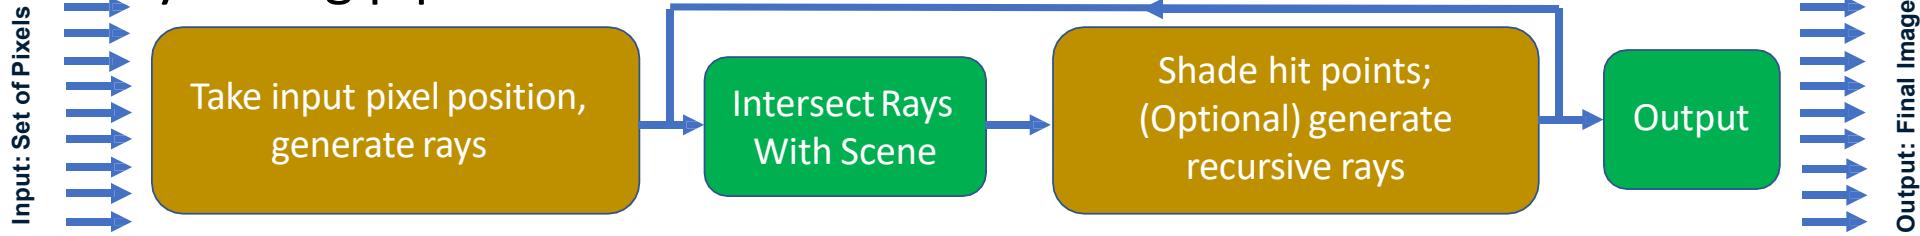


Rasterization vs Raytracing Pipeline

Rasterization pipeline

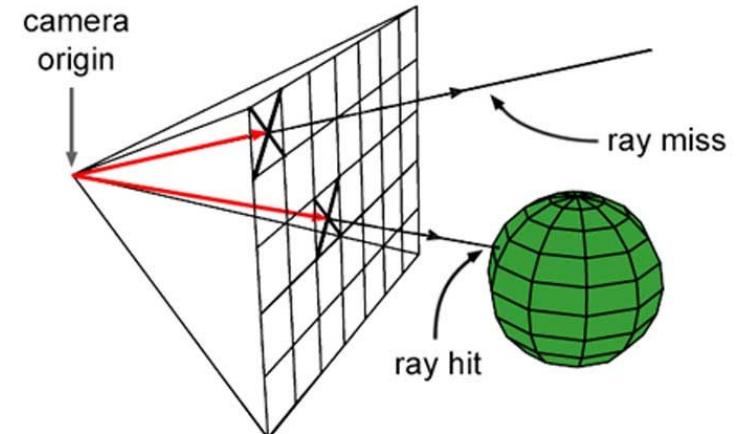
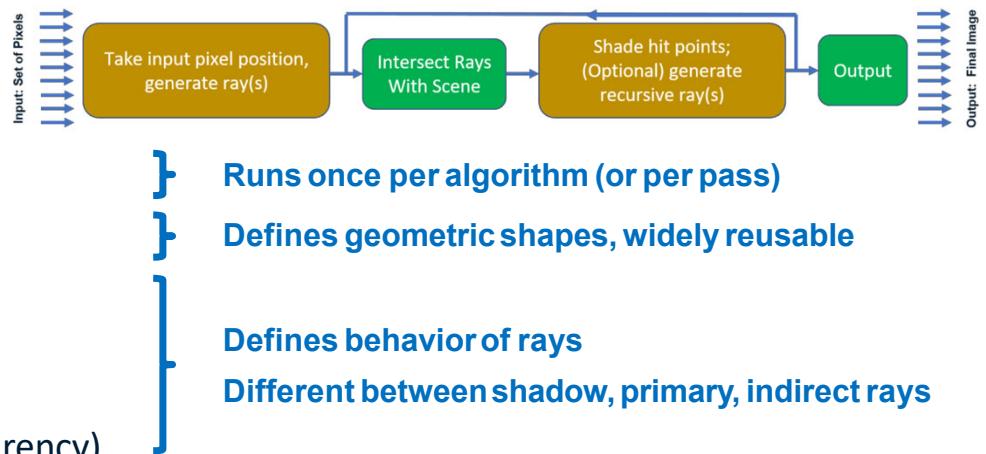


Raytracing pipeline



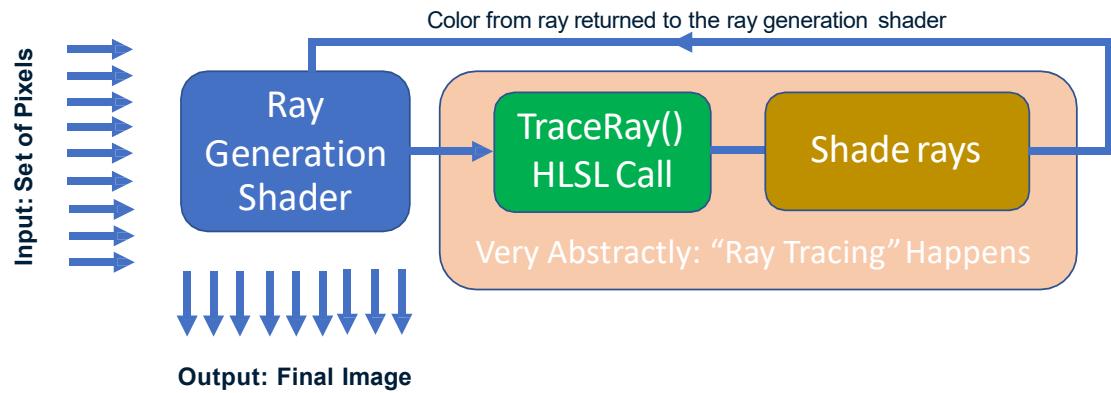
DirectX Ray Tracing Pipeline

- Pipeline is split into five new shaders:
 - **Ray generation shader** defines how to start ray tracing
 - **Intersection shader** define how rays intersect geometry
 - **Miss shader** define behavior when rays miss geometry
 - **Closest-hit shader** run once per ray (e.g., to shade final hit)
 - **Any-hit shader** run once per hit (e.g., to determine transparency)



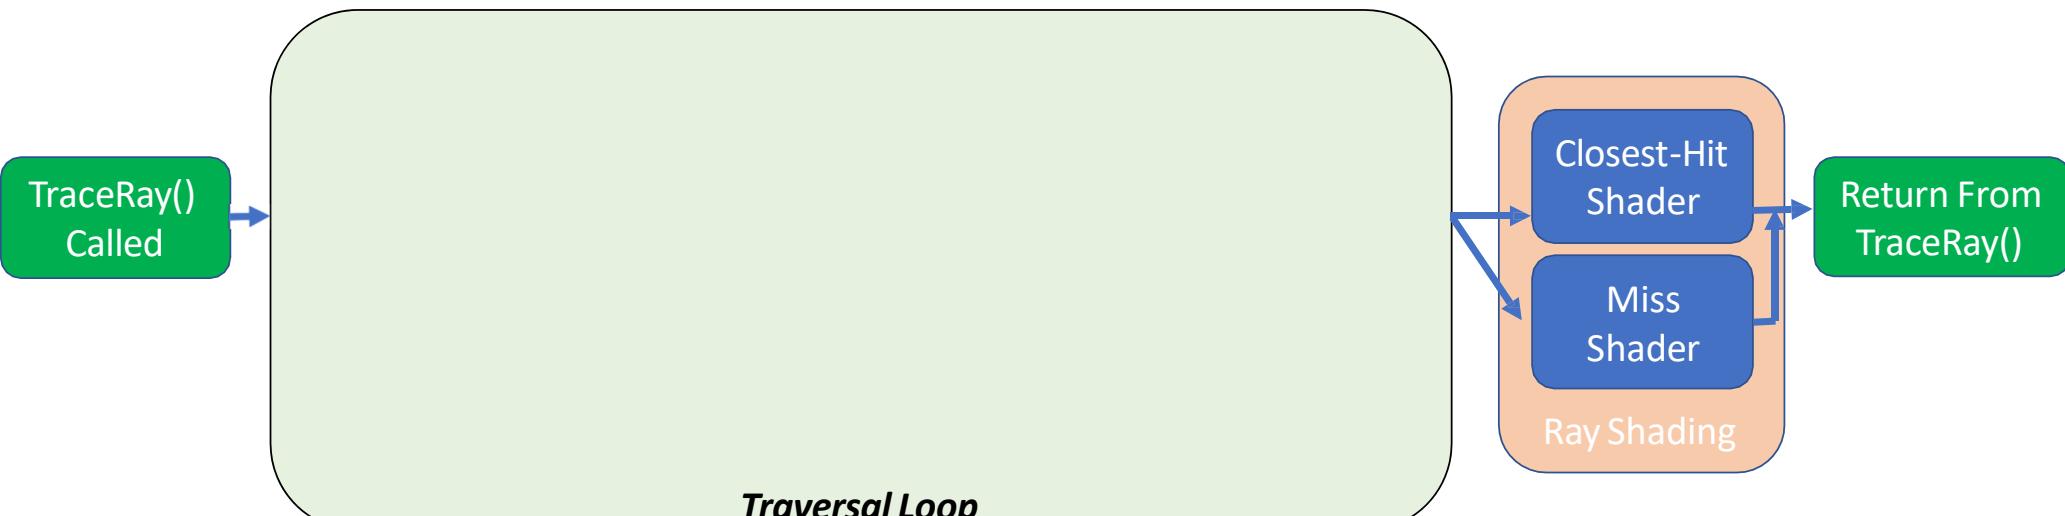
Ray Generation Shader

- Specify what rays to trace for each pixel
 - Launch rays by calling new HLSL TraceRay() intrinsic
 - Accumulate ray color into image after ray tracing finishes



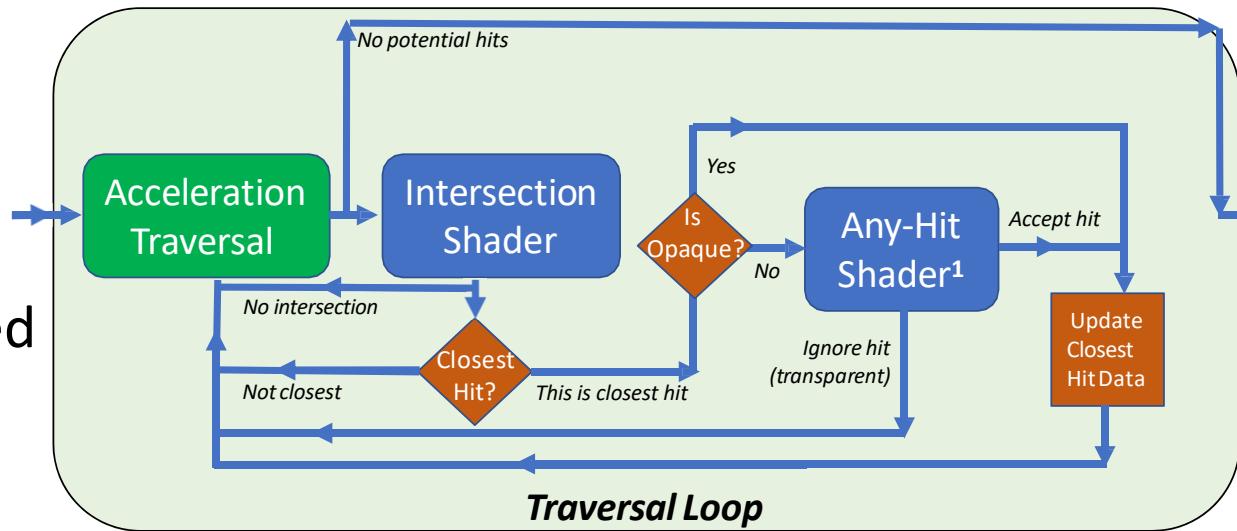
Tracing A Ray

- First, we traverse our scene to find what geometry ray hits
- When we find the closest hit, shade the point using **closest-hit shader**
- If ray misses all geometry (hits background), **miss shader** gets invoked



Traversal Loop

- If all geometry trivially ignored
 - Ray traversal ends
- For potential intersections
 - Call **intersection shader** is invoked
 - Triangles (special hardware), spheres, Bezier patches etc.
- If no intersection detected or not closest, continue traversal
- If hit
 - If transparent, run **any-hit shader** (call `IgnoreHit()` to continue)
 - Update closest hit data, continue to look for closer intersections



Starting a DXR Shader

- As any program, need an entry point where execution starts
 - Think `main()` in C/C++
- Shader entry points can be arbitrarily named
- Type specified by HLSL attribute: `[shader("shader-type")]`
 - Remember the **ray generation shader** is where ray tracing starts
- Starting other shader types look like this:
 - **RayPayload** is a user-defined (and arbitrarily named structure)
 - **IntersectAttrs** has data reported on hits (by intersection shader)

```
[shader("raygeneration")]
void PinholeCameraRayGen()
{ ... <Place code here> ... }
```

```
[shader("intersection")]
void PrimitiveIntersection ()
{ ... <Place code here> ... }
```

```
[shader("miss")]
void RayMiss(inout RayPayload data) // User-defined struct
{ ... <Place code here> ... }
```

```
[shader("anyhit")]
void RayAnyHit(inout RayPayload data,
               IntersectAttrs attrs)
{ ... <Place code here> ... }
```

```
[shader("closesthit")]
void RayClosestHit(inout RayPayload data,
                   IntersectAttrs attrs)
{ ... <Place code here> ... }
```

Ray Payload

- Ray payload is arbitrary user-defined struct
 - Contains intermediate data needed during ray tracing
 - Note: Keep ray payload as small as possible
 - Large payloads will reduce performance
 - Spills registers into memory
- A simple ray might look like this:
 - Sets color to blue when the ray misses
 - Sets color to red when the ray hits an object

```
struct SimpleRayPayload
{
    float3 rayColor;
};

[shader("miss")]
void RayMiss(inout SimpleRayPayload data)
{
    data.rayColor = float3( 0, 0, 1 ); // blue
}

[shader("closesthit")]
void RayClosestHit(inout SimpleRayPayload data,
                    IntersectAttribs attribs)
{
    data.rayColor = float3( 1, 0, 0 ); // red
}
```

Intersection Attributes

- Communications intersection information needed for shading
 - E.g., how do you look up textures for your primitive?
- Specific to each intersection type
 - One structure for triangles, one for spheres, one for Bezier patches
 - DirectX provides a built-in for the fixed function triangle intersector
 - Could imagine custom intersection attribute structures
- Limited attribute structure size: max 32 bytes

```
struct BuiltinIntersectionAttrs {
    // Barycentric coordinates of hit in
    float2 barycentrics; // the triangle are: (1-x-y, x, y)
}

struct PossibleSphereAttrs {
    // Giving (theta,phi) of the hit on
    float2 thetaPhi; // the sphere (thetaPhi.x, thetaPhi.y)
}

struct PossibleVolumeAttrs {
    // Doing volumetric ray marching? Maybe
    float3 vox; // return voxel coord: (vox.x, vox.y, vox.z)
}
```

Data Needed for GPU Raytracing

Besides our shader, what data is needed on GPU to shoot rays?

- We need somewhere to write our **output**
- Where are we looking? – **camera data**
- Need to know about our **scene geometry**
- Also need information how to shade scene
 - Depends on material format
 - Depends on shading models

```
RWTexture<float4> outTex;  
  
// HLSL “constant buffer”, populated by host C++ code  
cbuffer RayGenData {  
    float3 wsCamPos; // World space camera position  
    float3 wsCamU, wsCamV, wsCamW; // right/up/forward  
};  
  
// Our scene’s ray acceleration structure, setup via the C++ DirectX API  
RaytracingAccelerationStructure sceneAccelStruct;
```

Ray Generation Shader Code

```
[shader("raygeneration")]
void PinholeCamera() {
    uint2 curPixel      = DispatchRaysIndex().xy;           What pixel are we currently computing?
    uint2 totalPixels   = DispatchRaysDimensions().xy;        How many rays, in total, are we generating?
    float2 pixelCenter  = (curPixel+float2(0.5,0.5))/totalPixels; Find pixel center in [0..1] x [0..1]
    float2 ndc          = float2(2,-2)*pixelCenter+float2(-1,1); Compute normalized device coordinate (as in raster)
    float3 pixelRayDir  = ndc.x*wsCamU+ndc.y*wsCamV+wsCamZ; Convert NDC into pixel's ray direction (using camera)
    RayDesc ray; Setup our ray
    ray.Origin          = wsCamPos;
    ray.Direction        = normalize( pixelRayDir );
    ray.TMin             = 0;
    ray.Tmax             = 999999999;
    SimpleRayPayload payload = { float3(0, 0, 0) }; Setup our ray's payload
    TraceRay( new intrinsic function in HLSL, can be called from ray generation, miss, and closest-hit
              sceneAccelStruct, Our scene acceleration structure
              RAY_FLAG_NONE, Special traversal behavior for this ray? (Here: no)
              0xFF, mask 0xFF = test all geometry (could ignore some geometry)
              HIT_GROUP, NUM_HIT_GROUPS, MISS_SHADER, Which intersection, any-hit, closest-hit, and miss shaders to use?
              ray, What ray are we shooting?
              payload ); What is the ray payload? Stores intermediate, per-ray data
    outTex[curPixel] = float4( payload.color, 1.0f ); Write ray query result into our output texture
}
```

RayDesc is a new HLSL built-in type:

```
struct RayDesc {
    float3 Origin; // Where the ray starts
    float TMin; // Min distance for a valid hit
    float3 Direction; // Direction the ray goes
    float TMax; // Max distance for a valid hit
};
```

Hit and Miss Shader Code

- Returns red if rays hit geometry

```
[shader("closesthit")]
void RayClosestHit(inout SimpleRayPayload data,
                     BuiltinIntersectionAttribs attrs)
{
    data.color = float3( 1, 0, 0 ); // red
}
```

- Returns blue on background

```
[shader("miss")]
void RayMiss(inout SimpleRayPayload data)
{
    data.color = float3( 0, 0, 1 ); // blue
}
```

- This is a complete DirectX Raytracing shader

- Both intersection shader and any-hit shader are optional
 - Shoots rays from app-specified camera

Performance optimization

- Exploit ray coherence
 - Trace rays with similar direction at once
 - Coherent rays performing similar operations and memory access
 - Arbitrary reflection/refraction rays are bad
- Adaptive raytracing
 - Trace only where necessary
 - Shoot rays on demand
 - Filter/denoise the results

Hybrid Rasterization + Raytracing

- Pure raytracing is usually too expensive
- Game engines use a mixture
 - Main rendering done with rasterization
 - Special effects added with raytracing
- DirectX offers easy interoperability between raster and raytracing
 - Raytracing and rasterization shaders can share code and data types
 - E.g., HLSL material shaders - directly usable for a hybrid raytracing
 - Raytrace primary rays to fill G-buffer in deferred rendering
 - Raytracing of shadows, ambient occlusion, reflections

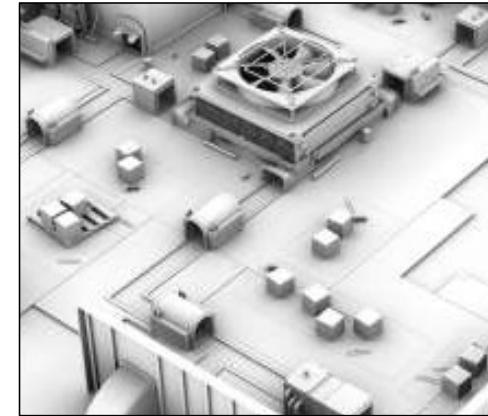
Raytraced Shadows

- Launch ray towards the light
 - If ray misses → not in shadow
 - Skip closest hit shader
 - Use miss shader
- Soft shadows
 - Choose random direction from cone
 - Cone width defines penumbra
 - Sample 1 ray per frame
 - Accumulate and filter over time (like TAA)



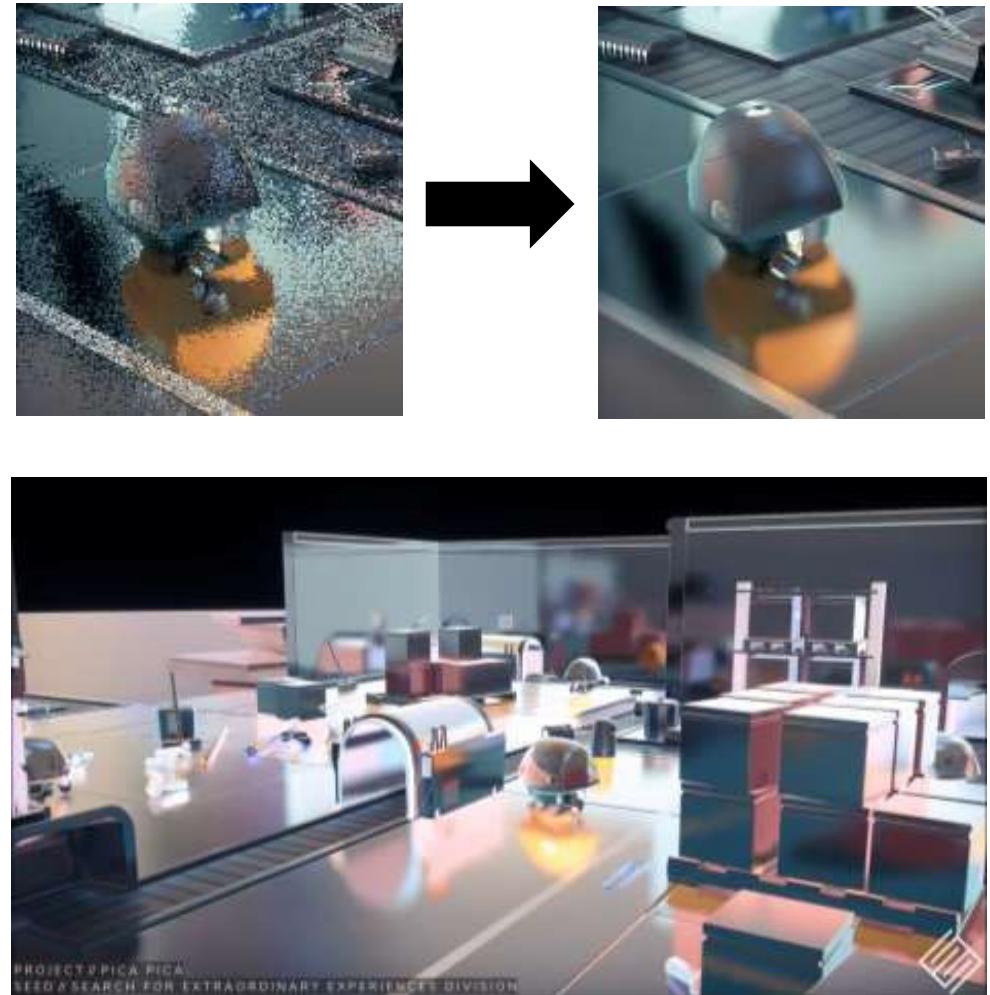
Raytraced Ambient Occlusion

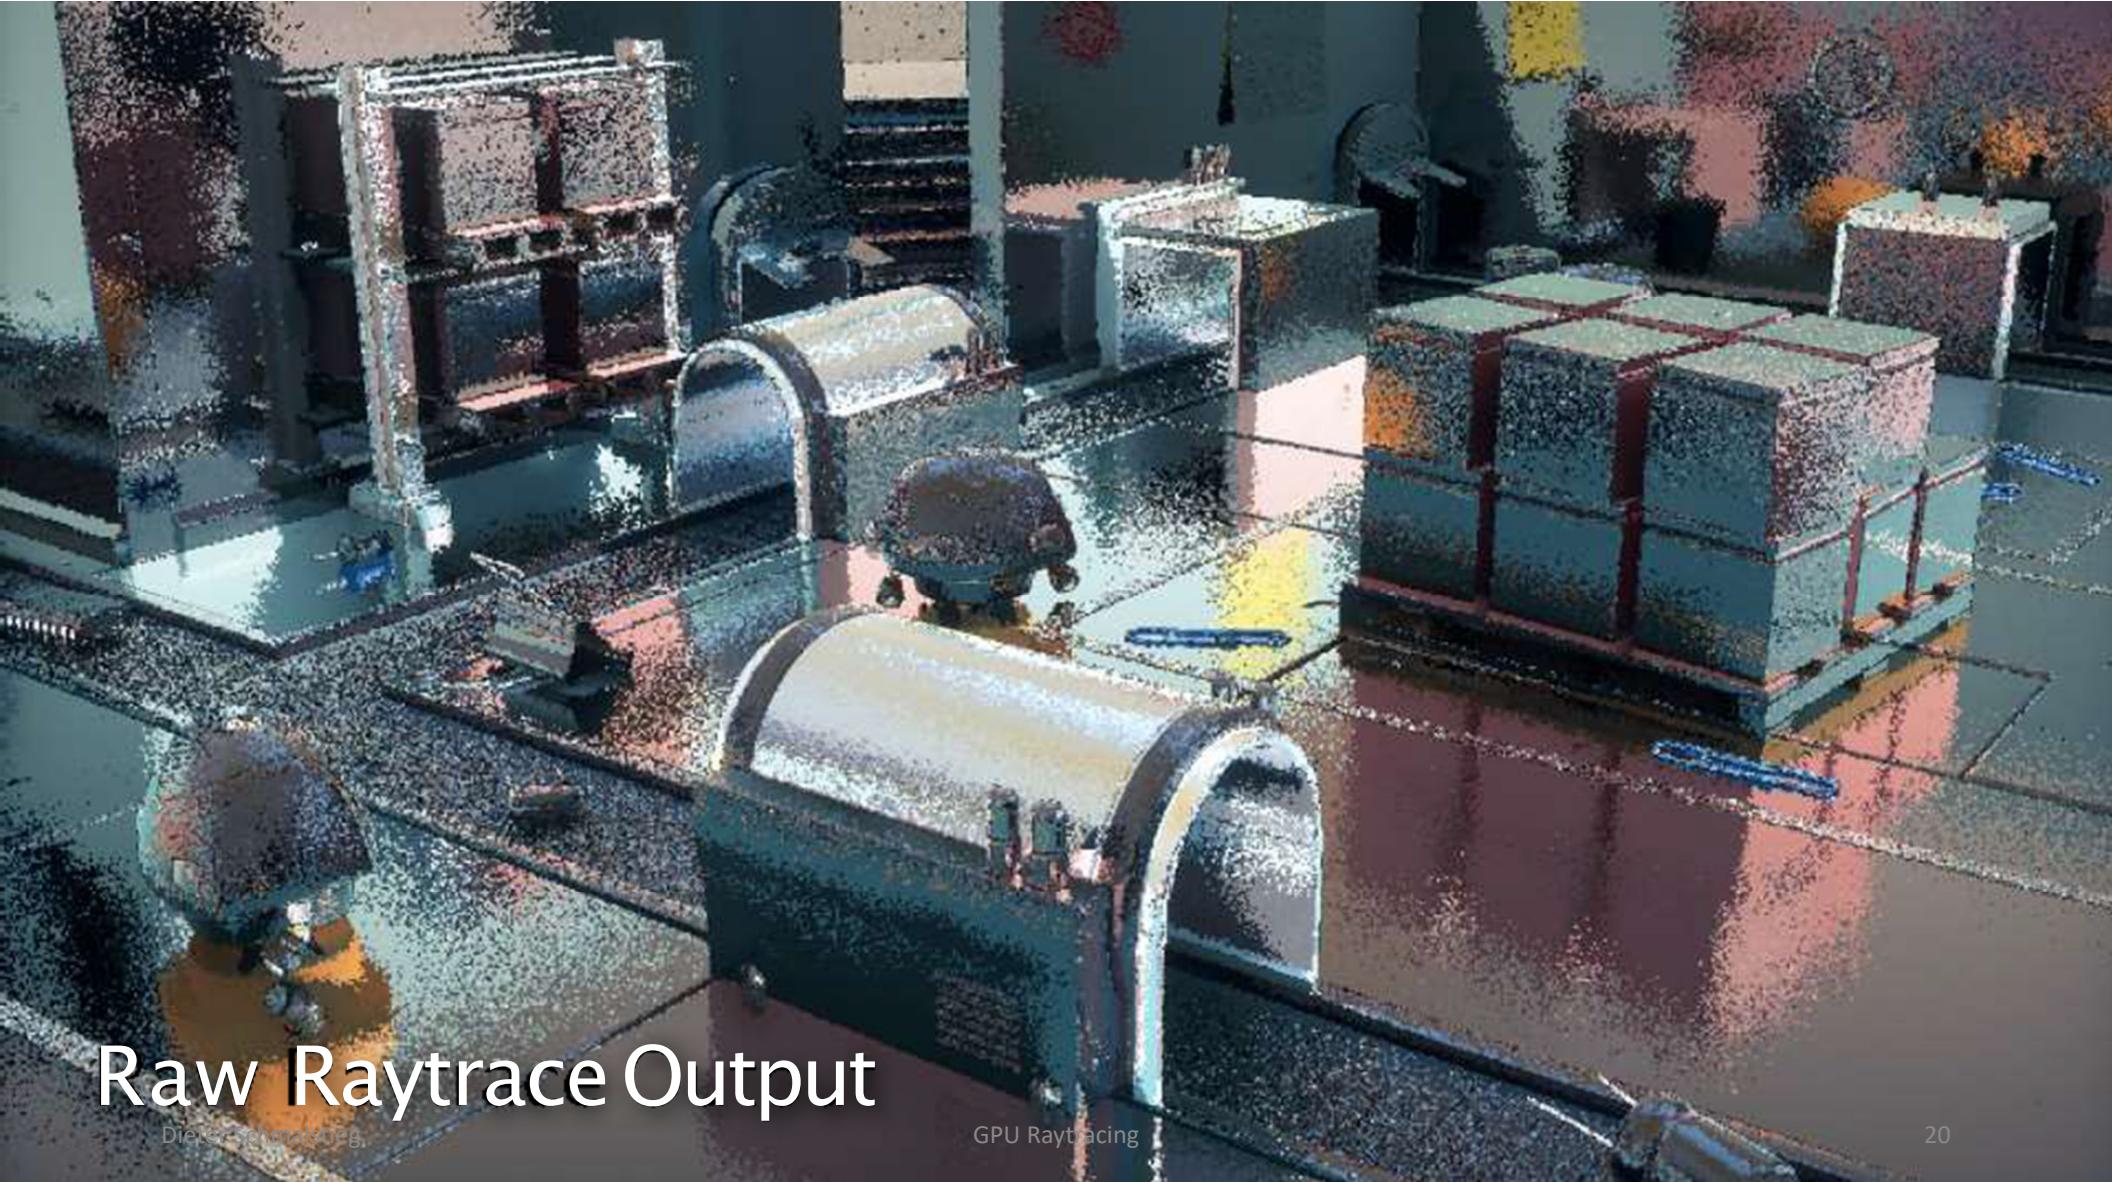
- Randomly chose directions over hemisphere at a surface point
- Like shadows, miss shader tells us if a point is occluded over not
- Compute average



Raytraced Reflections

- Launch rays from G-Buffer
 - Trace at 50% resolution
 - 25% ray/pixel for reflection
 - 25% ray/pixel for reflected shadow
- Reconstruct at full resolution
- Combine reflection methods
 - Screen-space reflections (SSR)
 - + Raytraced reflections (at SSR gaps)
 - + Environment map (at remaining gaps)
- Both spatial + temporal filtering





Raw Raytrace Output

Dieser Schmalsieg.

GPU Raytracing

20



Spatial Reconstruction

Dieter Schmalstieg

GPU Raytracing

21



+Temporal Accumulation

Dieter Schmalstieg

GPU Raytracing

22



+Bilateral cleanup

Dieter Schmalstieg

GPU Raytracing

23



+Temporal Anti-Aliasing

Dieter Schmalstieg

GPU Raytracing

24

Comparison



Real-time raytracing

Pathtracing, ~15 seconds accumulation