



Visibility

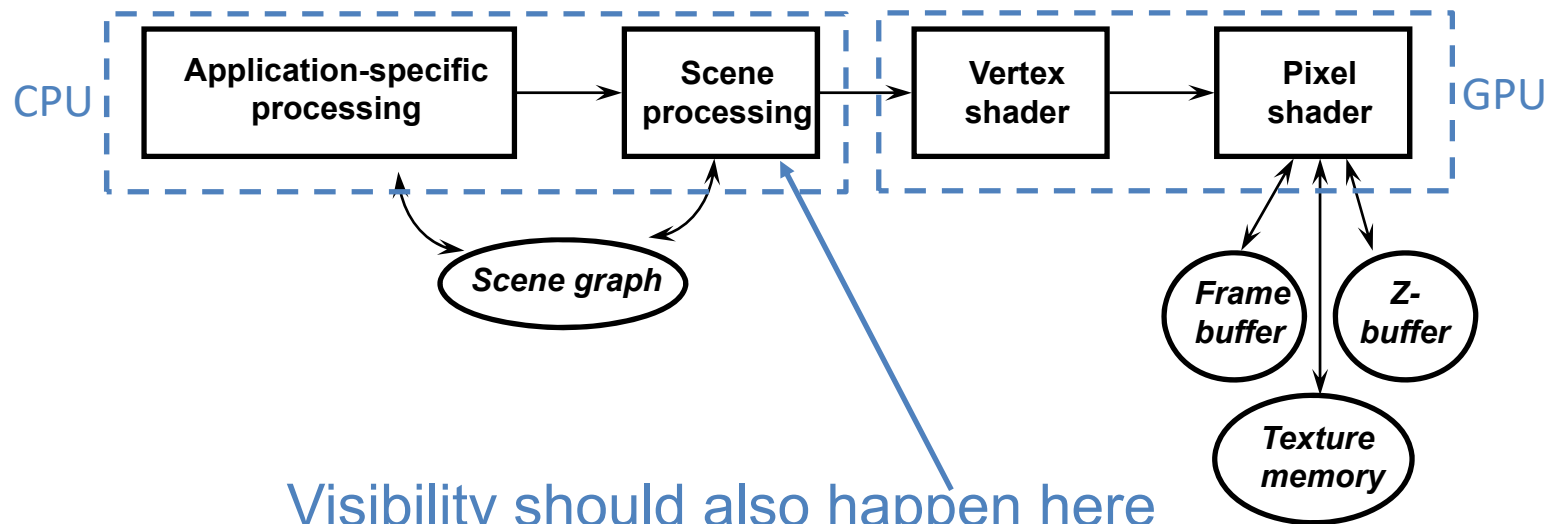
Recap: Depth Buffering

- Hardware to determine per-pixel visibility
 - 2D buffer for storing z-values per pixel
 - Pixel shader result only stored if z-value test passed
 - Allows drawing of unsorted geometry
- Sorting still greatly improves performance
 - *Early-z testing*: run depth test *before* pixel shader
 - Drawing objects close to camera early increases chances that a later pixel shader is not called

Why is the Z-Buffer Not Enough?

(Even with depth sorting:) Z-buffer...


- ...Does not eliminate depth-complexity (overdraw)
- ...Does not eliminate vertex processing of occluded polygons




Visibility should also happen here
 → we need occlusion culling before vertex shader

View Frustum Culling

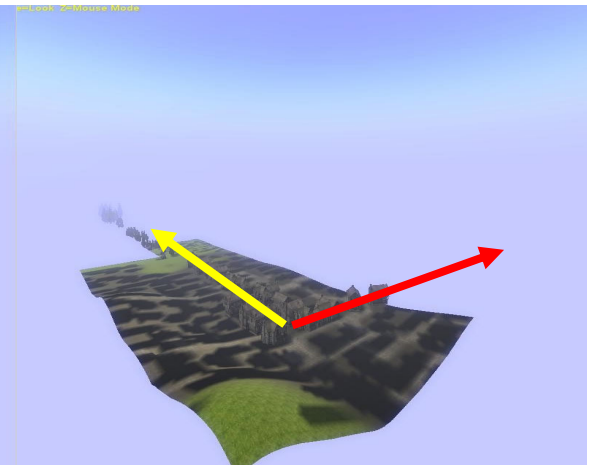
- View-frustum culling
- Occlusion culling
- Backface culling



Dieter Schmalstieg





Visibility



Bounding Volume Hierarchy

- Most important spatial data structure
- Common bounding volume types
 - Spheres
 - Axis-aligned bounding box (AABB)
 - Oriented bounding box (OBB)
- Encloses the bounded object
- Hierarchical data structure
 - Tree (binary or n -ary)
 - Leaves: store objects
 - Interior nodes: stores bounding volume enclosing all contained bounding volumes




Frustum Culling with BVH

```
Function Cull(node)
{
    if not (intersect(node, frustum) = EMPTY)
        if node = LEAF then draw(node)
        else for all children C of node
            Cull(C)
}
```


Occlusion Culling

- View-frustum culling
- **Occlusion culling**
- Backface culling




Backface Culling

- View-frustum culling
- Occlusion culling
- **Backface culling**




Backface Culling

- If an object is watertight, we cannot see the interior
- We only must draw those primitives facing the camera
- Can save up to 50% of primitives on average
- Simple to implement using dot product of normal \mathbf{n} + view vector \mathbf{v}




Visibility Culling Result

Result of frustum + occlusion + backface culling




Exactly Visible Set

- All primitives that are visible
- No more, no less




(a)

(b)



Potentially Visible Set

- PVS = all primitives that are visible
- And a bit more
 - E.g., visible within the next n frames
 - Exact hidden surface removal done later by z-buffer
- How much more?
- What do we want to do with the PVS?



PVS Classification


- The exact visible set (EVS) is unknown: $O(N^9)$...
- Potentially visible set (PVS) = set of objects that *could* be visible
- PVS can be
 - Aggressive, $PVS \subseteq EVS$
 - Conservative, $PVS \supseteq EVS$ (preferred)
 - Approximate, $PVS \sim EVS$
- PVS can be precomputed
 - But we need discretize viewpoints into view *regions*

Naïve Occlusion Culling

- Select some promising occluders
 - Large objects
 - Close to camera → large in screen space
- Test all other objects against occluders
 - Remove the objects that are occluded
 - Maybe speed up by testing bbox instead of object
- Why does it not work well for real-life scenes?


Visibility from a Point

- Terms: occluder, occludee, shadow volume, umbra




Visibility from a Point

- Complete shadow volume for occluders $occ_1, \dots, occ_n = \mathbf{union}$ of all individual shadow volumes



Occluder Fusion

- **Occluder fusion:** combined effect of multiple occluders



Visible without occluder fusion
Invisible with occluder fusion

Occlusion Culling in Practice


- Use two two spatial data structures
 - 1. Scene data structure (SDS)**
 - Stores the objects in the scene
 - 2. Shadow volume data structure (SVDS)**
 - Generated by occluders (selected from scene), or
 - Generated by virtual occluders (synthesized)
- Cull SDS using SVDS

Simple Algorithm for Point Visibility


- Shadow volume data structure (*SVDS*) = empty
- For each occluder occ_i
 - Calculate shadow volume SV_i
 - Add SV_i to *SVDS*
- For every object o_j
 - Test o_j against the *SVDS*
 - Cull o_j if occluded

Spatial Data Structures

- For quickly finding/culling large portions of scene with a single test
- All kinds of (hierarchical) data structures used
- E.g., bounding boxes, bounding volume hierarchy, grids, quadtree, octree, k-d tree, BSP tree





Simple 3D grid



Example: Octree

- 3D equivalent of quadtree
- Hierarchical subdivision of a cube into 8 octants

Region in 3D space




Simple Algorithm, Revisited

- Shadow volume data structure (*SVDS*) = empty
- For each occluder occ_i
 - Calculate shadow volume SV_i
 - Add SV_i to *SVDS*
- For every object o_j
 - Test o_j against the *SVDS*
 - Cull o_j if occluded




What shall we use for *SVDS*?

Cells and Portals

- Indoors
 - Most rooms (*cells*) occluded by walls
 - Store *portals* (windows, doors) instead of occluders (walls)
- Cells and portals form nodes and edges of a *portal graph*




Adjacency graph



Screen-Space Portals

- In runtime, find the portals of the current cell that are in the frustum
- Traverse through all found portals to the adjacent cells and find all portals that are visible to the camera through the original portal



Regular Depth Buffer

- *Depth Prepass* in deferred rendering
 - Pass 1: Rasterize the geometry,
only write depth + object id to G-buffer
 - Pass 2: Read + collect all visible ids from G-buffer
- Can be expensive
 - For large scenes with many primitives
 - For high framebuffer resolutions


Virtual Occluders

- Expensive to use objects with many primitives as occluders
- Cannot use *bounding* volumes of objects, since this is not a conservative test
- But we can use *bounded* volumes completely contained inside objects as *virtual occluders*
- Virtual occluders can be simple boxes




Hierarchical Depth Buffer

- Regular depth buffer can be expensive for high resolution
- Replace depth buffer with a depth *pyramid*
 - Bottom of the pyramid: full-resolution depth buffer
 - Higher levels: smaller resolution depth buffers, where a pixel represents max. z-value of group of pixels on lower level
- Hierarchically rasterize polygon, starting from highest level
 - If polygon is further than recorded pixel, early exit
 - If polygon is closer, hierarchically test lower levels
 - If bottom of pyramid is reached and polygon still closer, propagate the value up the pyramid




Visibility from a Region (Example in 2D)




- If we want to precompute visibility, we must discretize viewpoints into view regions
- Only works for static scenes

Umbra and Penumbra


- **Umbra** = full shadow, penumbra = half-shadow
- **Umbra** is a simple in/out classification
- **Penumbra** additionally encodes which parts of the viewing region are visible




Visibility from a Region



Visibility from a Region




Visibility from a Region




Visibility from a Region

- Area XXX is always occluded → complete shadow volume is **more** than the union of individual shadow volumes






Virtual Occluders for Regions

- Shadow volume data structure = empty
- For each occluder occ_i in front-to back order
 - Expand occluder inside existing shadow volume as far as possible
 - Calculate shadow volume SV_i
 - Add SV_i to shadow volume data structure
- Test the scene against the SVDS



Virtual Occluders from Fusion

- Occluder fusion via intermediate data structure
- Search inside umbra for adjunct occluders and add it to umbra
- Virtual occluder is guaranteed to be occluded from cell
- Virtual occluder grows large



V. Koltun, Y. Chrysanthou, and D. Cohen-Or, "Virtual Occluders: An Efficient Intermediate PVS Representation," Rendering Techniques 2000: Proc. 11th Eurographics Workshop Rendering, pp. 59-70, June 2000.

Questions?